

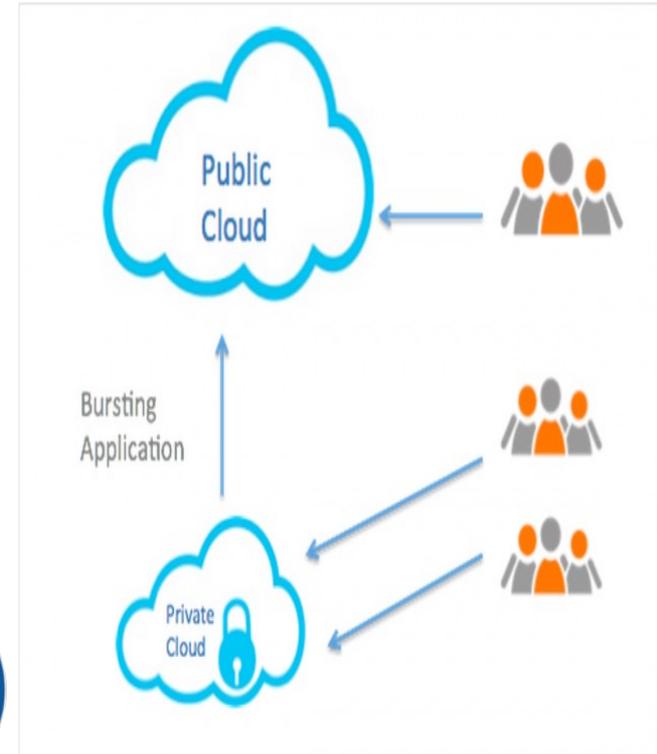
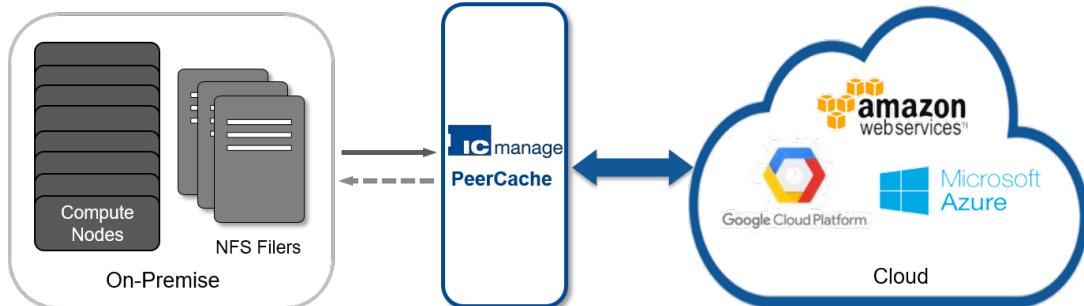


# **History & Building Blocks of Cloud Computing**

## (Mastering Cloud Computing: Chapter#1)

# What is cloud bursting?

Cloud bursting is an application deployment model in which an application runs in a private cloud or data center and bursts into a public cloud when the demand for computing capacity spikes. The advantage of such a hybrid cloud deployment is that an organization only pays for extra compute resources when they are needed.



# Distributed Systems

A distributed system is a collection of independent computers that appears to its users as a single coherent system

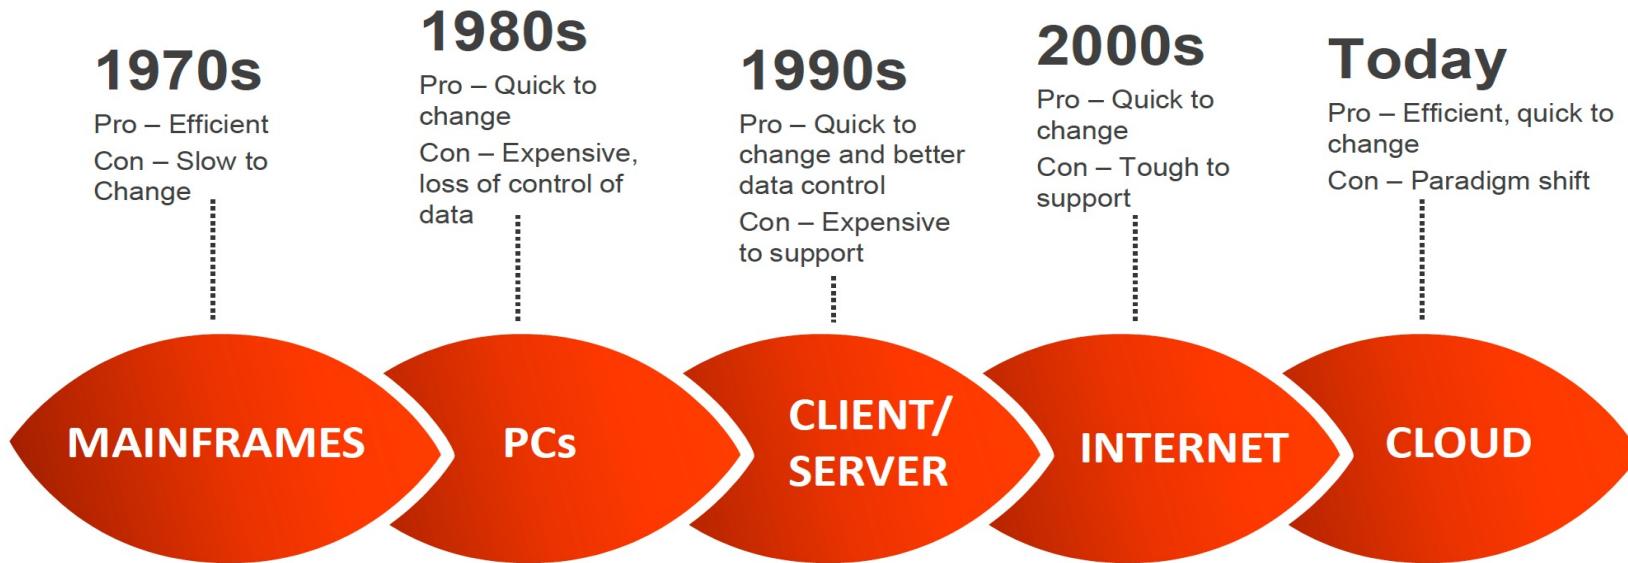
- Tanenbaum

- **To share resources for better utilization.**

CEAS offers a course in Distributed Systems!

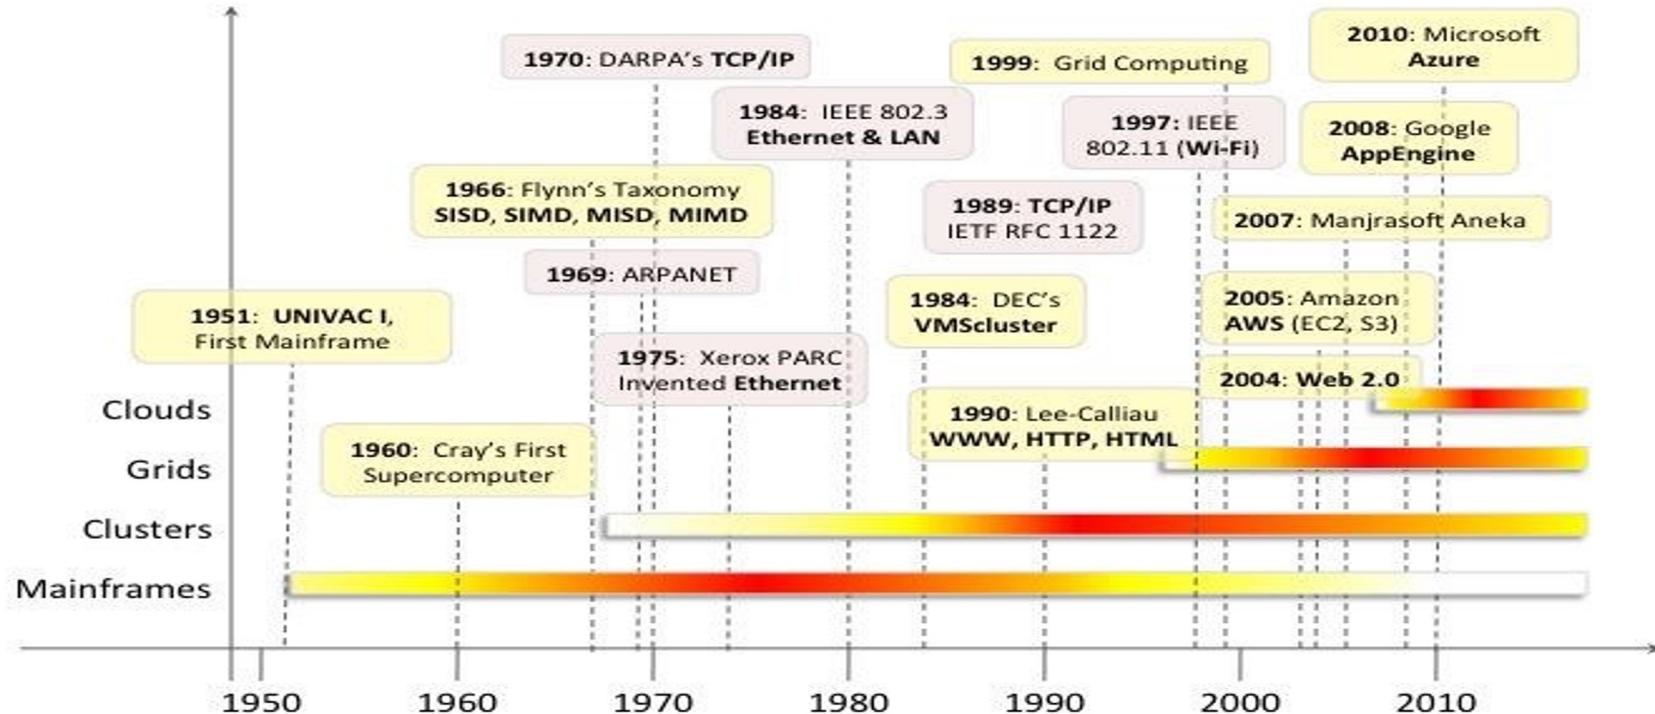
# How did we get here?

**Persistent Goal: Efficient systems that respond to the business**



“Cloud computing is based on the time-sharing model we leveraged years ago before we could afford our own computers. The idea is to share computing power among many companies and people, thereby reducing the cost of that computing power to those who leverage it. The value of time share and the core value of cloud computing are pretty much the same, only the resources these days are much better and more cost effective.” – David Linthicum

# Cycle (*what is old is new again...*)



# Mainframes (1951 onward)

- Multiple processing units
- Powerful, reliable, IO optimized
- Time Sharing systems
- Replacement of components while running - always on
- Still used for transaction processing: banking, airline ticketing, registration

- IBM z990
- 2003
- 256 GB RAM
- 1k+ VM Linux
- 11k SSL con/sec



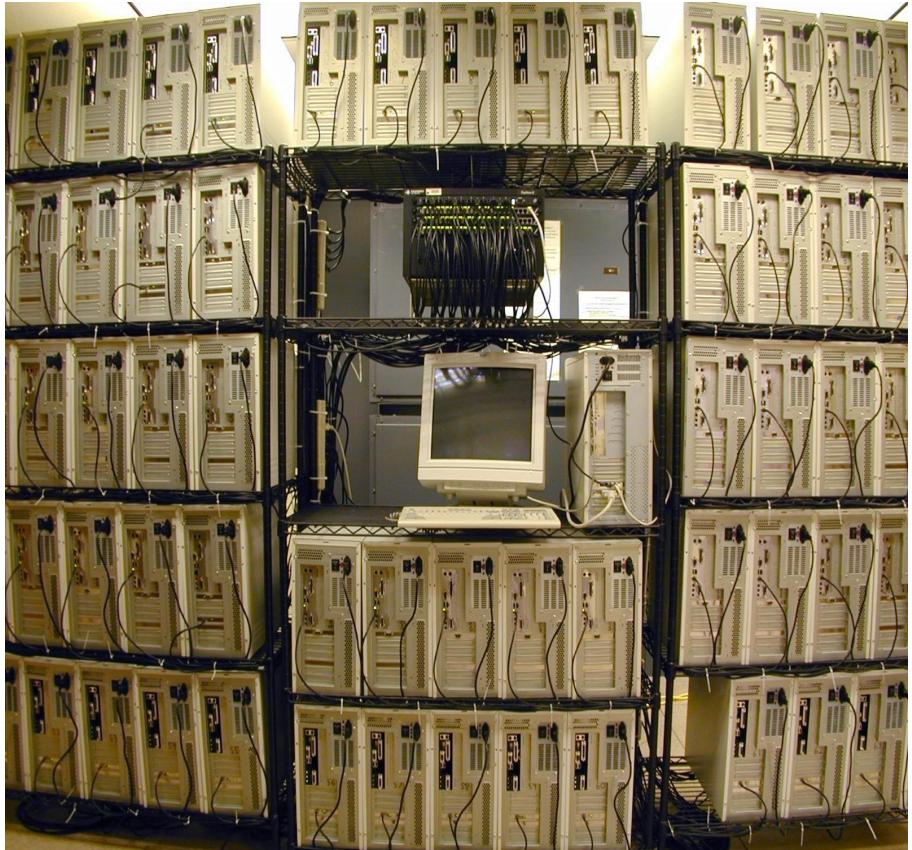
- IBM EC12
- 2012
- 101 CPUs
- 5.5Ghz hex core
- Integrated SSDs
- 6k Linux VMs

# Clusters (*1968 onward*)

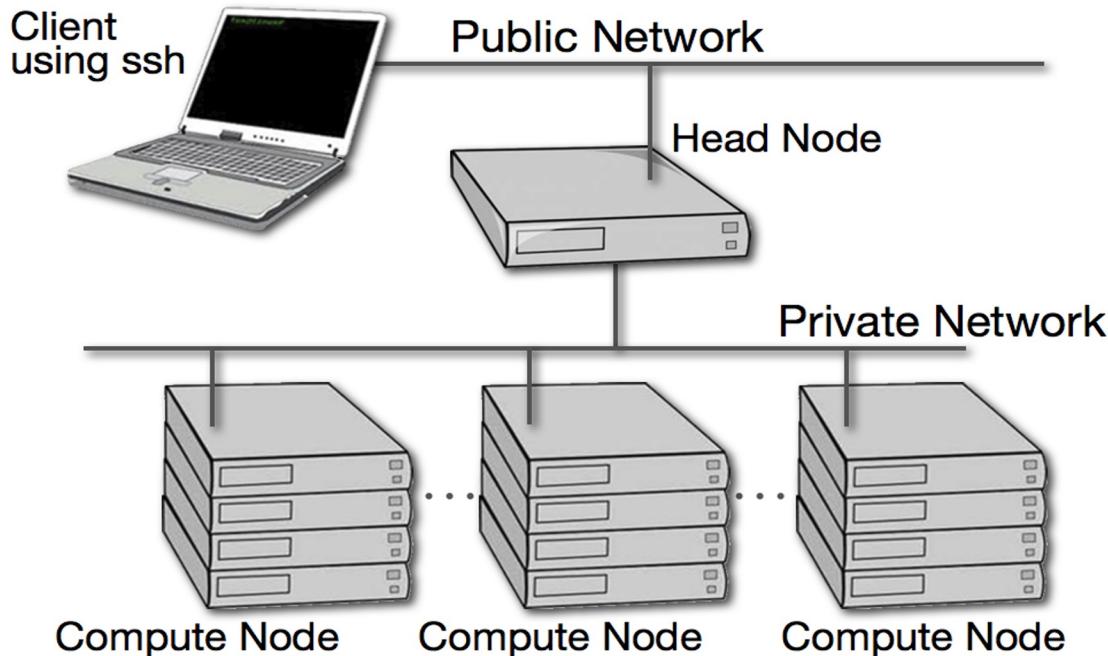
- Networked commodity (cheap) machines
- Physically close (room/building/LAN)
- Good for processing work, not IO
- Easy to expand
- Some use spare CPU cycles
- [Condor](#), [Beowulf Clusters](#), [Message Passing Interface \(MPI\)](#)
- [Ohio Supercomputer](#)

# What is a Beowulf Cluster

- A [Beowulf](#) Cluster is one class of a cluster computer
- Uses Commercial Off The Shelf (COTS) hardware built in 1994 at [NASA](#)
- Typically contains both master and slave nodes
- Not defined by a specific piece of hardware
- As of 2014 Beowulf systems operate worldwide, chiefly in support of [scientific computing](#).



# Cluster Components



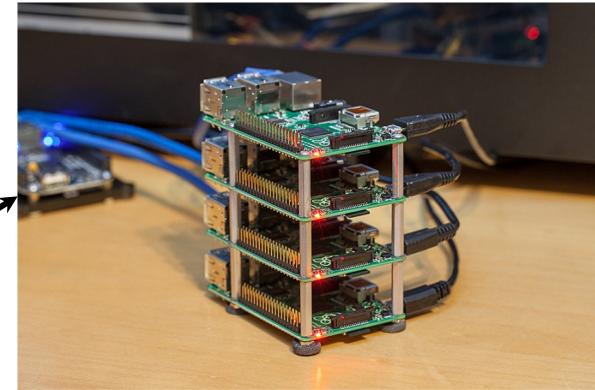
- Processors
- OS
- Network / Interfaces
- Cluster middleware
  - Beowulf...
  - Rock...
- Execution Environment
  - PVM/MPI

**Compute node can also be called as worker node**

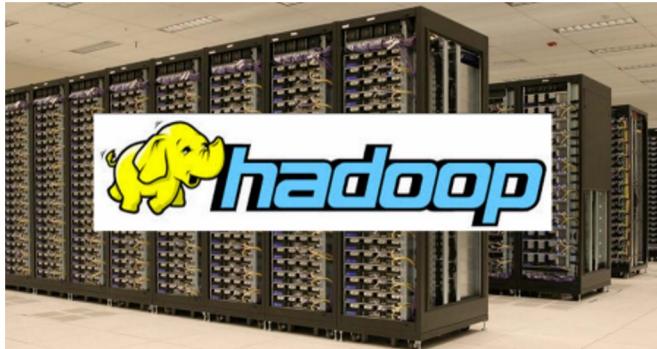
# Different Clusters



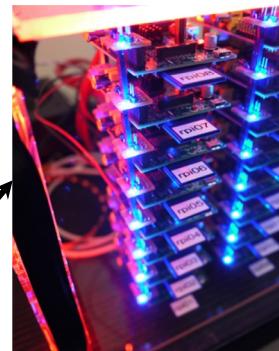
Playstation 3  
Supercomputing  
Cluster



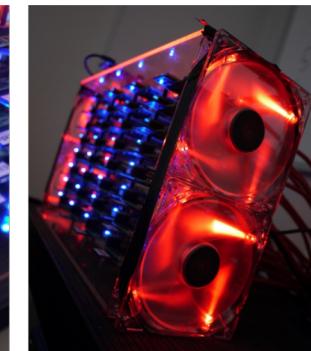
Raspberry Pi  
Cluster



Hadoop  
Cluster



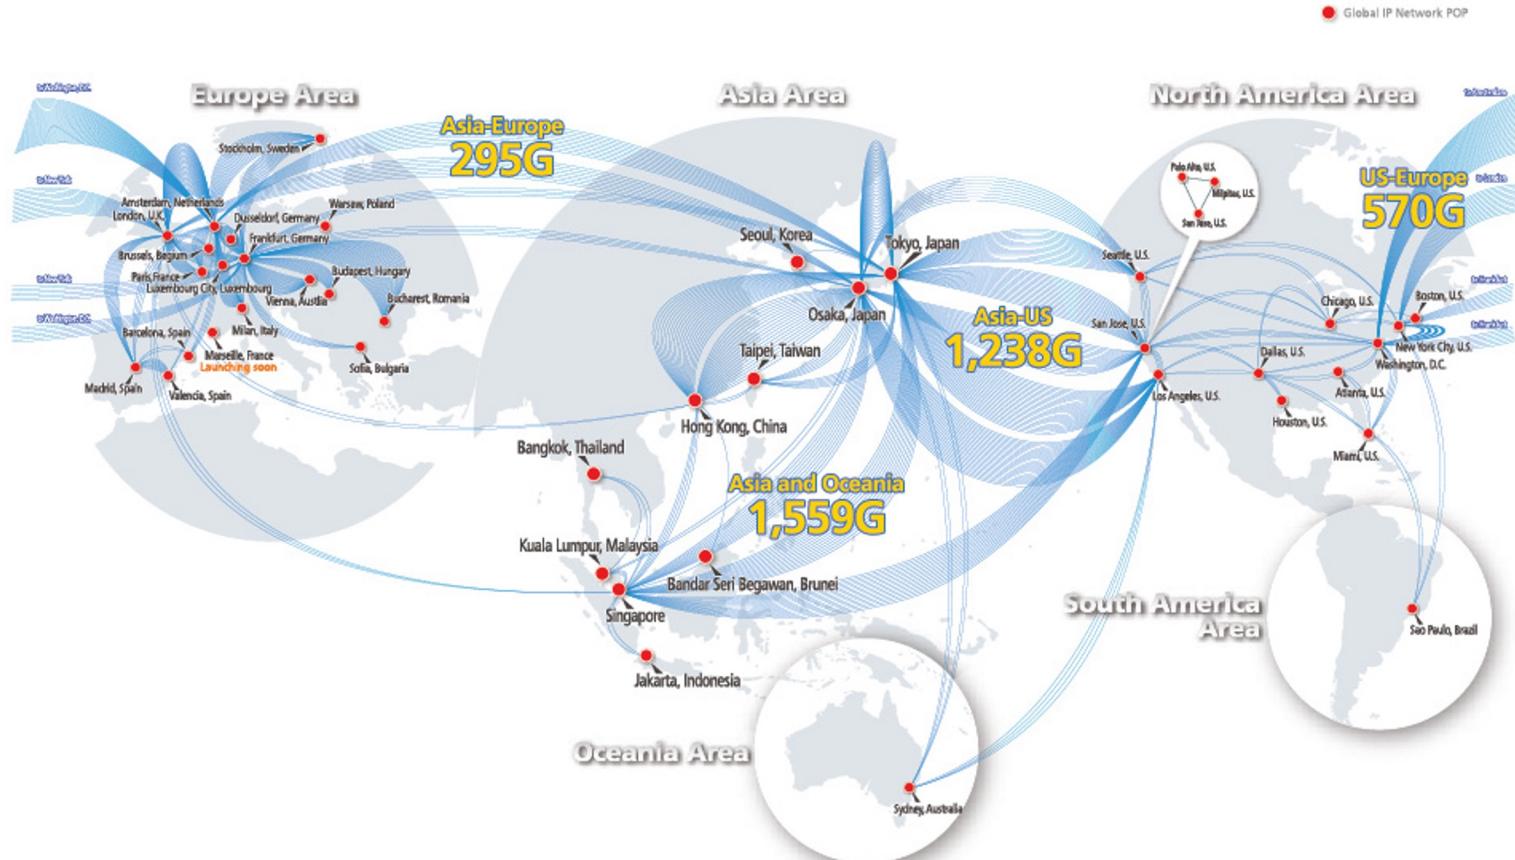
Raspberry Pi  
Beauwolf Cluster



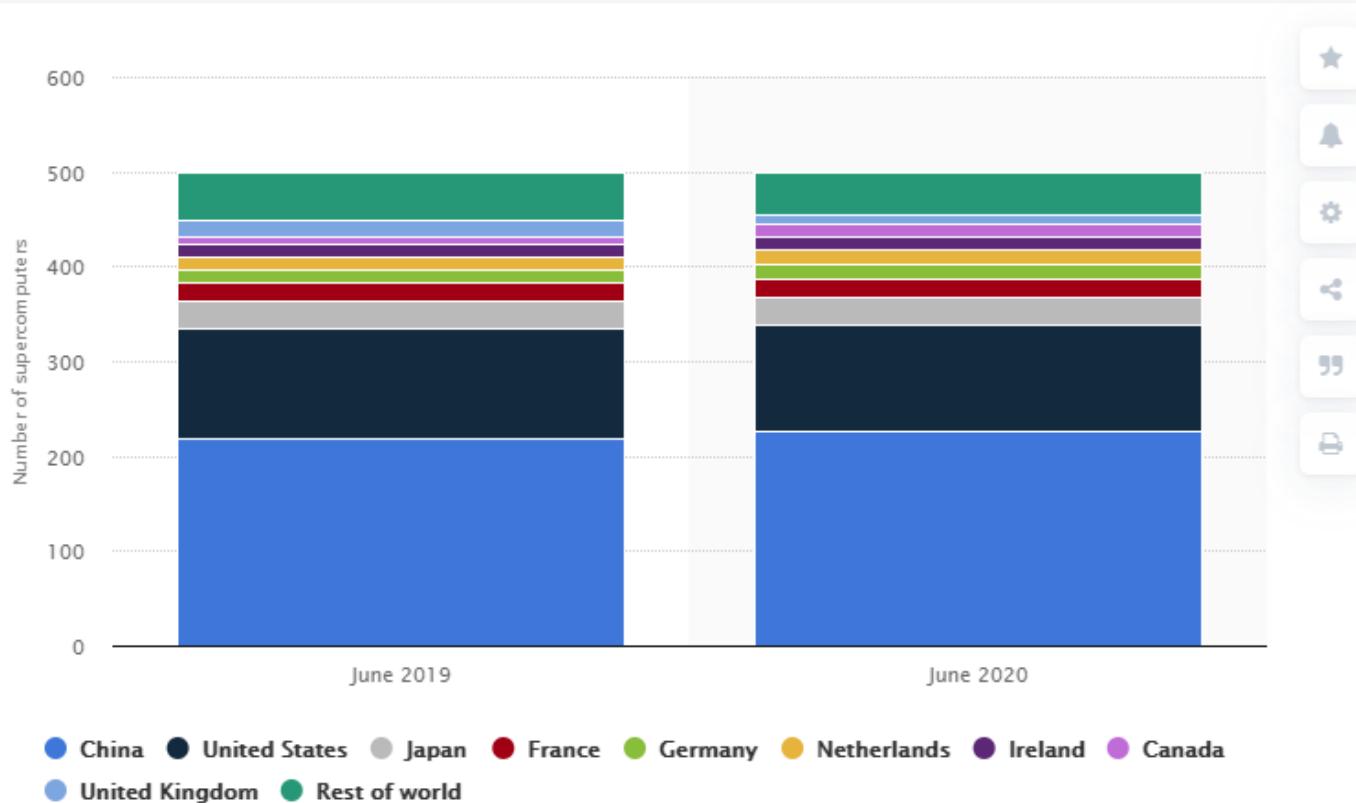
# Grid Computing (*1990 onward*)

- Like cluster, but heterogeneous nodes
- Large physical distances
- Utility computing idea
- Needed high bandwidth connections (Internet)
- Good for processing work, not IO
- [SETI@Home](#), [BOINC](#)

# Building Blocks (Network)

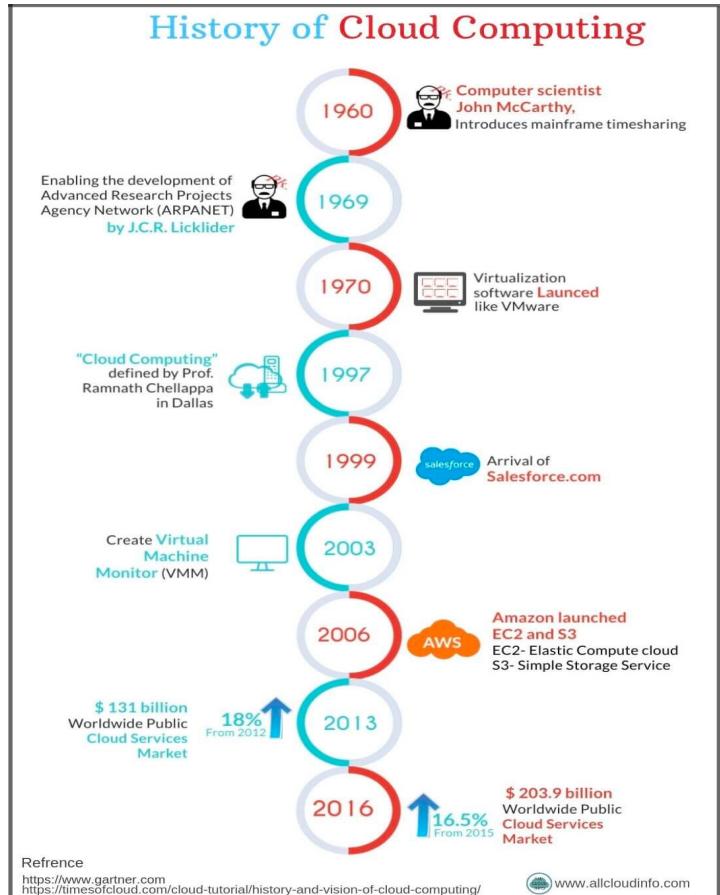


# Building Blocks (*Computing Power*)



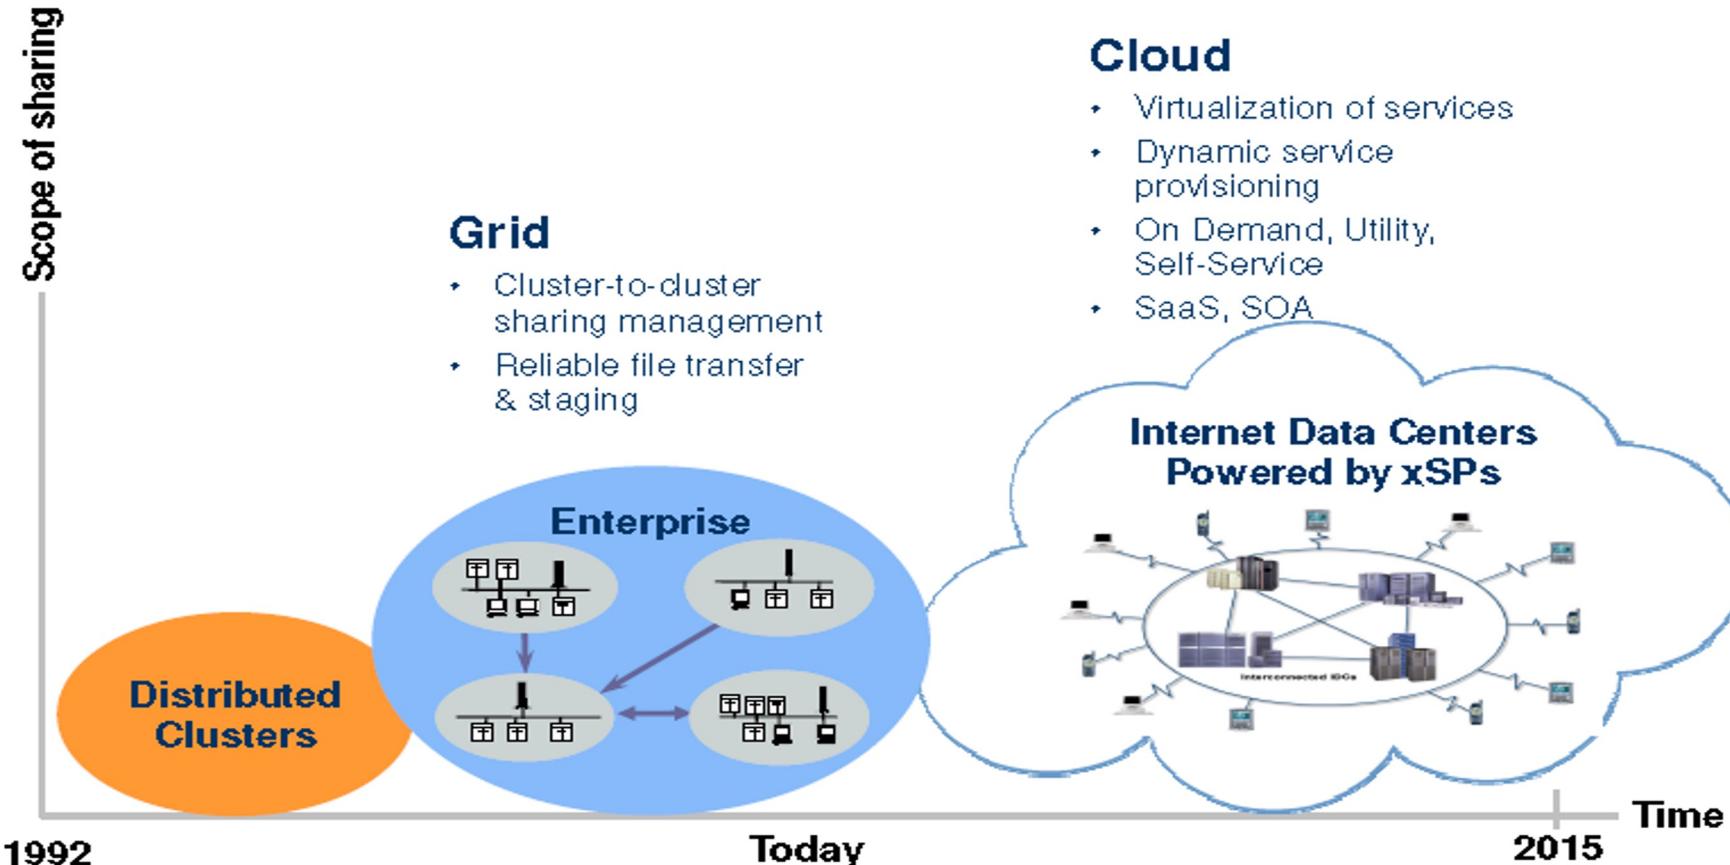
# Building Blocks (*Data Storage*)

1. The Citadel – Tahoe Reno, Nevada
2. Range Range International Information Group – Langfang, China
3. Switch Switch SuperNAP – Las Vegas, Nevada
4. DFT Data Center – Ashburn, Virginia
5. UTAH Data Center – Bluffdale, Utah
6. Microsoft Data Center – West Des Moines, Iowa
7. Lakeside Technology Center – Chicago, Illinois
8. Tulip Data Center – Bangalore, India
9. QTS Metro Data Center – Atlanta, Georgia
10. Next Generation Data Europe – Wales, UK



# But what about latency!?!

- A semi full of hard drives has **HUGE** bandwidth!
- We get sold with bandwidth, but low latency is what we're after.
- But the speed of light is fixed!
  - Significant drop in sales for > 500 ms response
  - Google goal is < 200ms
  - NY to LA is 74ms RTT
- Grids may be good for some workloads, but **HORRIBLE** for others (IO and communication)
- Mainframe->Cluster->Grid, latency between nodes increases.
- Data location matters, and how your app uses data.



# Cloud Computing

- Next evolution after grid computing
- Infinite capacity
- Resilient to failures
- Always on
- Built using commodity machines
- Pay-per-use (Utility vision)

# Key Technologies that enabled Cloud

The implementation and wider use of distributed systems theory to build distributed databases and file systems allowed cloud computing to take off.

- Huge datasets (multiple petabytes)
- Vector clocks ([Lamport timestamps](#))
- Paxos - Protocol to arrive at consensus in a network of unreliable processors ([Wiki](#))

# Virtualization

- Abstract core computing elements away
  - Processor
  - Storage
  - Networking
- Hardware virtualization (VMware, VirtualBox, XEN, EC2, etc....)
  - Most performance issues solved
- Process virtualization (Google AppEngine, Azure, Java)

# Web 2.0

- Web 1.0? - Static pages
- Web 2.0 is:
  - “Web as platform” - John Battelle and Tim O'Reilly
  - Interactivity & flexibility - Allow users to change a site's content!
  - Asynchronous JavaScript and XML (AJAX)
  - Web Services
- Not a ‘next version’ of the web, but a way of using HTTP/HTML
- Examples: Google Docs, Facebook, YouTube, Wikipedia

# Web 2.0. Future of Internet Web 3.0

1900-early 2000s  
Static platform

**Web 1.0**



Current Centralized  
Social-media driven

**Web 2.0**

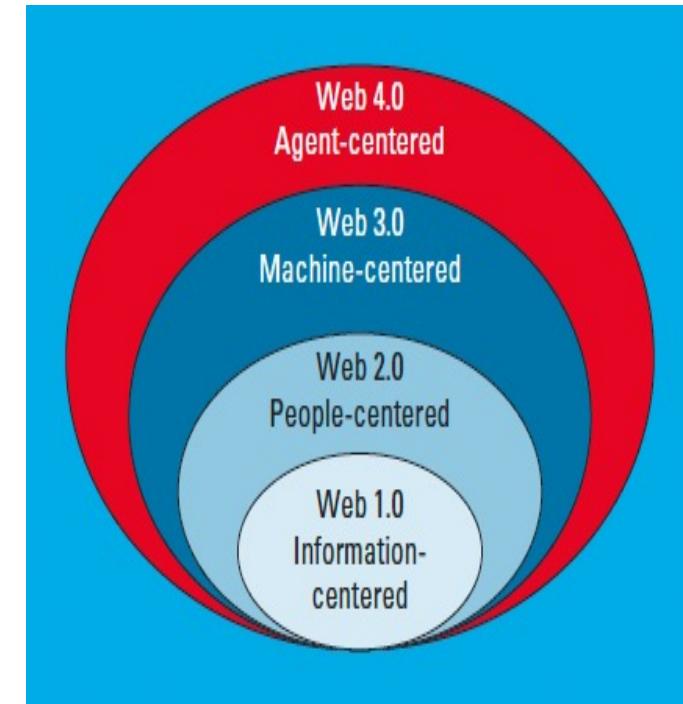


Decentralized  
Autonomous p2p

**Web 3.0**



Crawl	Walk	Run
Mostly Read-Only	Wildly Read-Write	Portable & Personal
Company Focus	Community Focus	Individual Focus
Home Pages	Blogs / Wikis	Lifestreams / Waves
Owning Content	Sharing Content	Consolidating Content
Web Forms	Web Applications	Smart Applications
Directories	Tagging	User Behavior
Page Views	Cost Per Click	User Engagement
Banner Advertising	Interactive Advertising	Behavioral Advertising
Britannica Online	Wikipedia	The Semantic Web
HTML / Portals	XML / RSS	RDF / RDFS / OWL



# Service-Oriented Computing (SOA)

- A component that can perform any function (web services).
  - Loosely coupled
  - Reusable
  - Programming language independent
  - Location transparent
- By layering services we can build a service-oriented architecture

## Important attributes:

- Quality of Service (QoS)
  - Service attributes, response times, security, uptime, etc...
- New software delivery model
  - Can sell components, not entire programs
  - Access through the internet
  - HTTP
  - Web Service Description Language (WSDL)
  - Simple Object Access Protocol (SOAP)

# Utility-Oriented Computing

- Storage, compute, applications, infrastructure, all on a pay-per-use basis.

## Old idea..

- job queueing systems, OS time-slicing, all developed in mainframes to charge for use.
- Buying services/products online common now.

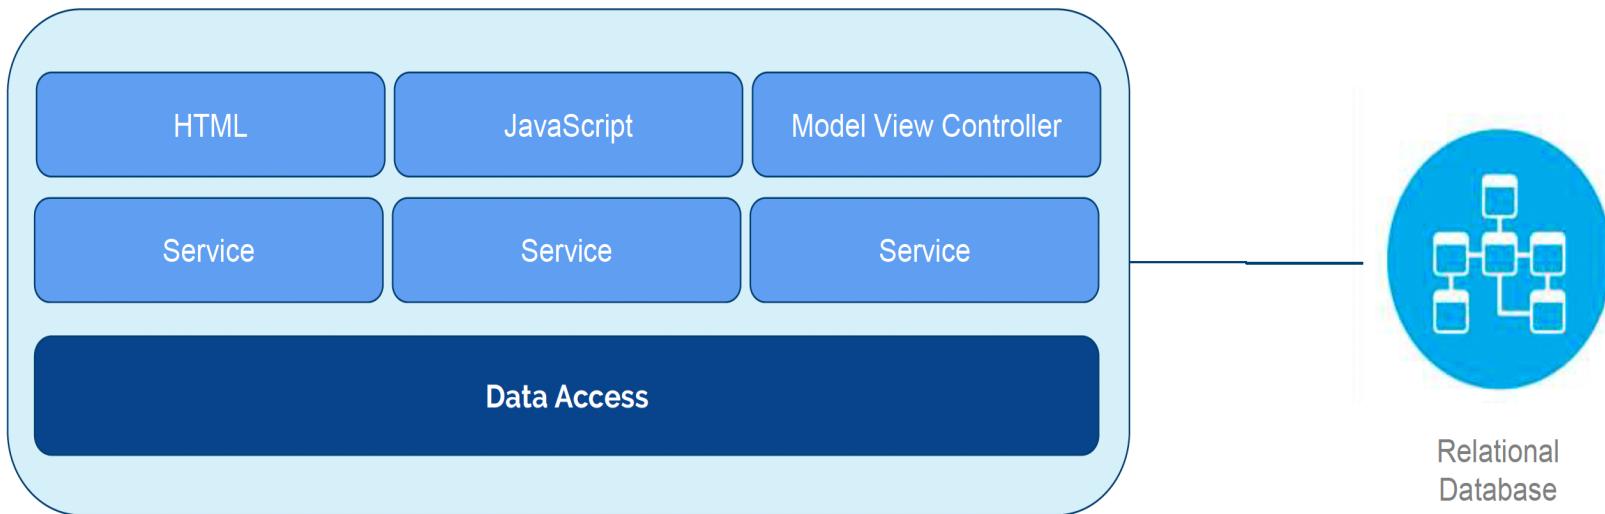
# Current Platforms

- Amazon Web Services (AWS) - Current leader in IaaS: Elastic Compute Cloud (EC2), Simple Storage Service (S3), many many others.....
- Google Cloud Platform (GCP) - PaaS - Python, Java, Go
- MS Azure - IaaS & PaaS
- Hadoop - Open Source framework for MapReduce
- Force.com & Salesforce.com
- Serverless Architecture

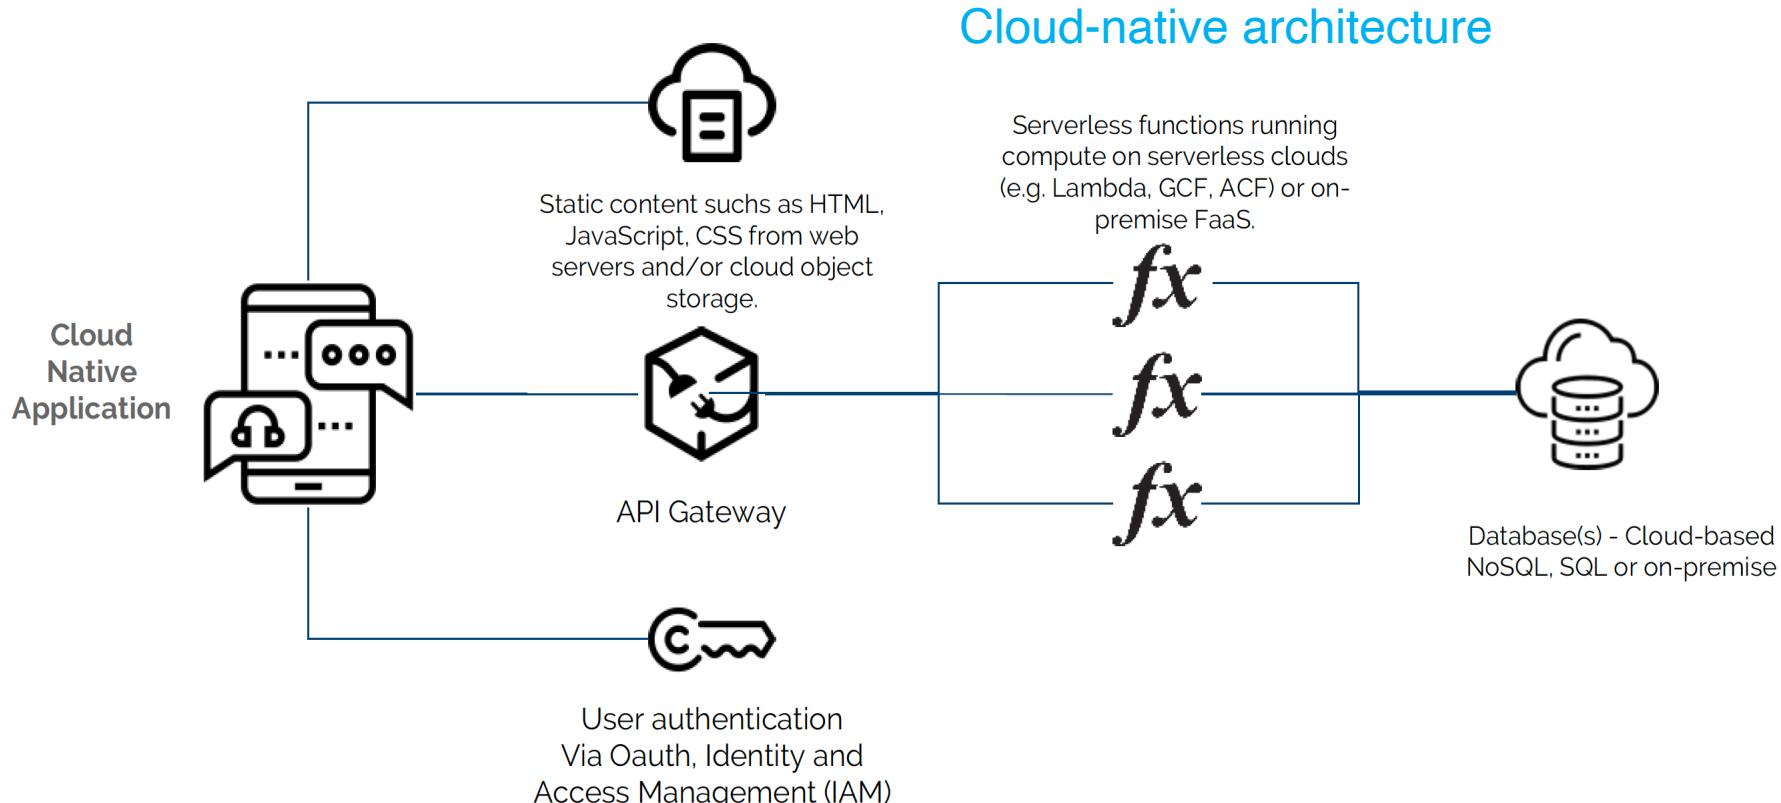
# Serverless Architectures

- Microservice architecture (decompose service into microservice)
- CloudFoundry
- Amazon Lambda
- BlueMix
- Predix
- Microsoft Services Framework

# Monolithic Architectures



# Serverless Architectures



# Serverless is the future (sort of)

Continuous integration, continuous deployment is going to change and that's where things like microservices, containers are playing a massive role in the outer loop, but one of the things that I think is going to completely change how we think about logic is **serverless**.

Serverless computation will fundamentally, not only change the economics of what is backing computing, but it's going to be the core of the future of distributed computing. So this application paradigm shift of intelligent cloud and intelligent edge is going to be pervasively changing everything

....  
**Satya Nadella**

Keynote - Microsoft Build 2017 conference



# Serverless is the future (sort of)

The torrid pace of adoption and innovation in the **serverless** (Lambda) space has totally blown us away,

In particular, Lambda, AWS's main serverless service, has “**grown like crazy**,” with hundreds of thousands of active customers using it in the last 30 days. That’s 300%-plus year-over-year growth.

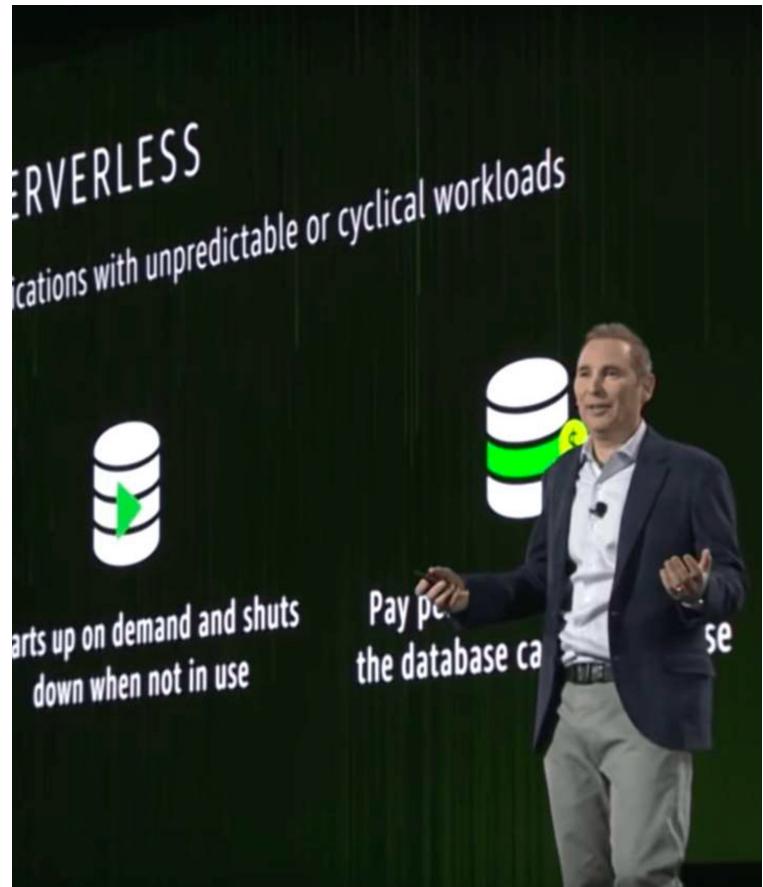
“When we launched Lambda, the first serverless compute service, it was a **watershed moment**,”

...large generations of customers will skip instances and containers and go right to serverless -- in fact, if Amazon.com were starting today, it would go serverless.

Andy Jassy

Amazon Web Services Inc. Chief Executive

Re:Invent 2017



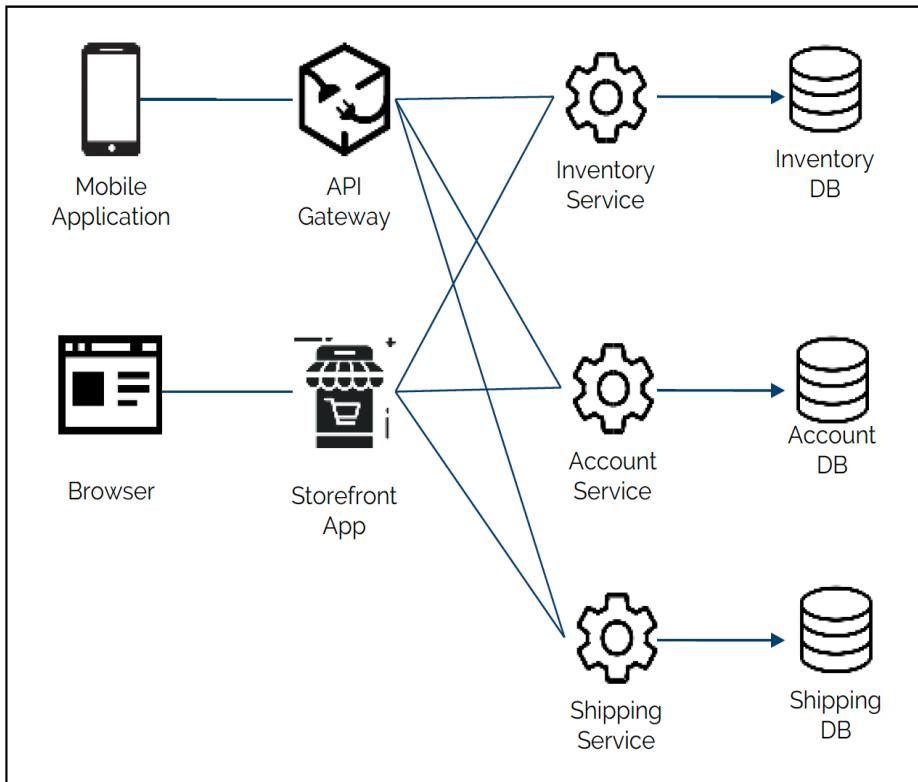
# Where serverless is a good idea?

- ✓ Abstraction of backend infrastructure completely
- ✓ Execution environment for single purpose functions
- ✓ Hosted in public cloud or on-premise
- ✓ Serverless functions have a runtime in a stateless container
  - E.g. Node.js, JavaScript, Go, Python, Java
- ✓ Event-driven startup triggers/instant scale out or in.
- ✓ Micro-billing instead of per-hour/month billing.
- ✓ Criteria for using serverless
  - Asynchronous, concurrent
  - Infrequent or has sporadic demand
  - Stateless, ephemeral
  - Highly dynamic in terms of changing business requirements

# Microservices

- ✓ Single purposes functions that deliver accomplish a single task
- ✓ Can run on a server, container or in a serverless infrastructure
- ✓ Microservices are typically combined to deliver a cloud native application via presentation layer.

## Microservices in a Storefront Application



# Some Security Implications

- ✓ Provisioning in software - automated, online, hands-off
- ✓ Full-system imaging & snapshotting, assists forensics investigation
- ✓ Partitioning - Better isolation between applications than virtual web hosting servers
- ✓ Software networking - pros / cons
- ✓ Insecure Deployment Configuration
- ✓ Broken Authentication
- ✓ Overprivileged Function Permissions and Roles
- ✓ Inadequate Function Monitoring and Logging
- ✓ Insecure Third-Party Dependencies
- ✓ Insecure Application Secrets Storage
- ✓ Denial-of-Service and Financial Resource Exhaustion
- ✓ Business Logic Manipulation, Function Event-Data Injection
- ✓ Improper Exception Handling and Verbose Error Messages
- ✓ Legacy/Unused Functions & Cloud Resources
- ✓ Cross-Execution Data Persistence

# Web 2.0

- Web 1.0? - Static pages
- Web 2.0 is:
  - “Web as platform” - John Battelle and Tim O'Reilly
  - Interactivity & flexibility - Allow users to change a site's content!
  - Asynchronous JavaScript and XML (AJAX)
  - Web Services
- Not a ‘next version’ of the web, but a way of using HTTP/HTML
- Examples: Google Docs, Facebook, YouTube, Wikipedia

# Web 2.0. Future of Internet Web 3.0

1900-early 2000s  
Static platform

Current Centralized  
Social-media driven

Decentralized  
Autonomous p2p

Web 1.0



Web 2.0



Web 3.0

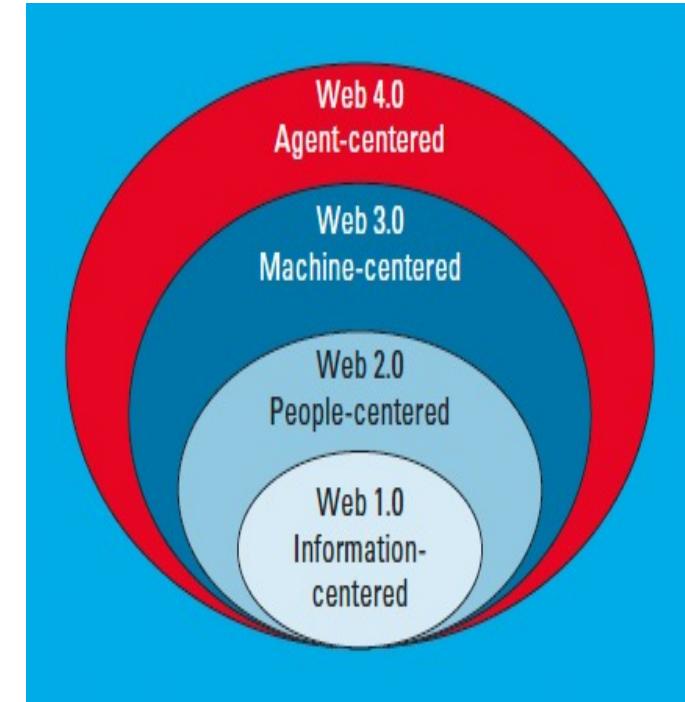


Mostly Read-Only	Wildly Read-Write	Portable & Personal
Company Focus	Community Focus	Individual Focus
Home Pages	Blogs / Wikis	Lifestreams / Waves
Owning Content	Sharing Content	Consolidating Content
Web Forms	Web Applications	Smart Applications
Directories	Tagging	User Behavior
Page Views	Cost Per Click	User Engagement
Banner Advertising	Interactive Advertising	Behavioral Advertising
Britannica Online	Wikipedia	The Semantic Web
HTML / Portals	XML / RSS	RDF / RDFS / OWL

Crawl

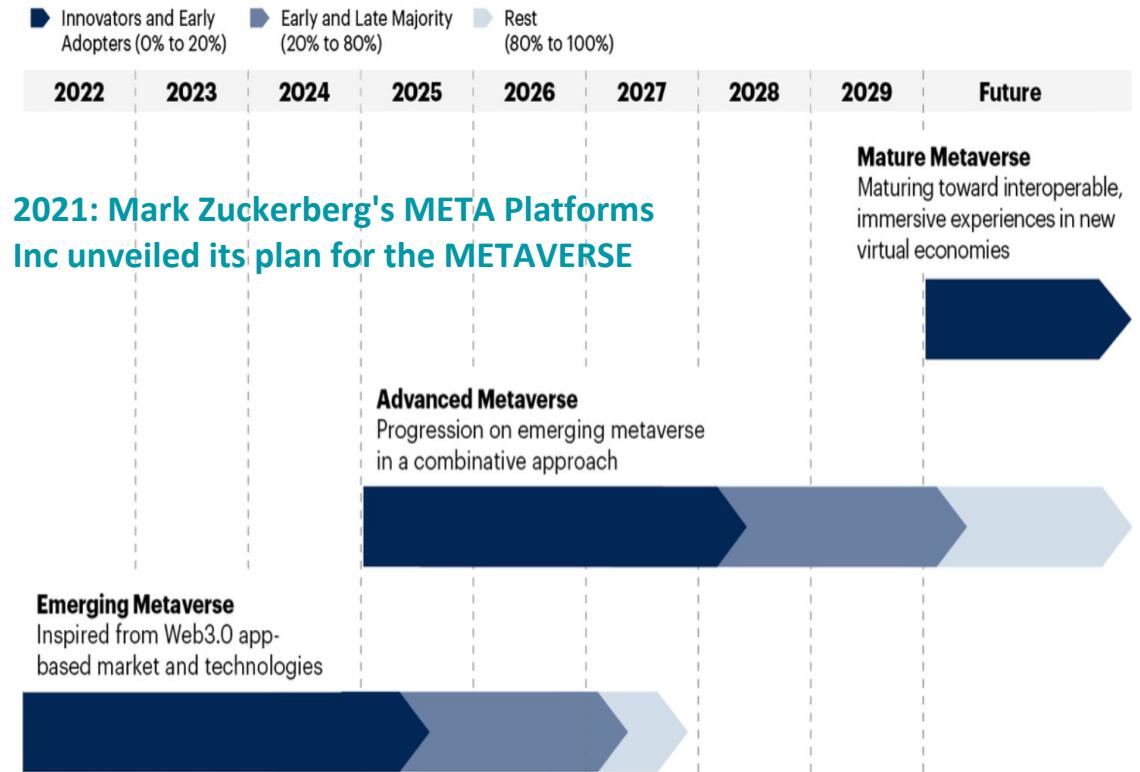
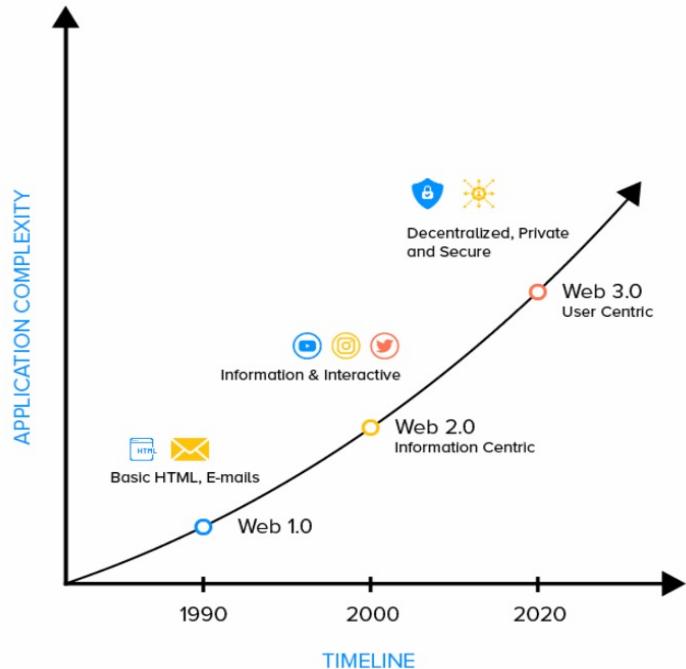
Walk

Run



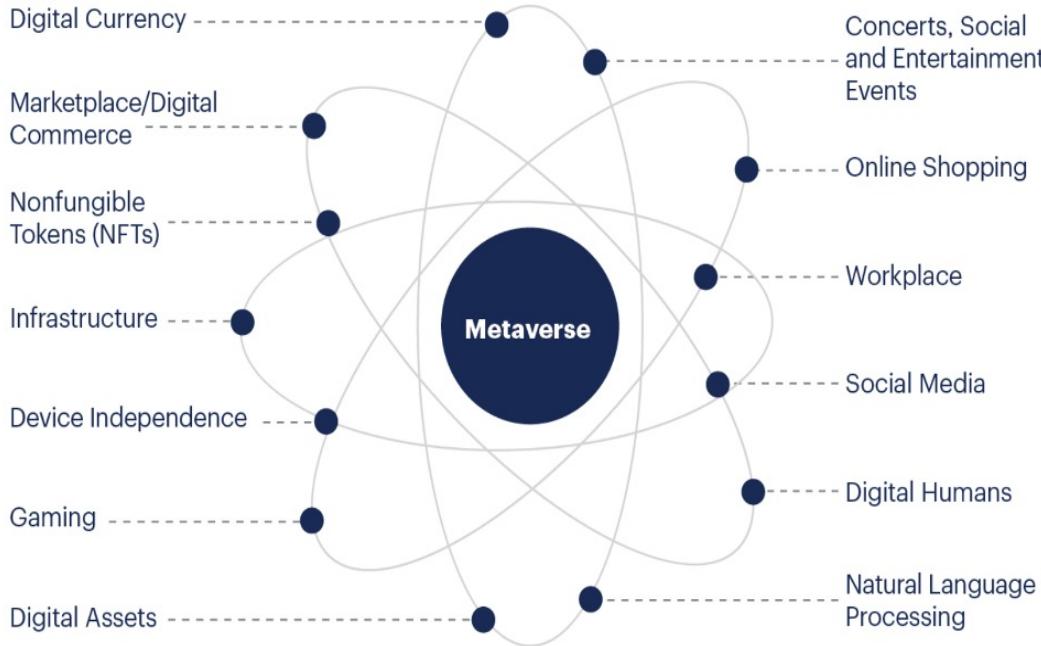
# What is after Cloud Computing?

The technology to create a Metaverse is still in its infancy and a 'Metaverse' does not yet exist. Even what exactly it will be like when it does is largely undefined.



Source: Gartner

# Future of METAVERSE



[gartner.com](http://gartner.com)

Source: Gartner

© 2022 Gartner, Inc. and/or its affiliates. All rights reserved. CTMKT\_1635001

**Gartner®**

- A collective virtual shared space, created by the convergence of virtually enhanced physical and digital reality.
- Metaverse is persistent, providing enhanced immersive experiences, and once matured it will serve as a device-independent ecosystem for further innovation.



[The metaverse explained in 14 minutes | Matthew Ball](#)

# Moving fast in cloud, but is it fast enough?

Build Planes in the Air <https://www.youtube.com/watch?v=JVSiooNywfg>

cowboys herding cats <https://www.youtube.com/watch?v=Pk7yqlTMvp8>

