


# Rashmi Kansakar

# Computing Eras

- Sequential - 1940s+
  - Parallel and Distributed - 1960s+
- 
- But... CS typically teaches sequential
  - Parallel programming is hard!
  - Moore's Law (modern CPUs) now require us into write parallel programs.

# Computing Eras...




# Von Neumann Architecture

Named after the Hungarian mathematician **John von Neumann** who first authored the **general requirements for an electronic computer** in his **1945** papers

Since then, virtually all computers have followed this basic design

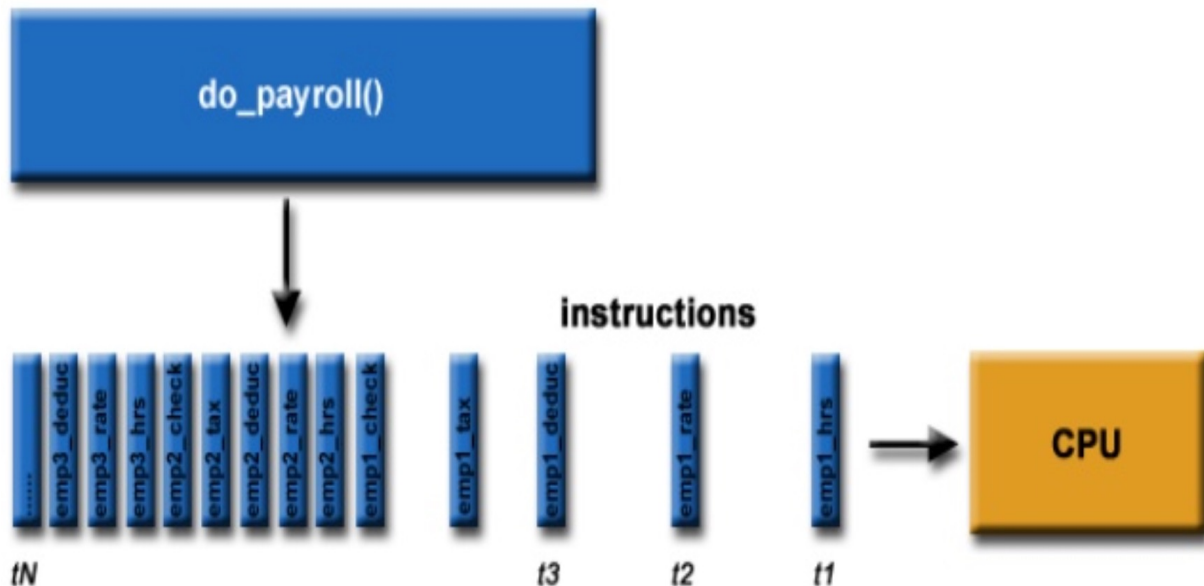
Comprises of four main components:

- **RAM** is used to store both program instructions and data
- **Control unit** fetches instructions or data from memory, Decodes instructions and then *sequentially* coordinates operations to accomplish the programmed task.
- **Arithmetic Unit** performs basic arithmetic operations
- **Input-Output** is the interface to the human operator



# What is Serial Computing?


**Traditionally, software written for serial computation architecture**



- ✓ Run on a single computer having a single Central Processing Unit (CPU)
- ✓ A problem is broken into a discrete series of instructions
- ✓ Instructions are executed one after another
- ✓ Only one instruction may execute at any moment in time

# What is Parallel Computing?

**Parallel computing is the simultaneous use of multiple compute resources to solve a computational problem**



- ✓ Breaking the problem into independent parts
- ✓ Each processing element execute its part of algorithm simultaneously with the others
- ✓ Computational problem solved in less time with multiple compute resources than with a single compute resource
- ✓ Compute resources can be a single computer with multiple processors or Several networked computers

# Parallel Vs. Distributed

- Parallel - tightly coupled system
  - Computation divided among processors sharing **common memory**
  - Homogeneous components. Each processor same type & capacity
  - Defined loosening with [InfiniBand](#) & distributed memory
  
- Distributed – Architecture that allows computation is broken down into units and executed concurrently
  - Parallel a subtype - distributed more general
  - Different nodes, processors, or cores
  - Heterogeneous components
  - E.g GRID computing, Internet computing systems

## Renewed interest...

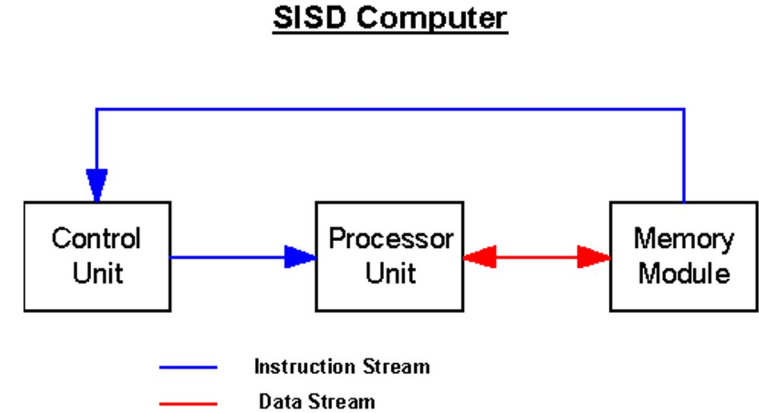

- Larger computation tasks
- CPU have reached physical limits
- Hardware features ([pipelining](#), [superscalar](#), etc) require complex compilers - reached limits
- Vector processing effective, but applicability isolated
- Networking technology mature




# Hardware Architectures

- Single instruction, single data ([SISD](#))
- Single instruction, multiple data ([SIMD](#))
- Multiple instruction, single data ([MISD](#))
- Multiple instruction, multiple data ([MIMD](#))

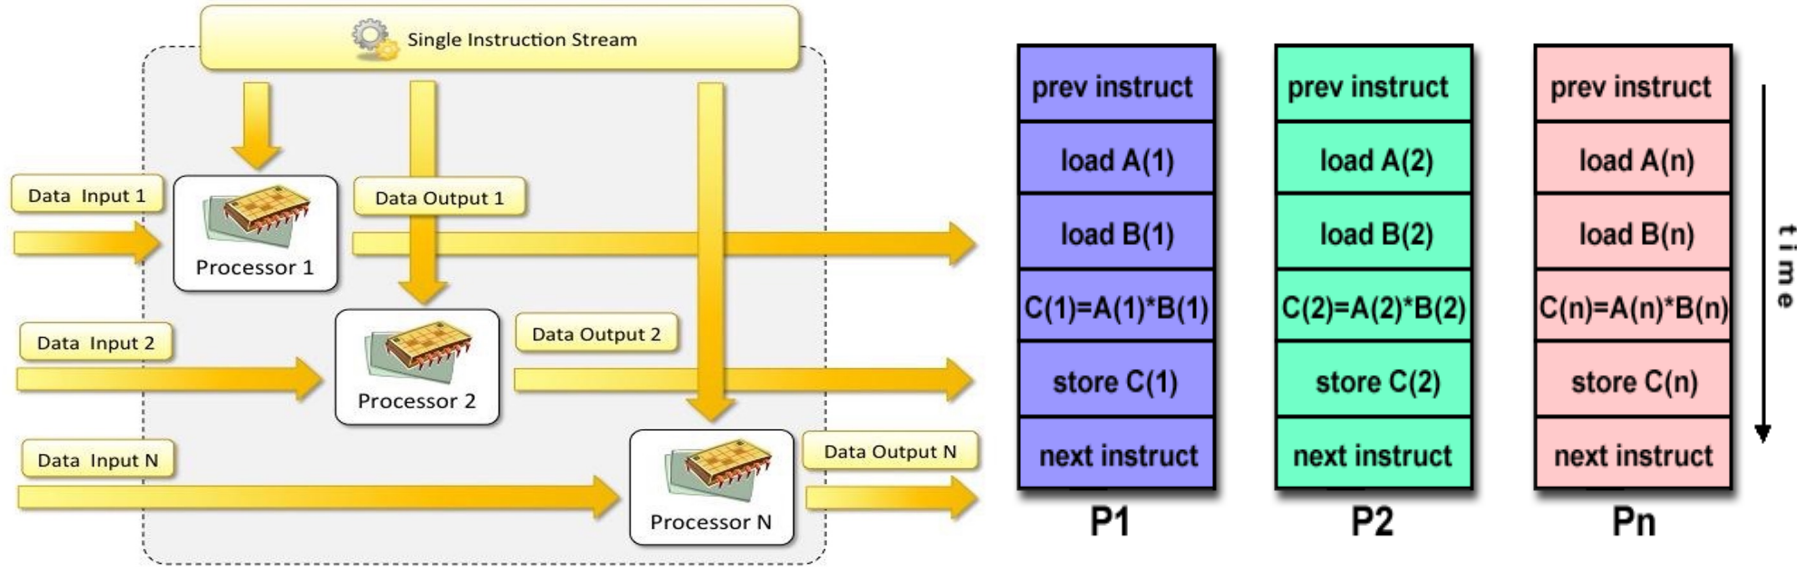
# Single Instruction Single Data (SISD)



- **Sequential computers** (no parallel instruction/data streams)
- **Single Instruction** is being acted by the CPU in 1 clock cycle
- **Single Data** being used as input during any 1 clock cycle
- Older generation mainframes, minicomputers
- 'Normal' computers, modern day PCs, Macs
- CS1, CS2, DS - typically programming

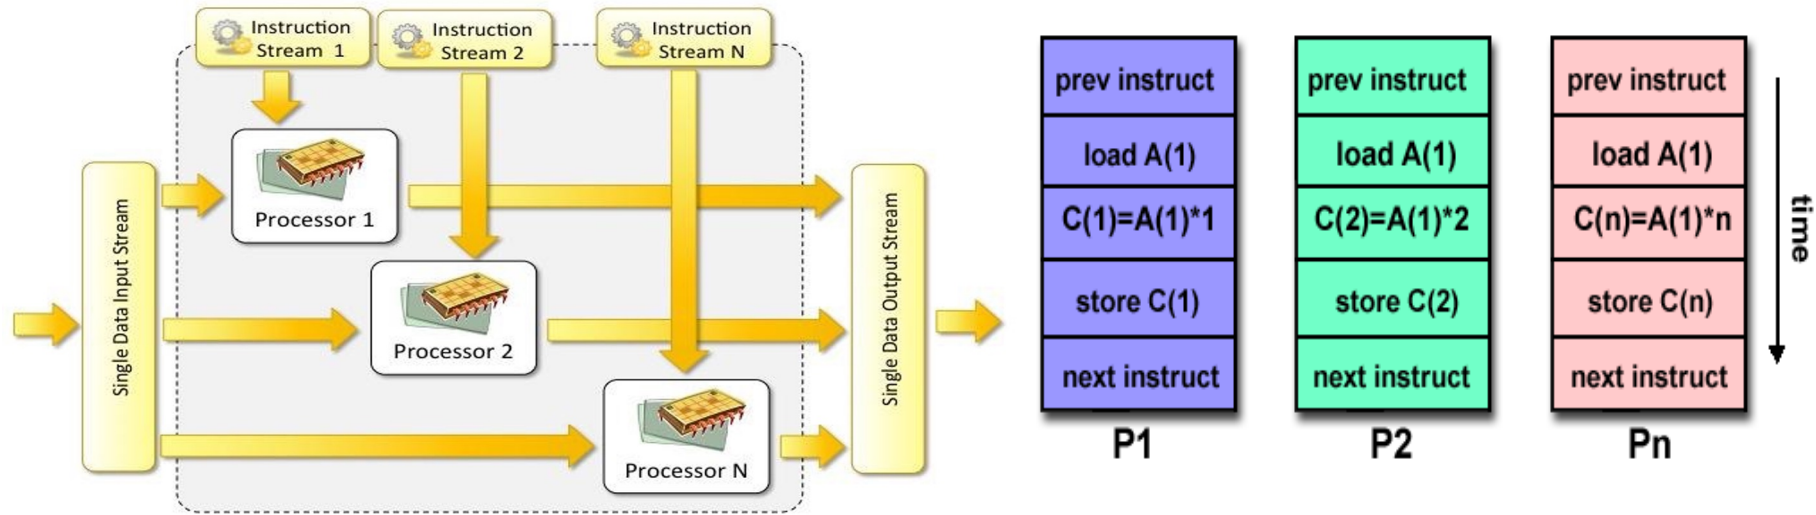


# Single Instruction Multiple Data (SIMD)



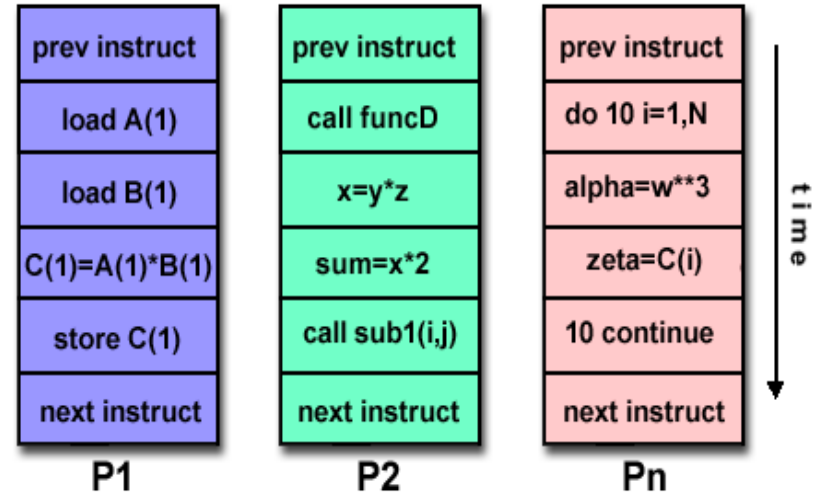
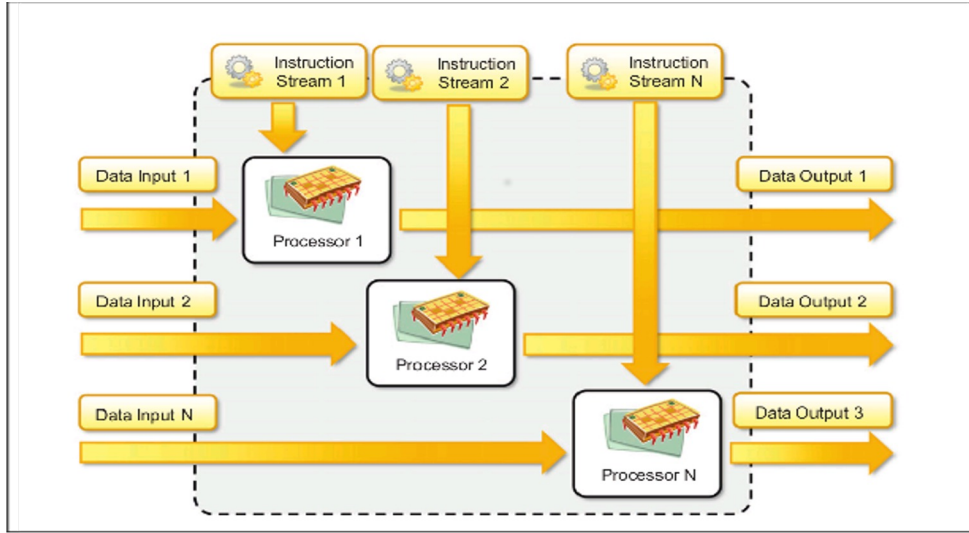
- Multiple data streams against a single instruction stream to perform operations
- Multiple Data: Each processing unit can operate on a different data element
- Vector Processor 1 instruction operate on 1D arrays of data called vectors
- Scientific workloads, vector and matrix operations
- GPUs ([CUDA](#)), Sony PS3 Cell processor ([1](#), [2](#)), [Cray's vector processor](#)

# Multiple Instruction Single Data (MISD)



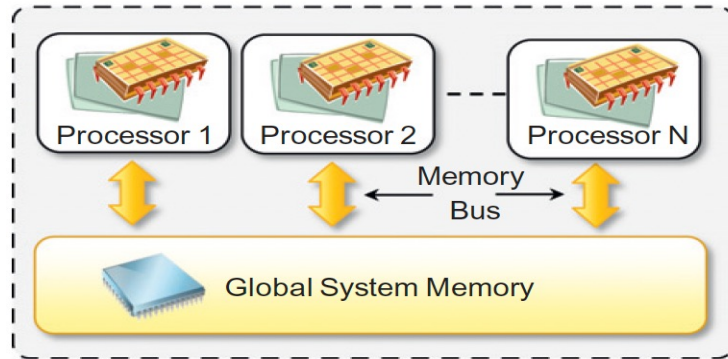
- Each processing unit operates on data independently via separate instruction streams. Example  $y = \sin(x) + \cos(x) + \tan(x)$
- A single data stream is fed into multiple processing units
- No commercial machines exist, though CPU superscalar and pipelining have a similar feel

# Multiple Instruction Multiple Data (MIMD)



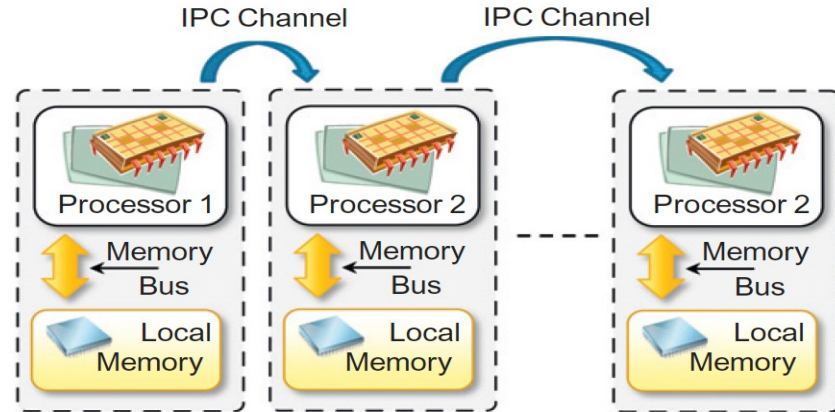
- Multiple autonomous processors simultaneously executing different instructions on different data
- Multiple Instruction: Every processor may be executing a different instruction stream
- Multiple Data: Every processor may be working with a different data stream
- Asynchronous transfers are generally faster than synchronous transfers
- Most supercomputers, networked parallel clusters, grids, clouds, multi-core PCs

# Memory Architectures in MIMD



## Shared Memory

- All PEs are connected to a single global memory and have access to it
- Tightly-coupled multiprocessor systems
- Communication thru shared memory
- Shared MIMD easier to program, less tolerant to failure, and harder to scale
- Failure affect entire entire system

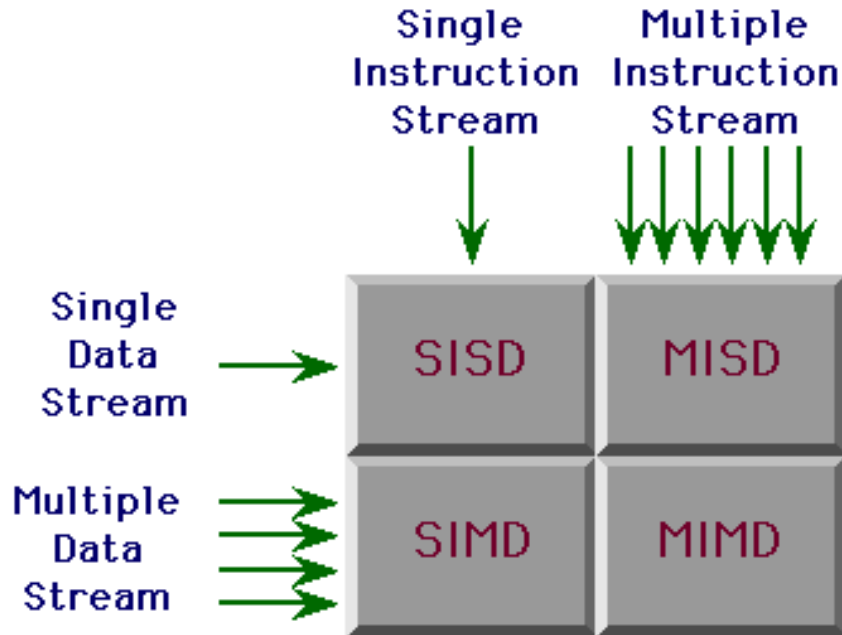


## Distributed-memory

- All PEs have local memory. Loosely-coupled multiprocessor systems
- Cost effectiveness: can use commodity, off-the-shelf processors
- Failure can be isolated. Popular
- Each processor can access its own memory without interference

# Hardware Architectures

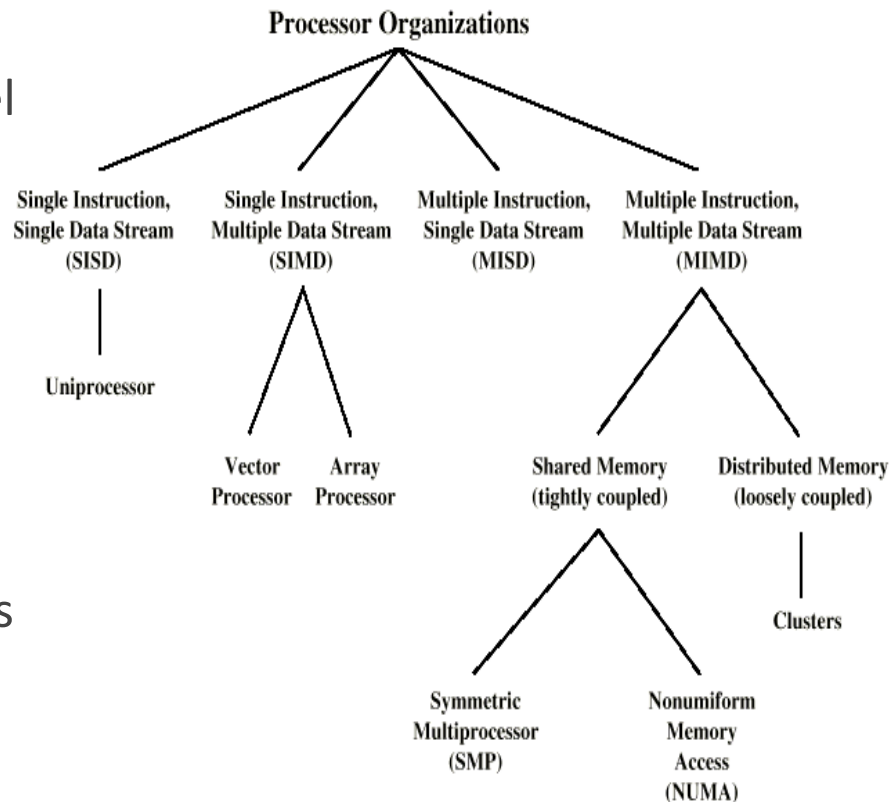
**Flynn's taxonomy:** Based upon the number of concurrent instruction and data streams available in the architecture



- [SISD](#)  
Uniprocessors
- [SIMD](#)  
Vector Processors  
Parallel Processing
- [MISD](#)  
Maybe Pipelined Computers
- [MIMD](#)  
Multi-Computers  
Multi-Processors

# Parallel Processor Architectures

- The computers of today, and tomorrow, have tremendous processing power that require parallel programming to fully utilize.
- There are significant differences between sequential and parallel programming, that can be challenging.
- With early exposure to these differences, students are capable of achieving performance improvements with multicore programming.





# How to Program in Parallel?

- Problem specific!
- Approaches:
  - ***Data parallelism***
    - MapReduce (data based)
  - ***Process parallelism***
    - Game/Cell Processor (code based)
  - ***Farmer-and-worker model***
    - Web serving (Apache) (thread based)

# Level of Parallelism

## ➤ Goal?

- Never have a processor idle!
- ‘Grain size’ important
  - How you break up the problem.


Grain Size	Code Item	Parallelized By
Large (task)	Separate (heavyweight process)	Programmer
Medium (control)	Function or procedure (thread)	Programmer
Fine	Loop or instruction block	Compiler
Very Fine	Instruction	Processor or OS

# Level of Parallelism...


Programmer




Large Level  
(Processes, Tasks)



Medium Level  
(Threads, Functions)



Fine Level  
(Processor,  
Instructions)



Very Fine Level  
(Cores, Pipeline,  
Instructions)

Compiler

Processor &  
OS

# Limit to Parallelism (*yet*)

## Linear speedup not possible

- Doubling # cores doesn't double speed
  - Communication overhead
- General guidelines:
  - Computation speed =  $\sqrt{\text{system cost}}$  or faster a system becomes the more expensive it is to make it faster
  - Speed of parallel computer increases as the  $\log(n)$  of processors ( *$n = \text{number of processors}$* )

# Parallel Overhead

- Sequential program runs in a single processor, and has single line of control
- Parallel programming is making many processors collectively work on a single program
- **Parallel Overhead:**
  - The amount of time required to coordinate parallel tasks, as opposed to doing useful work.
- **Parallel overhead can include factors such as:**
  - Task start-up time
  - Synchronizations
  - Data communications
  - Software overhead imposed by parallel languages, libraries,
  - operating system, etc.
  - Task termination time

# Why Use Parallel Computing ?

## ➤ **Save time and money**

- In theory, throwing more resources at a task will shorten its time to completion
- Parallel computers can be built from cheap, commodity Components

## ➤ **Solve larger problems**

- Many problems are so large and complex that it is impractical or impossible to solve them on a single computer, especially given limited computer memory

## ➤ **Provide concurrency**

- A single compute resource can only do one thing at a time.
- Multiple computing resources can be doing many things simultaneously

## ➤ **Limits to serial computing**

- Increasingly expensive to make a single processor faster
- Using a larger number of moderately fast commodity processors to achieve the same or better performance is less expensive

# The Future

Trends indicated by ever faster networks, distributed systems, and multi-processor computer architectures clearly show that PARALLELISM is the future of




- There has been a greater than 1000x increase in supercomputer performance, with no end currently in sight
- The race is already on for Exascale Computing!

1 exaFlops =  $10^{18}$ ,  
FLOATING POINT OPERATIONS PER SECOND  
FLOP machine will do one "operation" in a second.

# Moore's Law and beyond

**Moore's Law** originated around 1970;  
Simplified version of this **law** states that  
**processor speeds**, or overall processing  
power for computers will **double every two**  
**years**.



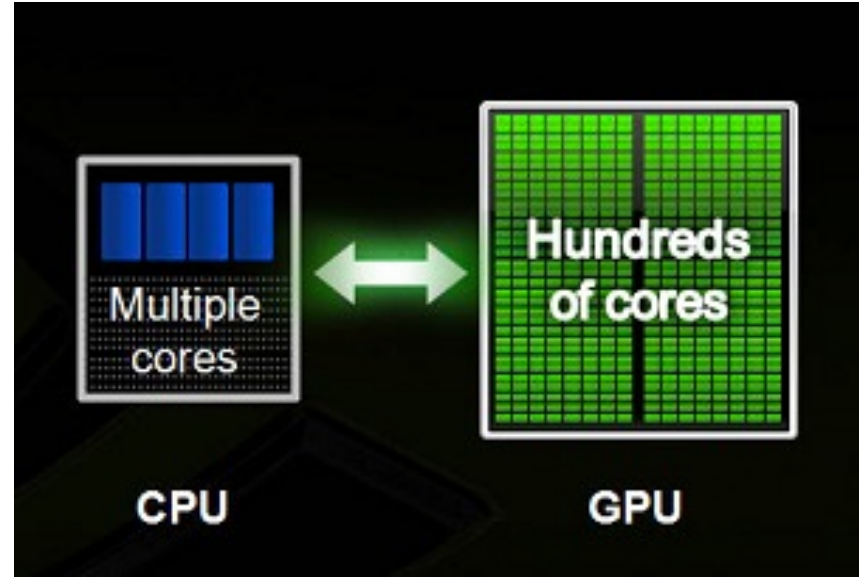

**Moore's Law** is no longer relevant,  
**GPUs** advancing faster pace than CPUs

Driven by  
massive computer demand  
for artificial intelligence  
which is the software glue  
binding services together



# Improving parallel processing CPU vs. GPU

- **CPU** is sometimes called the brains of a computer while a **GPU** acts as a specialized microprocessor.
- **CPU** is good at handling multiple tasks but a **GPU** can handle a few specific tasks very fast.
- **GPU** (graphical processing unit) is a programmable processor designed to quickly render high resolution images and video.
- **CUDA cores** are an Nvidia GPU's equivalent of CPU **cores**. They are optimized for running a large number of calculations simultaneously
- GPU render Billions of triangles per second
- **PARALLEL PROCESSING** in steroid!



**Intel XEON  
PLATINUM 9282**  
CPU Cores: 56 cores  
Retail: \$50,000 +  
Transistors : 8 Billion

**NVIDIA TITAN V**  
CUDA Cores: 5,120  
Tensor Cores: 640  
Retail: \$2,999  
Transistors : 21 Billion

# Beyond GPU to dedicated ML silicon

CPU	GPU
1. CPU stands for Central Processing Unit.	While GPU stands for Graphics Processing Unit.
2. CPU consumes or needs more memory than GPU.	While it consumes or requires less memory than CPU.
3. The speed of CPU is less than GPU's speed.	While GPU is faster than CPU's speed.
4. CPU contain minute powerful cores.	While it contain more weak cores.
5. CPU is suitable for serial instruction processing.	While GPU is not suitable for serial instruction processing.
6. CPU is not suitable for parallel instruction processing.	While GPU is suitable for parallel instruction processing.
7. CPU emphasis on low latency.	While GPU emphasis on high throughput.

# Comparing GPU & CPU

**MythBusters hosts: Adam and Jamie**  
**Paint the Mona Lisa in 80 Milliseconds!**

<https://www.youtube.com/watch?v=WmW6SD-EHVY>

# Why Deep Learning uses GPUs

- Artificial intelligence with PyTorch and CUDA.
- CUDA cores are an Nvidia GPU's equivalent of CPU cores. They are optimized for running large number of calculations simultaneously
- Discuss how CUDA fits in with PyTorch, and more importantly, why we use GPUs in neural network programming.

<https://www.youtube.com/watch?v=6stDhEA0wFQ>