

# Cloud Computing

## (SQL Vs NoSQL)

**Rashmi Kansakar**

1970: **NoSQL = We have no SQL**

1980: **NoSQL = Know SQL**

2000: **NoSQL = No SQL!**

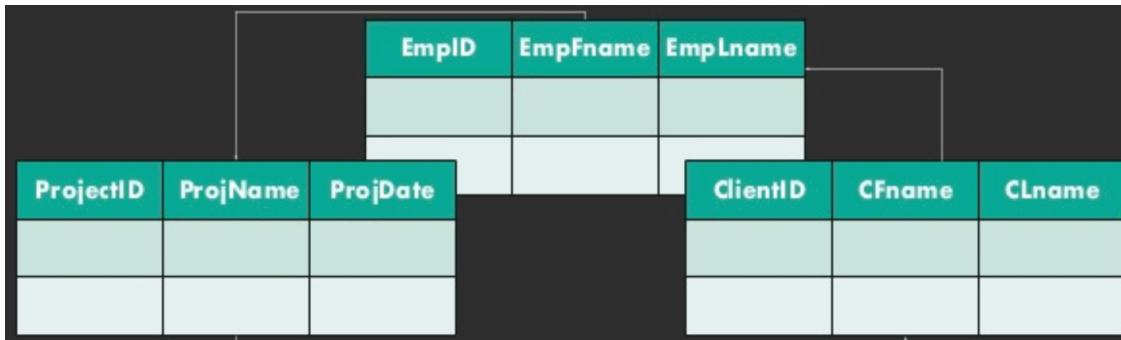
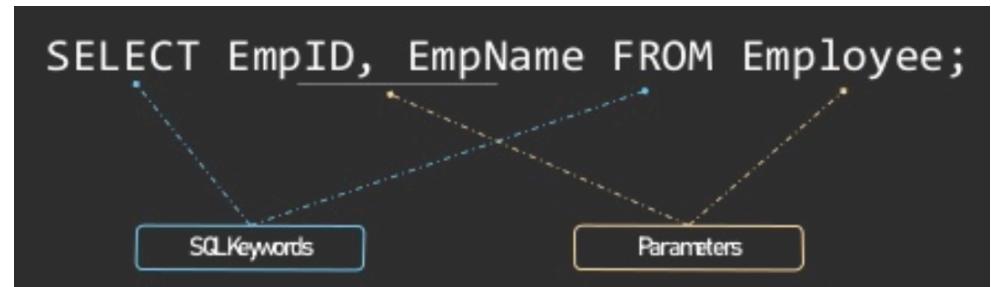
2005: **NoSQL = Not only SQL**

2013: **NoSQL = No, SQL!**

# History

IBM introduced **SEQUEL** in 1970

- Structured English Query Language → SEQUEL
- Later SEQUEL → SQL



The diagram illustrates three tables: Project, Employee, and Client. The Project table has columns `ProjectID`, `ProjName`, and `ProjDate`. The Employee table has columns `EmpID`, `EmpFname`, and `EmpLname`. The Client table has columns `ClientID`, `Cfname`, and `Clname`. The tables are shown as light green grids with black borders, and their column headers are highlighted in teal.

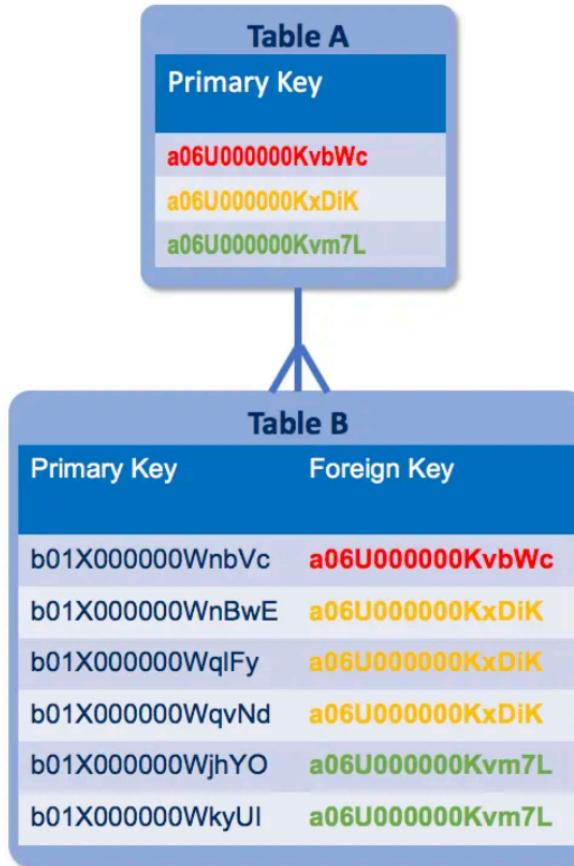
ProjectID	ProjName	ProjDate
EmpID	EmpFname	EmpLname
ClientID	Cfname	Clname

## Relational Database Management System

- Main focus on ACID
  - **Atomicity:** Each transaction is atomic. If one part of it fails, the entire transaction fails (*and is rolled back*).
  - **Consistency:** every transaction is subject to a consistent set of rules (*constraints, triggers, cascades*).
  - **Isolation:** No transaction should interfere with another transaction
  - **Durability:** Once a transaction is committed, it remains committed.

# Why ACID?

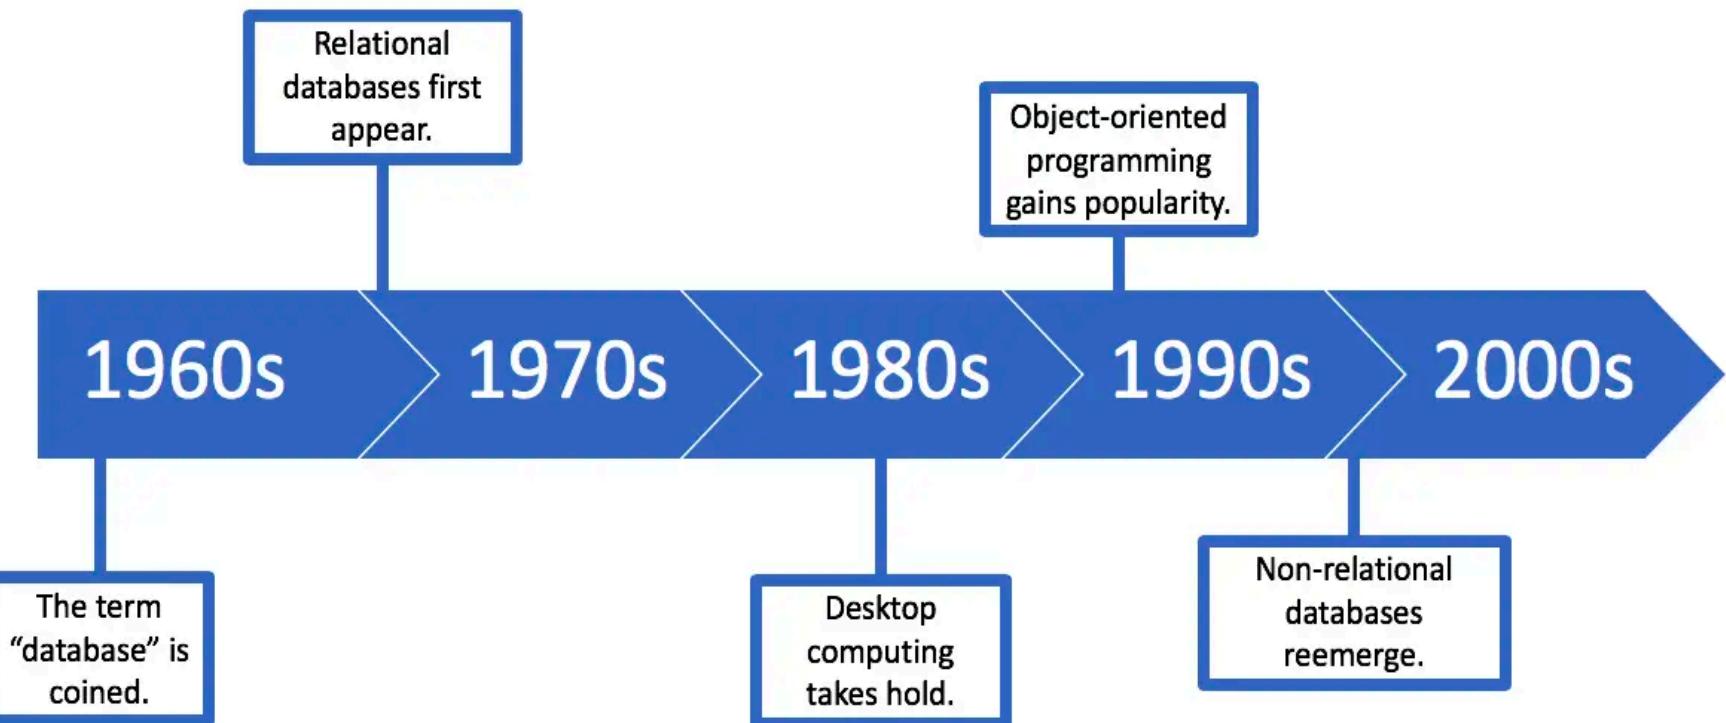
- ACID is important
  - Only when it's important.
  - Banking, Finance, System Safety etc.
  
- ACID adds overhead
  - Features like atomicity, isolation basically force database servers to use sequential evaluation.



## Other Features

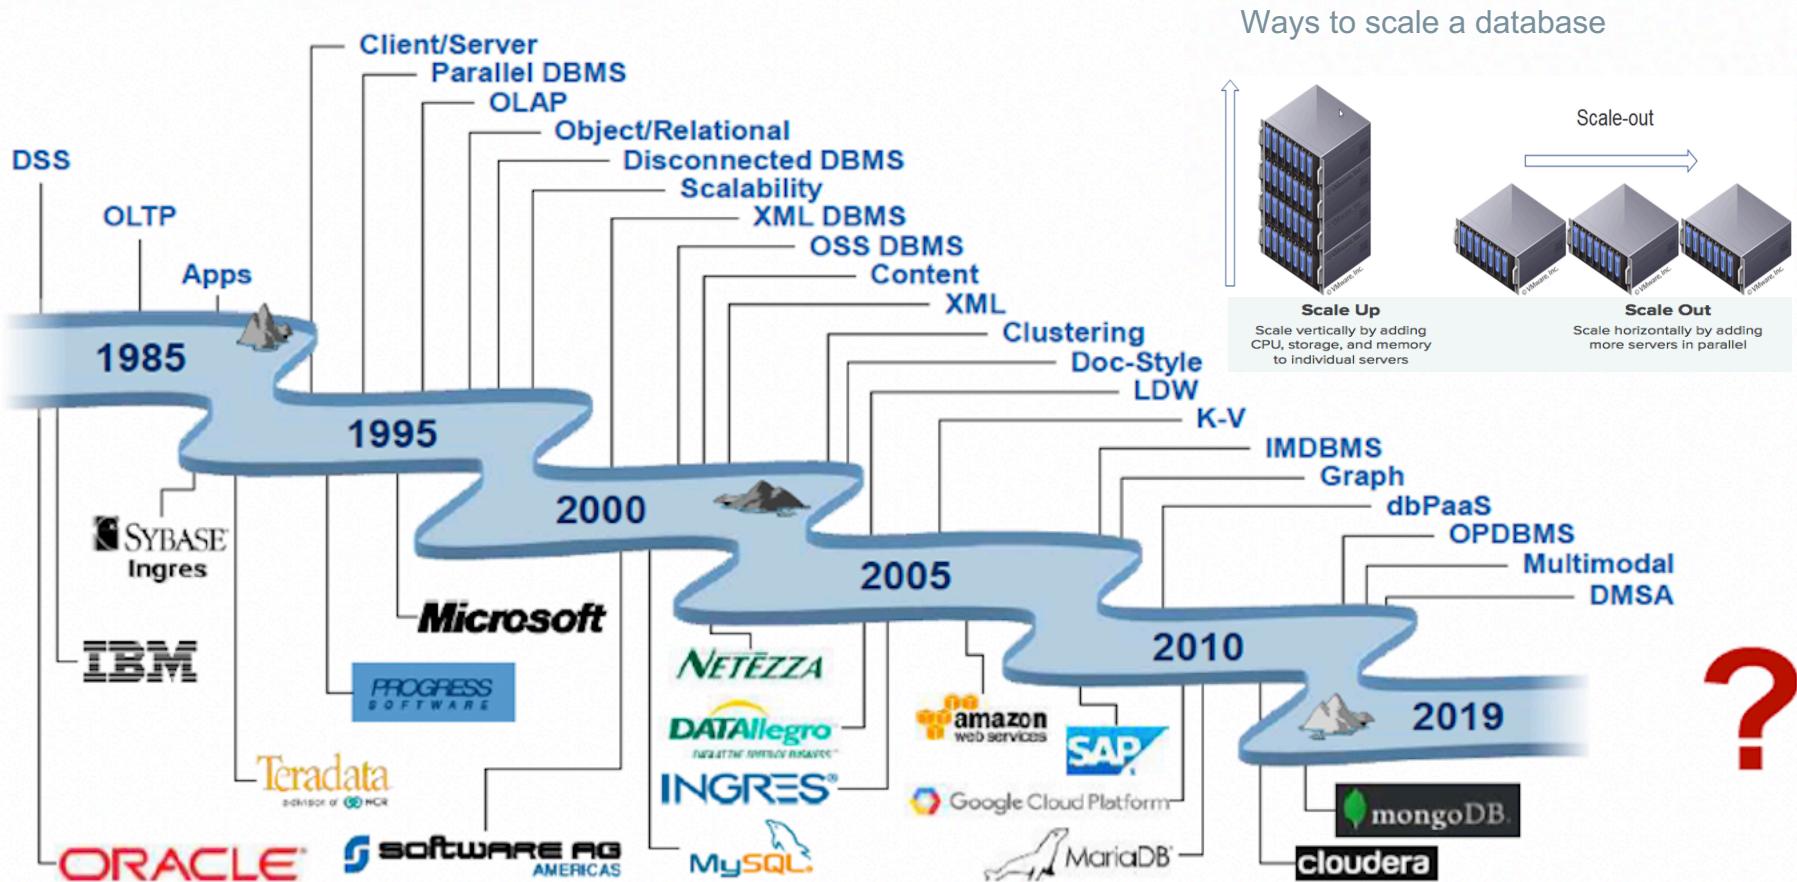
- Typed columns
- Normalized data
- Supports broad, unrestrained queries
- Has its own query optimizer
- Consistent way of accessing data (SQL)

# History of SQL



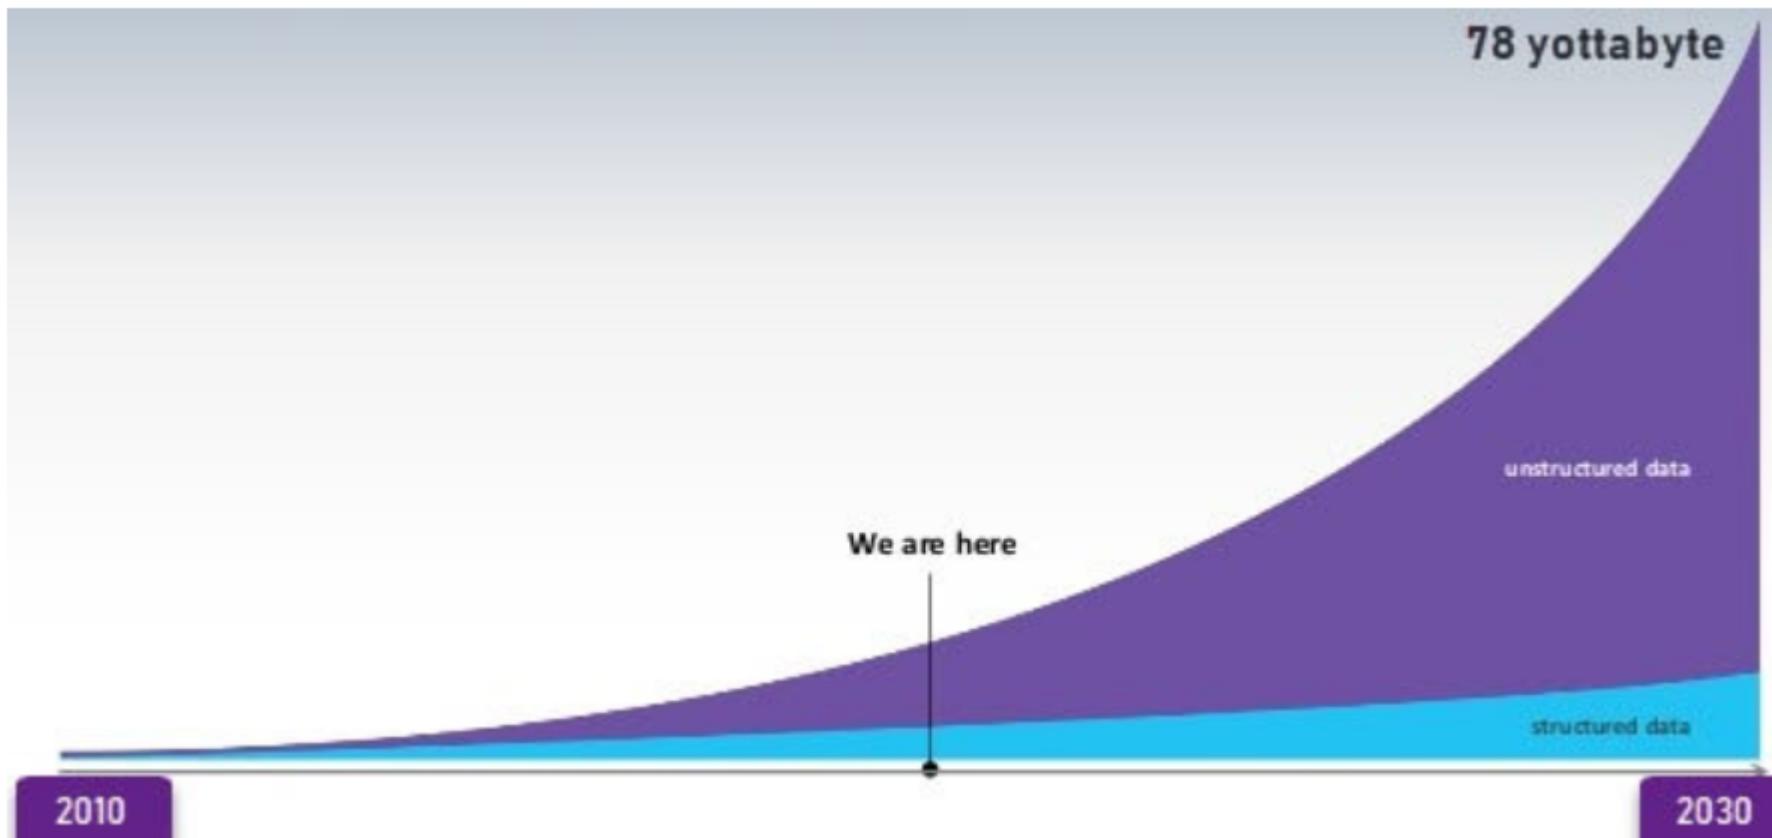
# Evolving Databases for modern application

Pass it on...  
Knowledge is power



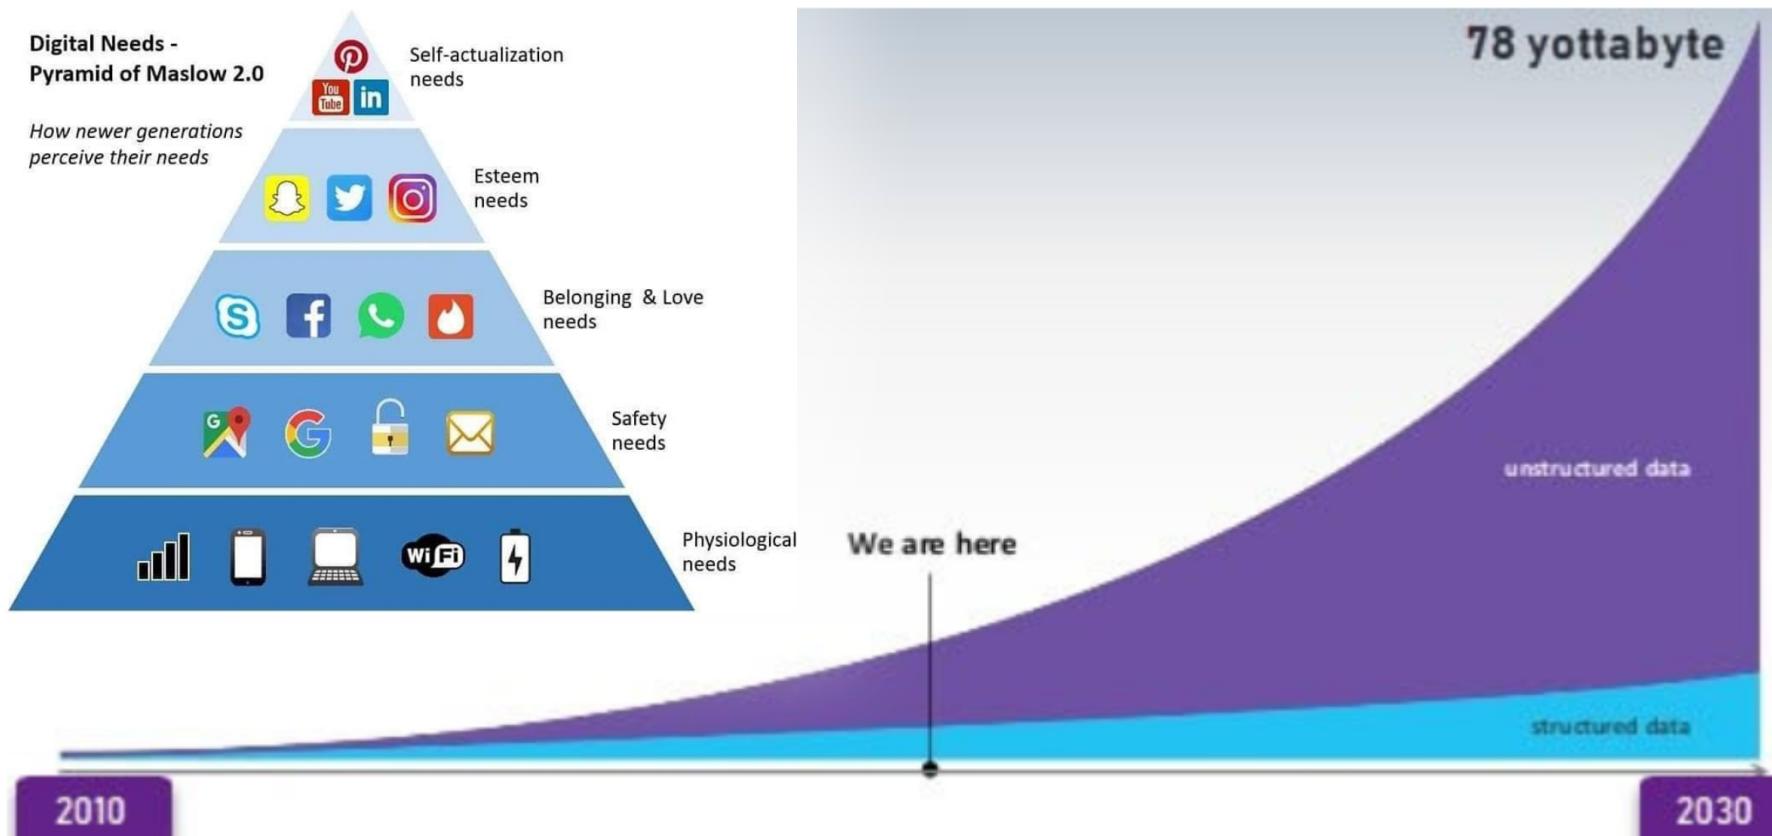
# Data Growth over the years

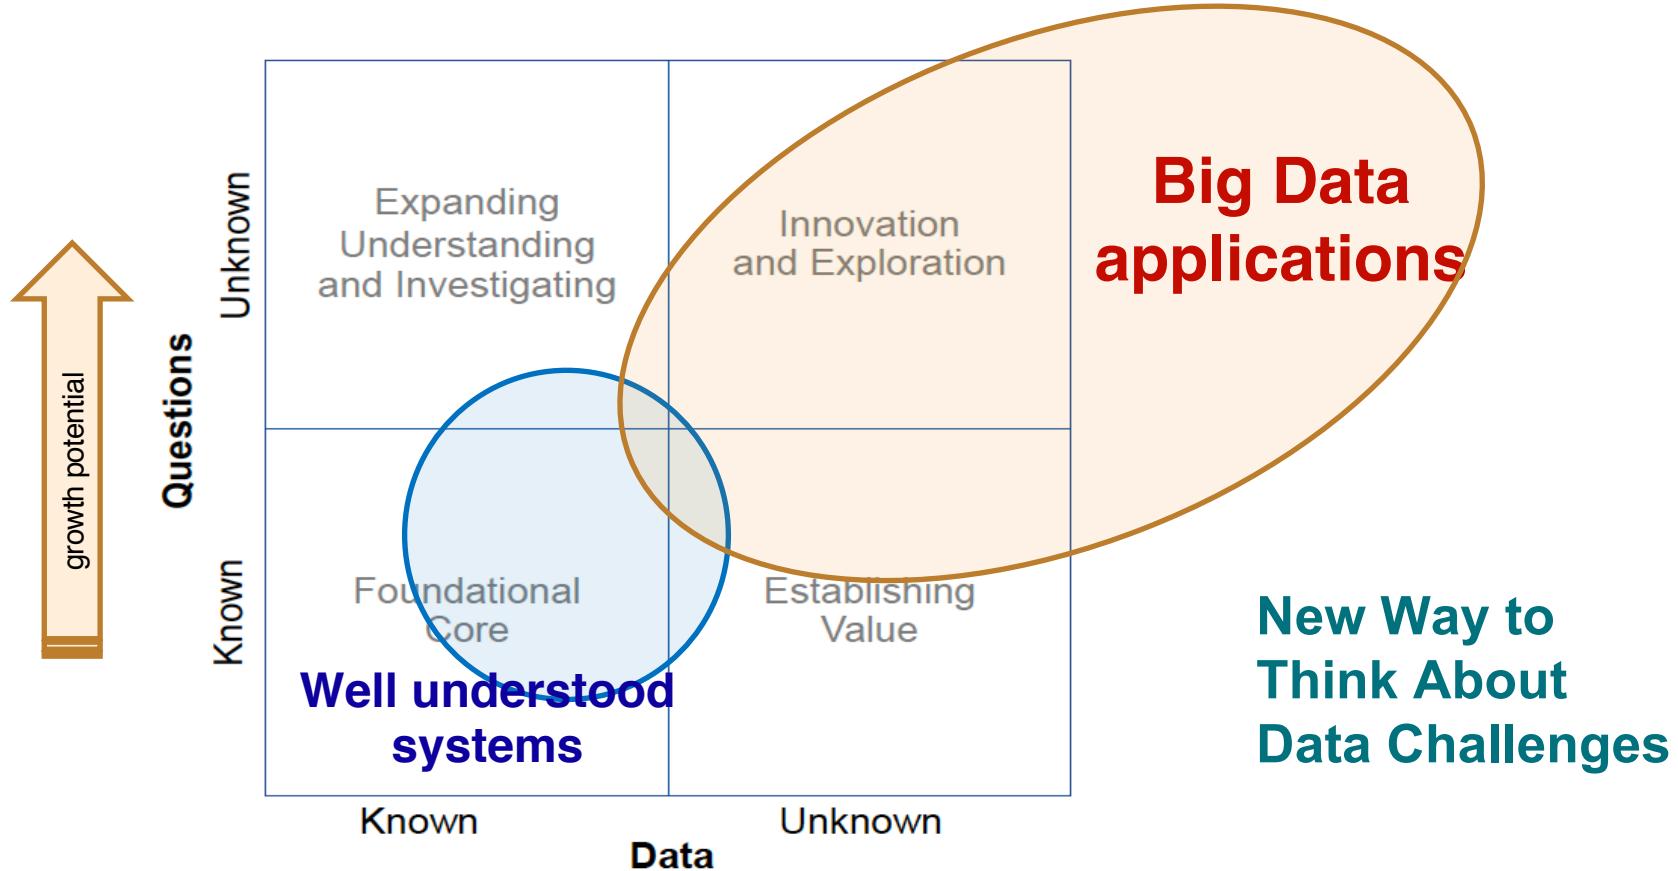
Pass it on...  
Knowledge is power



# Data Growth over the years

Pass it on...  
Knowledge is power



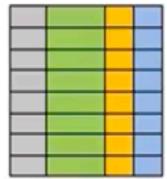


# Types of Databases

Pass it on...  
Knowledge is power

## SQL Databases

### Relational



ORACLE®  
MySQL®

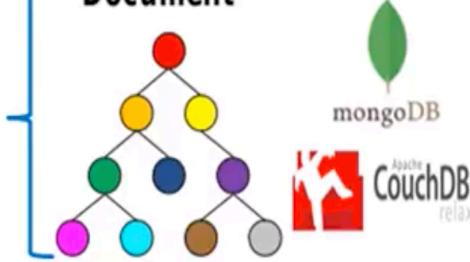
### Analytical (OLAP)



ORACLE®  
SAP

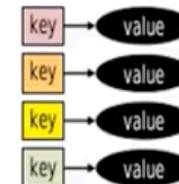
## NoSQL Databases

### Document



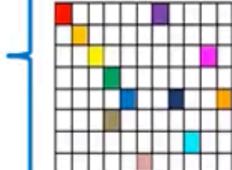
mongoDB  
Apache CouchDB  
relax

### Key-Value



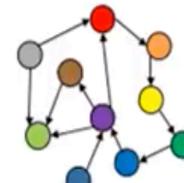
amazon DynamoDB  
redis  
riak

### Column-Family



cassandra  
APACHE HBASE

### Graph



neo4j

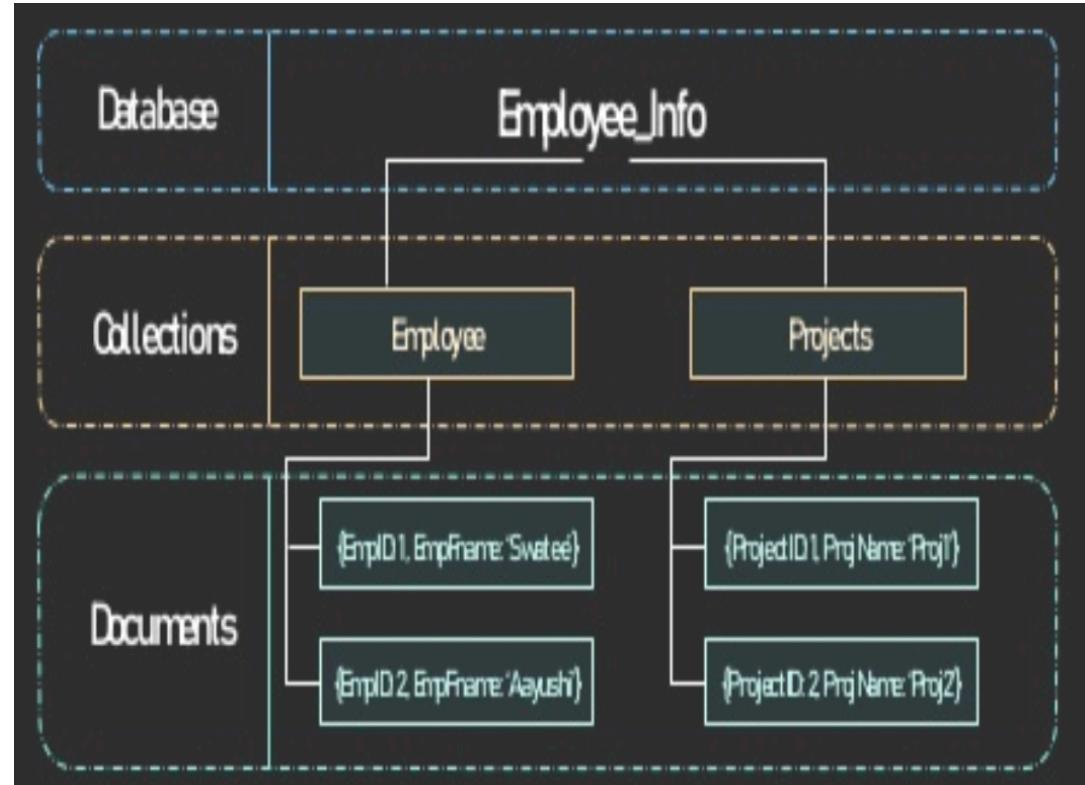
# When, Where & Why to Use NoSQL?

## RDBMS doesn't quite fit for \*some\* requirements

- Azure Data Lake (ADLS gen2), Google (BigTable) and Amazon(S3) decided to make their own to meet their needs
  - Distributed
  - Scalability
  - Control over performance characteristics
  - High availability
  - Low latency
  - Cheap

# Know NoSQL

- ✓ **NoSQL known as Not only SQL database.**
- ✓ **Provided mechanism for storage & retrieval of data.**
- ✓ **Next generation database.**
- ✓ **No specific schema & handle huge amount of data.**
- ✓ **Schema on Read vs Schema on Write**



# What is NoSql?

- Basically a large serialized object store\*
  - Retrieve objects by defined ID
- In general, doesn't support complicated queries\*
- Doesn't have a structured schema\*
  - Recommends denormalization
- Designed to be distributed (cloud-scale) out of the box
- Because of this, drops the ACID requirements
  - Any database can answer any query
  - Any write query can operate against any database and will “eventually” propagate to other distributed servers

## Opposite to ACID is BASE

### ➤ BASE

- Basically Available → guaranteed availability
- Soft state → state of system may change, even without a query (because of node updates)
- Eventually Consistent → the system will become consistent over time

## ACID Properties

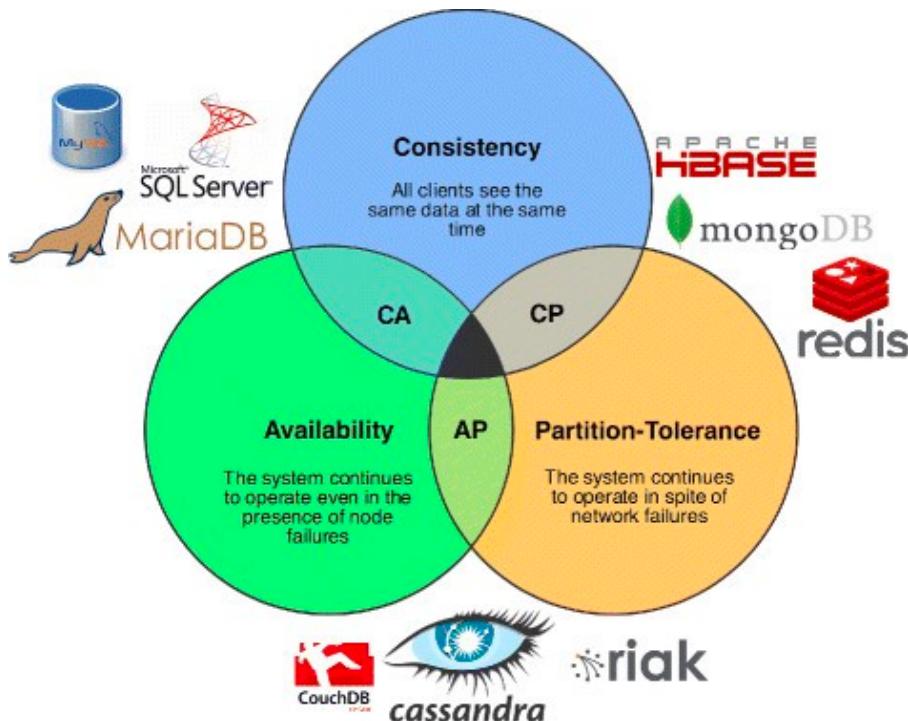
**A** Atomicity  
**C** Consistency  
**I** Isolation  
**D** Durability

## CAP Theorem

**C** Consistency  
**A** Availability  
**P** Partition Tolerance

# CAP Theorem

**Impossible for any shared data-system to guarantee simultaneously all of the following three properties**

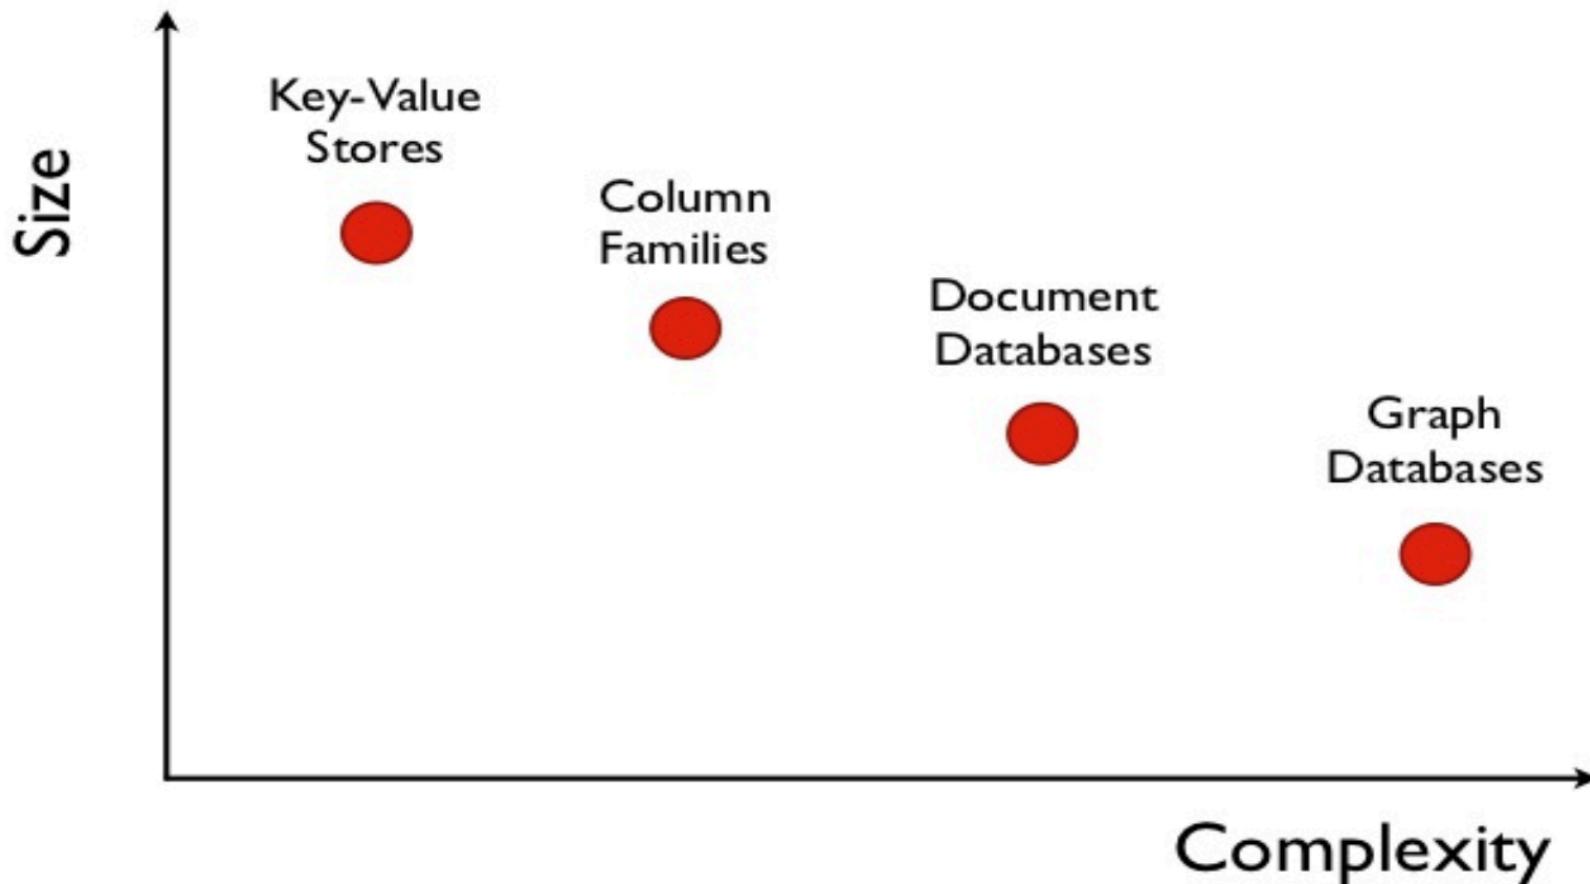


- **Consistency**: once data is written, all future read requests will contain that data
- **Availability**: the database is always available and responsive
- **Partition Tolerance**: if part of database is unavailable, other parts are unaffected.

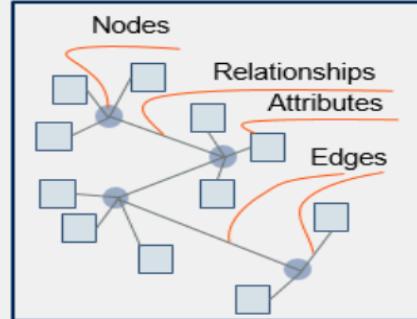
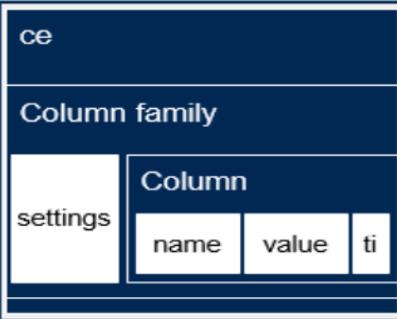
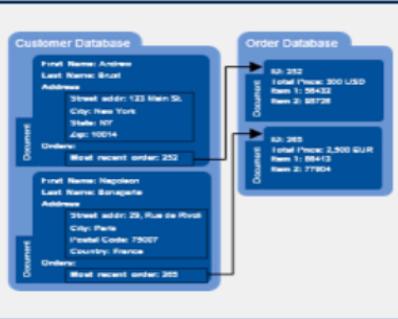
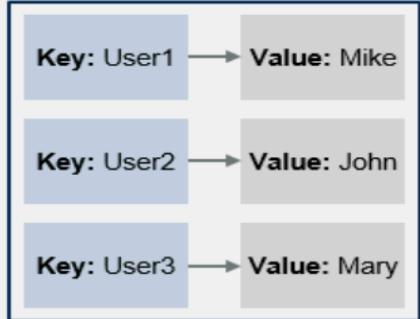
# NoSQL → Eventual Consistency

- Because of the distributed model, any server can answer any query
- Servers communicates amongst themselves at their own pace (*behind the scenes*)
- The server that answers your query might not have the latest data
- See a Facebook update from your friend delayed by some seconds may not make much difference

# NoSQL Data models



# NoSQL Database types



**Key-Value**

- Records generally are stored in key-data format
- Examples:
  - Aerospike
  - Microsoft Azure Table storage
  - Amazon SimpleDB
  - Basho

**Document**

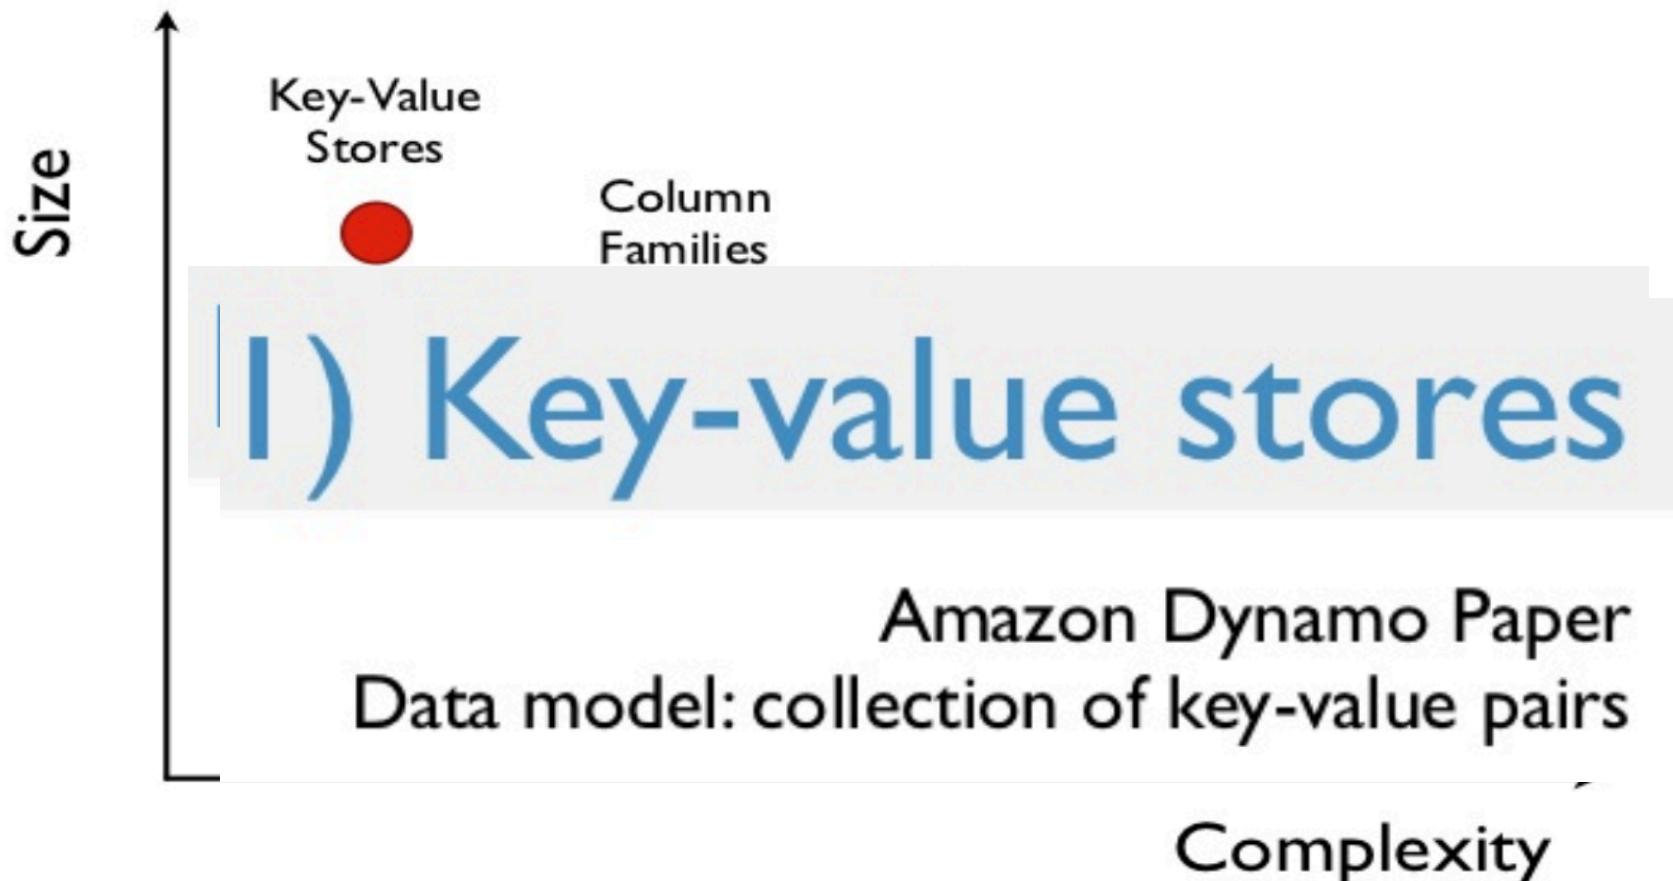
- Built to store document-centric data, such as JSON and/or XML documents
- Examples:
  - MarkLogic Server
  - Couchbase Server
  - Azure Cosmos DB
  - MongoDB
  - RavenDB

**Table-Style**

- Data model based on sparsely populated tables
- Examples:
  - Amazon DynamoDB
  - Apache HBase
  - DataStax Enterprise
  - Google Cloud Bigtable

**Graph**

- Accommodates graphs and relationship-centric data
- Examples:
  - Neo4j
  - Titan
  - Amazon Neptune



# Key-Value Data Store

	KEY	VALUE
Directory	Company MRH2 Labs	Phone # (408) XXX-XXX
Forwarding Table	IP Address 172.13.20.2	Interface: MAC Address Gi0/0: 38-F3-9C-AC-B0-99
Stock Trading	Order ID CP20183232049	Symbol, Exchange, Price MRH2,NSE,1390.05
Book Inventory	BookID:<ID> BookID: 1010	Hash of values Publisher: MRH Publications Author: Mohan ISBN: XX-XX-XX-XX

- The simplest type of NoSQL database
- Uses key-value pairs in a hash table
- Unique key as pointer to a value
- Logical group of keys are called a “bucket”
- Large fallbacks with running queries. Most do not support update operations that modify only part of the data

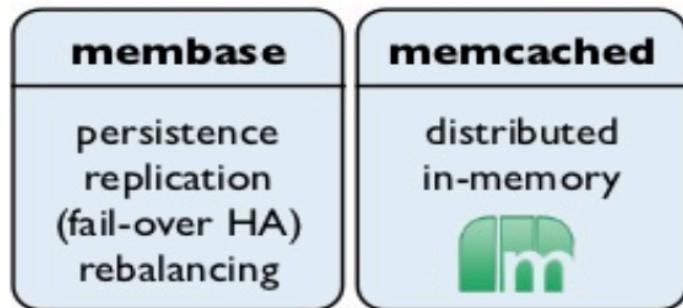


# Membase

Pass it on...  
Knowledge is power

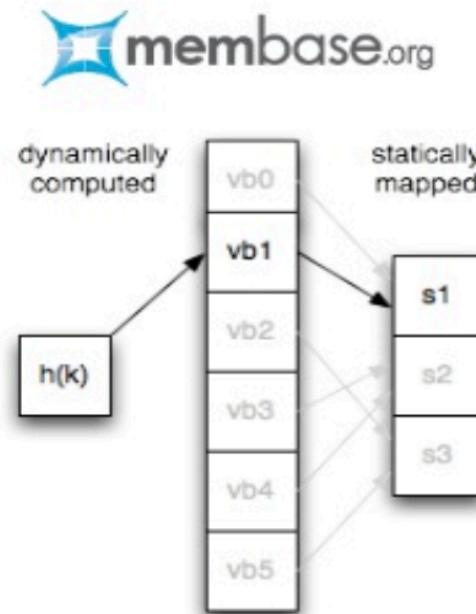
DHT (K-V), no SPoF

“VBuckets”



*Unit of consistency and replication  
Owner of a subset of the cluster key space*

<http://dustin.github.com/2010/06/29/memcached-vbuckets.html>



$h(k) \rightarrow vb1 \rightarrow s1$

*hash function + table lookup*

All metadata kept in memory (high throughput / low latency).  
Manual/Programmatic failover via the Management REST API.

LICENSE
Apache 2
LANGUAGE
C/C++ Erlang
API/PROTOCOL
REST/JSON memcached

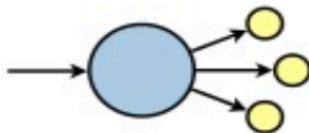
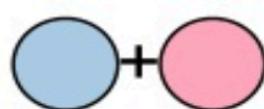
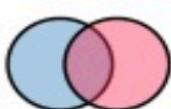
~~K-V store~~ “Data Structures Server”



redis

Map, Set, Sorted Set, Linked List

Set/Queue operations, Counters, Pub-Sub, Volatile keys



[http://redis.io/presentation/Redis\\_Cluster.pdf](http://redis.io/presentation/Redis_Cluster.pdf)



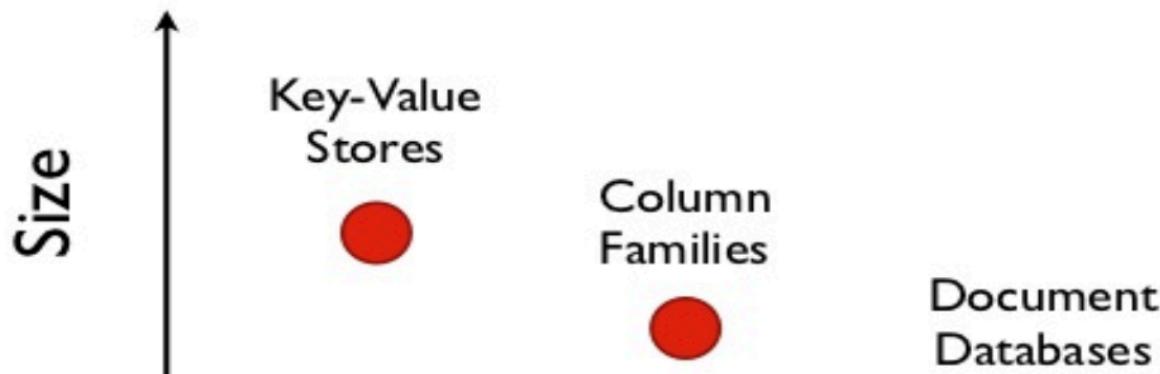
10-100K op/s (whole dataset in RAM + VM)



Persistence via snapshotting (tunable fsync freq.)

Distributed if client supports consistent hashing

LICENSE	 BSD
LANGUAGE	ANSI C
API	*
PROTOCOL	Telnet-like
PERSISTENCE	in memory bg snapshots
REPLICATION	master-slave



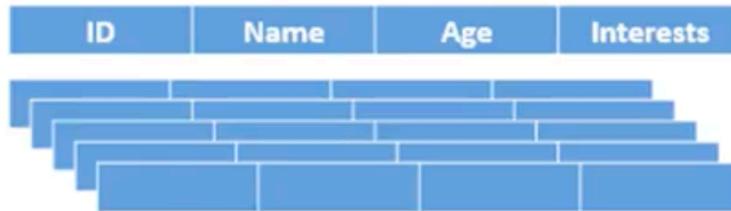
## 2) Column Families

Google BigTable paper

Data model: big table, column families →  
Complexity

# Column-Store Family Database

## Row Oriented (RDBMS Model)



## Column Oriented



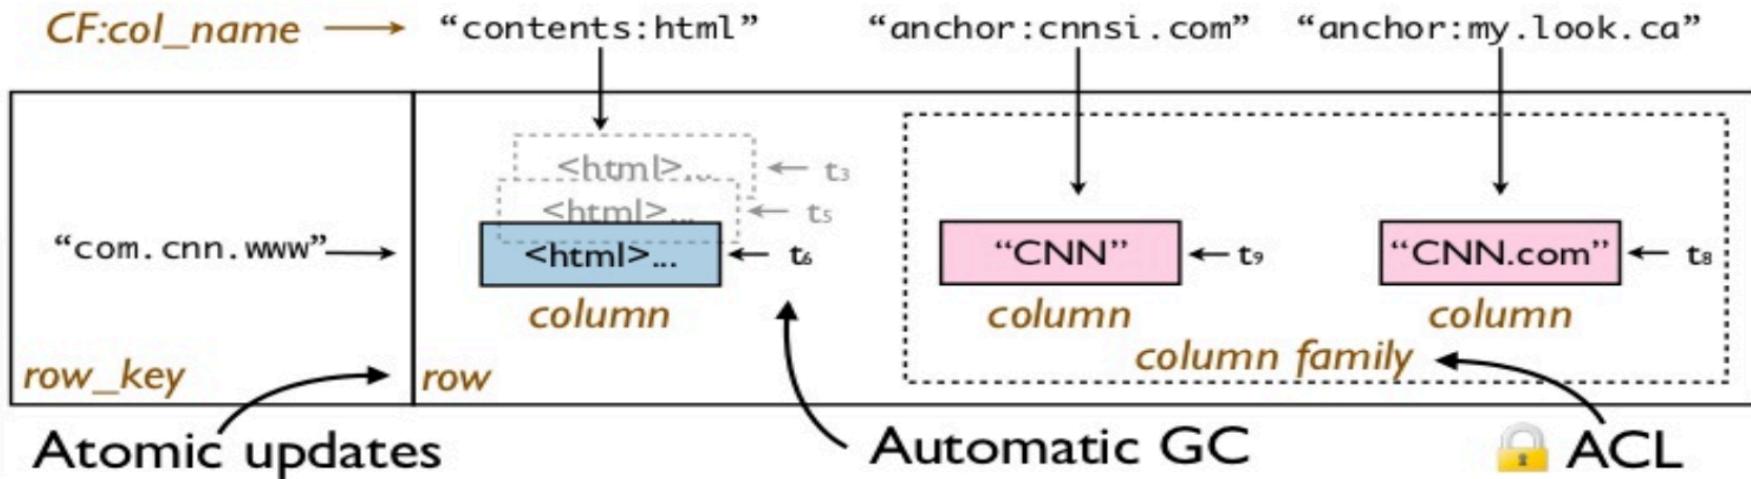
- Data is stored in cells that are grouped in columns of data rather than as rows of data
- Columns are grouped logically into column families. Column families can contain a virtually **unlimited** number of columns that can be created at **runtime** or the **definition** of the schema.
- These are generally for storing and processing very large amounts of data that is distributed across **many** machines.
- Faster access, search and data aggregation

# Google BigTable Paper

Sparse, distributed, persistent multi-dimensional sorted map indexed by (*row\_key*, *column\_key*, *timestamp*)



<http://labs.google.com/papers/bigtable-osdi06.pdf>



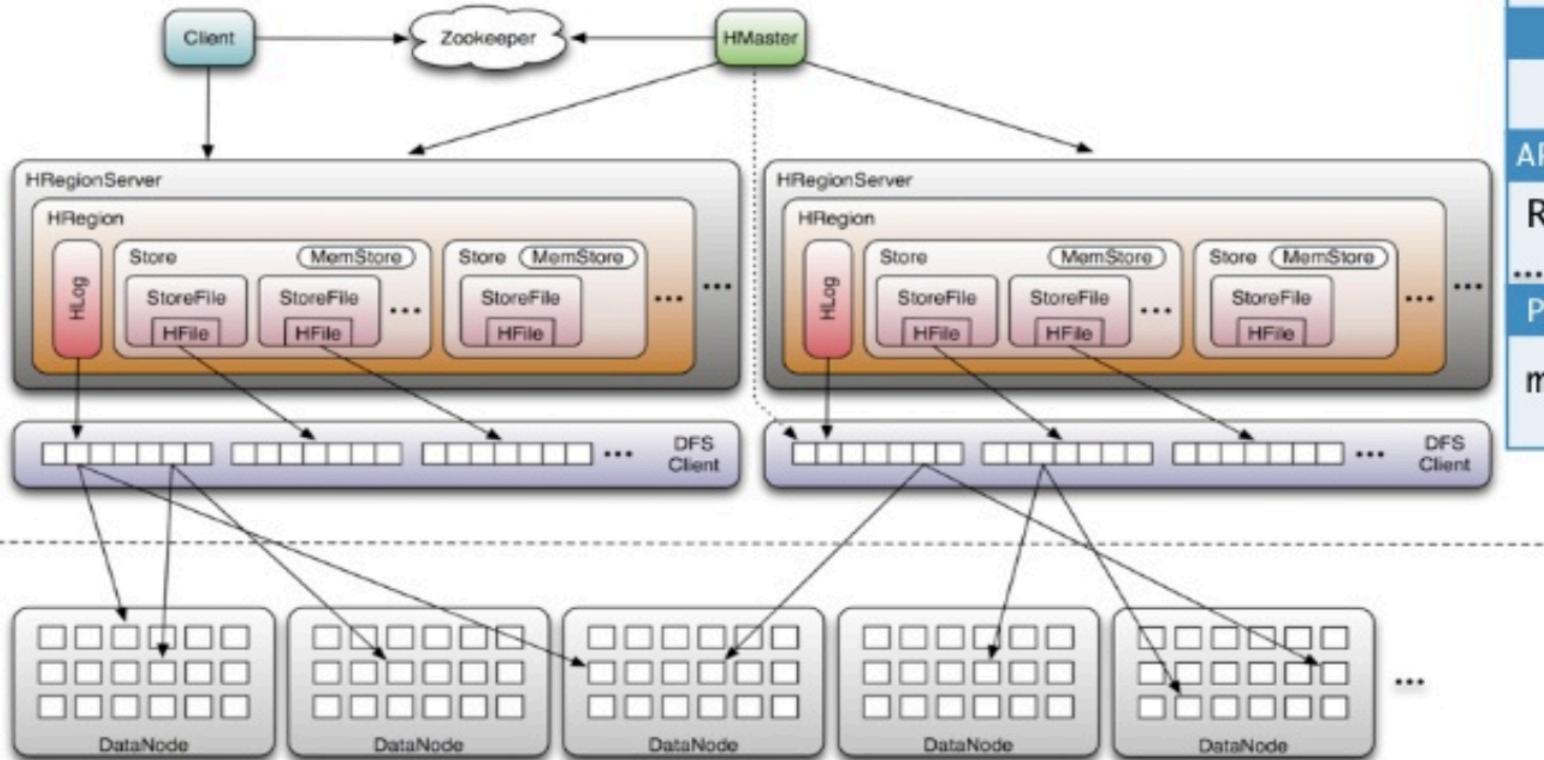
# HBase

Pass it on...  
Knowledge is power

## OSS implementation of BigTable



HBase

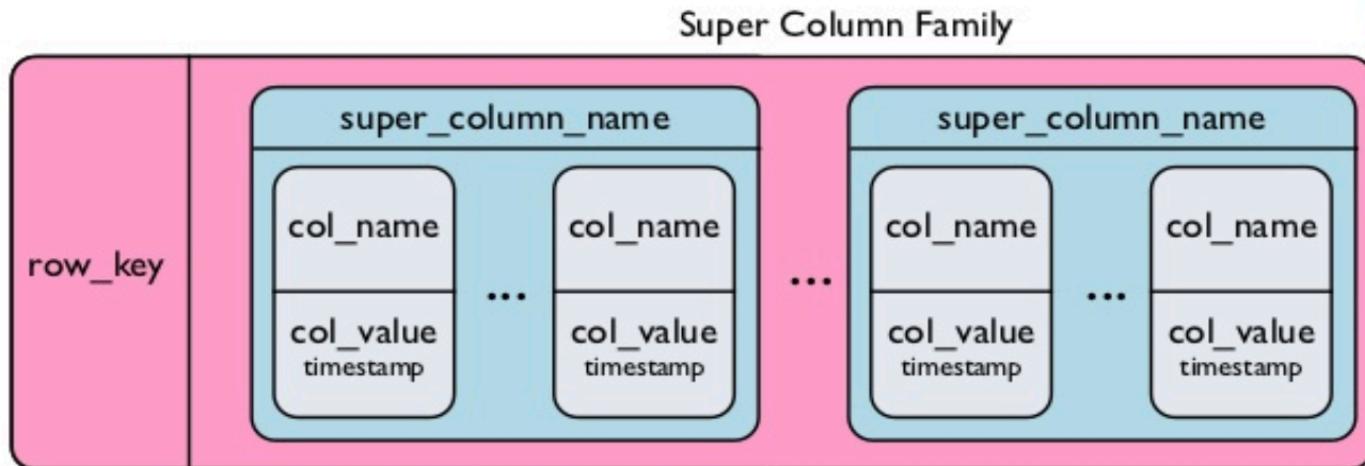


Hadoop

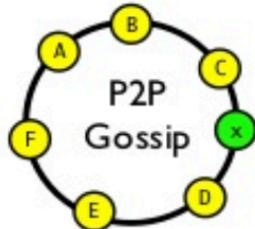
LICENSE
Apache 2
LANGUAGE
Java
API/PROTOCOL
REST HTTP ... Thrift
PERSISTENCE
memtable/ SSTable



## Data model of BigTable, infrastructure of Dynamo



```
keyspace.get("column_family", key, [ "super_column", ] "column")
```

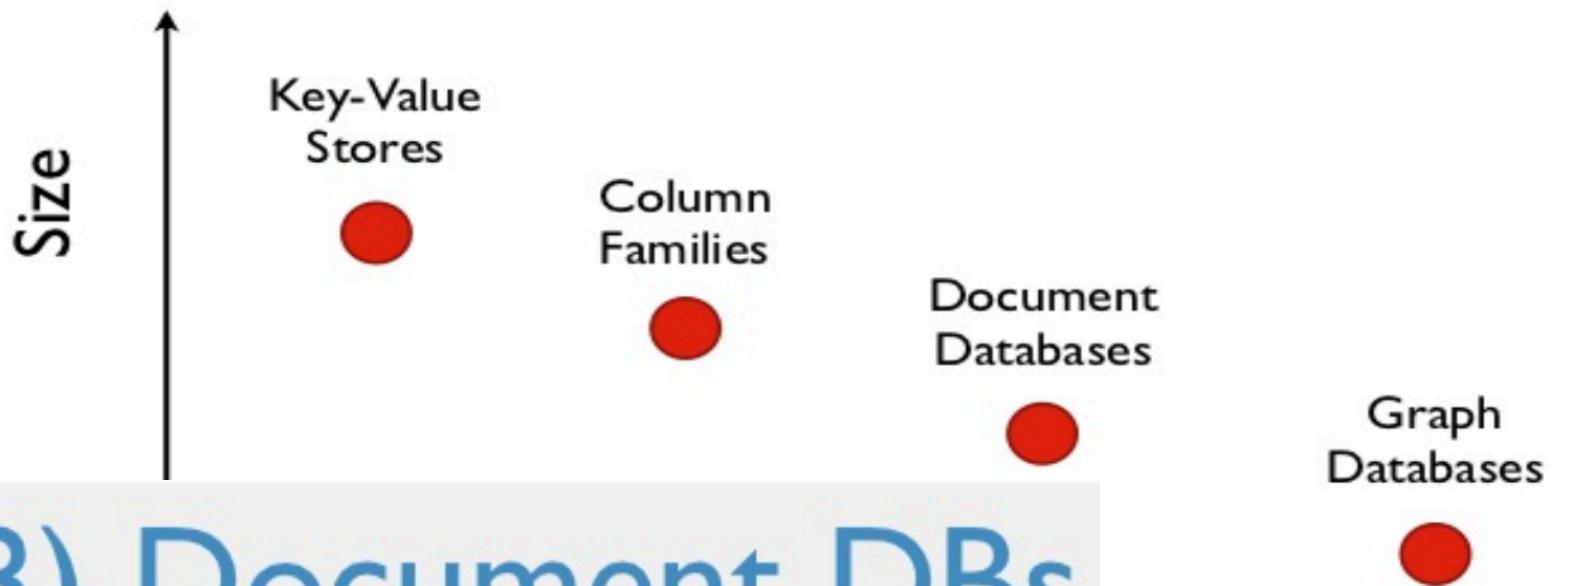


ALL  
ONE  
QUORUM

RandomPartitioner (MD5)  
OrderPreservingPartitioner  
Range Scans, Fulltext Index (Solandra)

LICENSE
Apache 2
LANGUAGE
Java
PROTOCOL
Thrift Avro
PERSISTENCE
memtable/ SSTable
CONSISTENCY
Tunable R/W/N

# NoSQL Data models



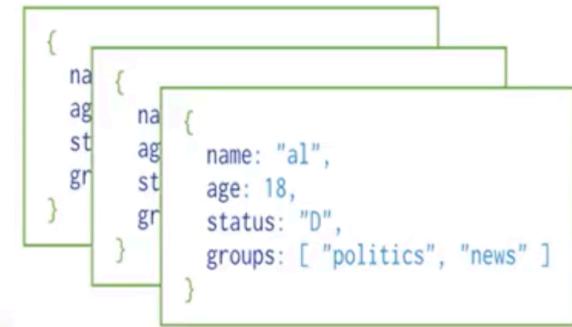
## 3) Document DBs

Data model: collection of K-V collections

Lotus Notes

Complexity

# Document Database



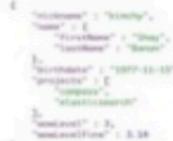
- Uses key value pairs compressed as a document store similar to a key-value store, but values stored (“documents”) provide some structure and encoding of the managed data.
- Can use multiple formats with document databases – **JSON, XML, BSON**
- Can perform nested queries and can do much more with your data when compared to a key-value store.
- **MongoDB** and **Couchbase** are the most popular document databases at this time

# CouchDB

<http://horicky.blogspot.com/2008/10/couchdb-implementation.html>

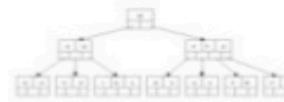


Pass it on...   
Knowledge is power

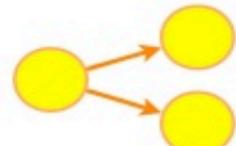
  
{"id": "1",  
 "name": "Nancy",  
 "age": 25,  
 "friends": ["John",  
 "Mike"],  
 "coordinates": {"lat": 37.8131, "lon": -122.4278},  
 "temperature": 72,  
 "humidity": 50,  
 "timestamp": 1283456789012,  
 "sequence": 23,  
 "sequenceTime": 1.33}

JSON docs

  
map-reduce "views"  
(materialised resultset)



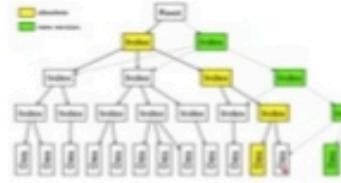
Storage + View Indexes (B+Tree)  
[by\_id\_index, by\_seqnum\_index]



Replication used  
as a way to scale  
transactions volume



Conflict Resolution  
at application level



MVCC (copy-on-modify)  
Volatile Versioning



Online Compaction  
(very primitive VACUUM)



Update validation /  
Auth triggers



Delayed commits  
(write performance)

## LICENSE



Apache 2

## LANGUAGE

Erlang

## API/PROTOCOL

REST/JSON

## PERSISTENCE

Append-only  
B+Tree

## CONCURRENCY

MVCC

## CONSISTENCY

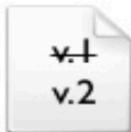
crash-only  
design

## REPLICATION

multimaster

bson\_encode()  
bson\_decode()

BSON serialisation  
(storage & transfer)



Update in place  
(no versioning, no  
append-only log)

GROUP BY

Map/Reduce  
(well, aggregation)



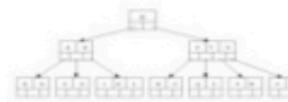
Auto-Sharding,  
Master-Slave,  
Auto-Failover



Geo-Spatial Indexes



No ACK on Updates  
(or ensure N replicas)



B-Tree Indexes  
(on different cols too)



Persistence via  
Replication +  
Snapshotting



LICENSE



AGPLv3

LANGUAGE

C++

API/PROTOCOL

REST/BSN  
\*

PERSISTENCE

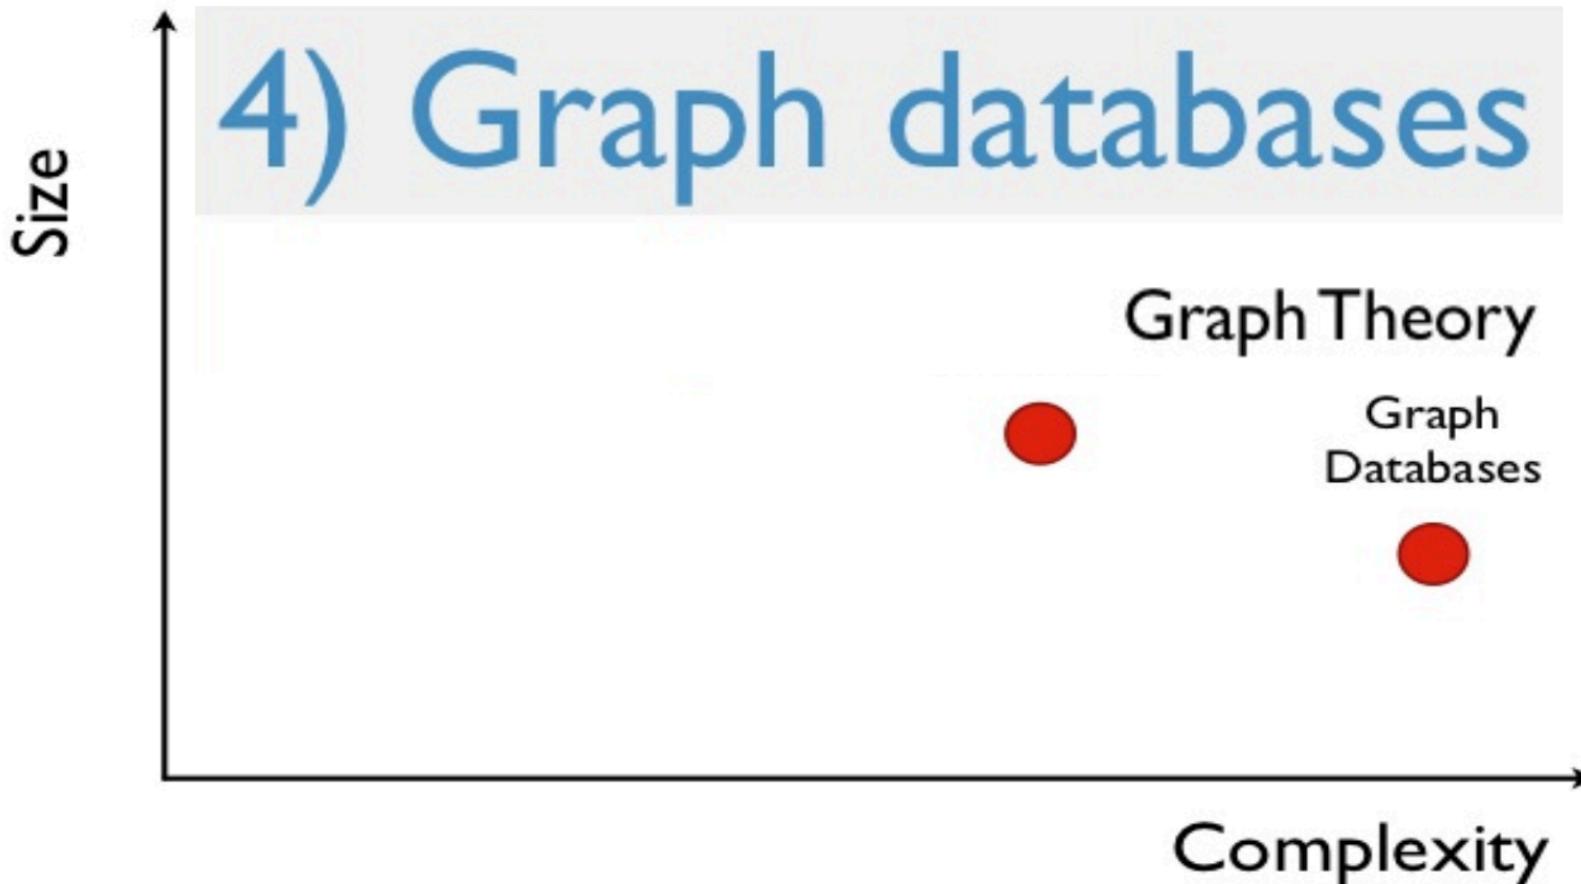
B+Tree,  
S snapshots

CONCURRENCY

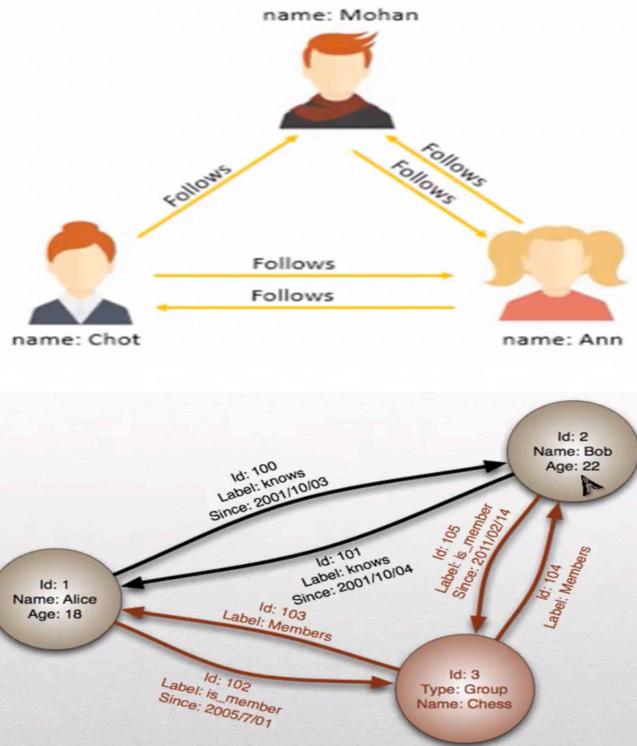
In-place  
updates

REPLICATION

master-slave  
replica sets



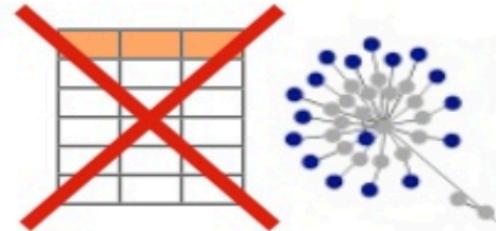
# Graph Database



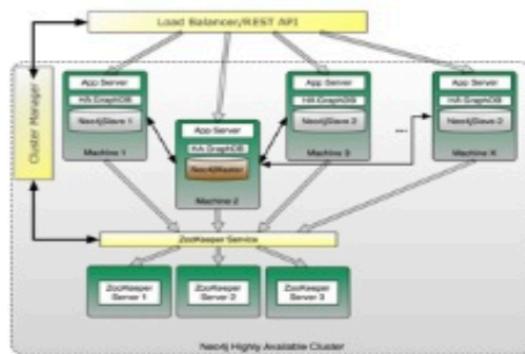
- Based on “graph theory”. Uses a flexible graphical representation as oppose to rigid tables and columns
- Uses “edges” and “nodes” instead of key-value pairs. This creates index-free adjacency
- Often used to store information about networks, such as social connections
- Compared to **RDBMS**, graph databases are usually faster for associative data sets and map more directly to the structure of object-oriented applications



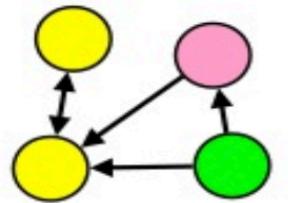
<http://docs.neo4j.org/chunked/stable/ha-how.html>



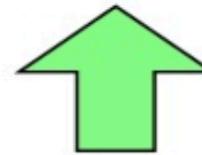
Graph Data Structure



HA cluster with ZooKeeper  
(nodes = exact replicas)



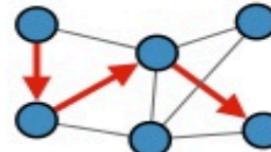
Nodes,  
Relationships,  
Properties on both



Vertical Scalability  
(1000s times faster,  
but not distributed)



Physical structure:  
LinkedList stored on disk



High-performance  
node traversal

LICENSE
AGPLv3 / Commercial
LANGUAGE
Java
API/PROTOCOL
REST Java SPARQL
PERSISTENCE
On-disk linked-list



SPARQL

# Azure Cosmos DB: multi-model DB

Pass it on...  
Knowledge is power

Azure Cosmos DB

The diagram illustrates the global reach and distributed nature of Azure Cosmos DB, showing connections between four continents.

Below the map, there is a horizontal navigation bar with four categories:

- Key-Value
- Column-Family
- Documents
- Graph

Icons representing various database models and APIs are displayed above the navigation bar:

- Table
- SQL
- JavaScript
- {LEAF}
- API for MongoDB
- Gremlin
- Cassandra
- Apache Spark
- etcd
- ETCD
- ...more APIs coming

Microsoft's  
globally  
distributed  
multi-model  
database  
service on  
cloud

# globally distributed multi-model DB

Pass it on...  
Knowledge is power

Azure Cosmos DB



Table SQL JavaScript API for MongoDB Gremlin Cassandra Spark etcd ...more APIs coming

Document SQL API (JSON) MongoDB API (BSON)

Key-Value Table API (replaces Azure Table Storage)

Graph Gremlin API (graph traversal language)

Columnar Cassandra API

Key-Value Column-Family Documents Graph

Microsoft's  
globally  
distributed  
multi-model  
database  
service on  
cloud



# Azure Cosmos DB

... more coming soon

Key-Value



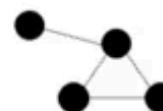
Column-family



Documents



Graph



Global distribution

Elastic scale out

Guaranteed low latency

Five consistency models

Comprehensive SLAs

A globally-distributed, multi-model database service

# Cosmos DB: Key Features

- **Global distribution** - transparently replicates your data, geo-redundancy, multi-master writes
- **Always On** – provides 99.999 high availability SLA
- **Elastic scalability** – transparent horizontal partitioning and multi-master replication, can scale storage and throughput separately
- **Predictable low latency:** latch-free and wire optimized engine, <10ms
- **5 consistency models:** strong, bounded staleness, session, consistent prefix, and eventual
- **Multi-model:** key-value pairs, Column family, Document and Graph
- **Automatic indexing** - not need any schema or any secondary indexes
- **Secure and enterprise ready** – data encrypted at rest and in motion, certified for a wide range of compliance standards

# Comparison SQL vs. NoSQL



RDBMS	NoSQL
Structured Data	Semi-structured and Unstructured Data
Scale Up (Vertical Scalability)	Scale Out (Horizontal Scalability)
Strict Schema	Schema-free or Dynamic Schema
ACID Compliance	Eventual Consistency
Custom High-Availability	Auto High-Availability

## Relational World

### STORE

ID	Branch	Type	Count	SalesAssistant	Product
2313	Boston	Shoes	36	324	756
2654	Worcester	Shoes	39	354	567
6546	Weymouth	ShoeCareProducts	127	654	445

### CUSTOMER

ID	GivenName	FamilyName
324	Jack	Bauer
354	Tony	Stark
654	Charles	Xavier

### PRODUCT

ID	Color	Size
756	red	20-28
567	blue	30-41
787	brown	36-48

### TRANSACTION

ID	Type	LeatherType
445	Cream	Velour
676	Brush	Smooth leather
987	Spray	all

Not all DB type Are equal. Right DB for Workload Type

## NoSQL World

### Document - JSON

```

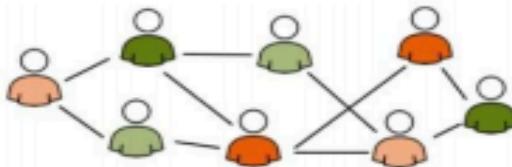
{
  "type": "sweater",
  "color": "blue",
  "size": "M",
  "material": ["wool", "cotton"],
  "form": "turtleneck"
}

{
  "type": "pants",
  "waist": 32,
  "length": 34,
  "color": "blue",
  "material": "cotton"
}

{
  "type": "diagonal screen size": 46,
  "hdmi inputs": 3,
  "wall mountable": true,
  "built-in digital tuner": true,
  "dynamic contrast ratio": "50,000:1",
  "Resolution": "1920x1080"
}

```

### GRAPHS



### KEY VALUES



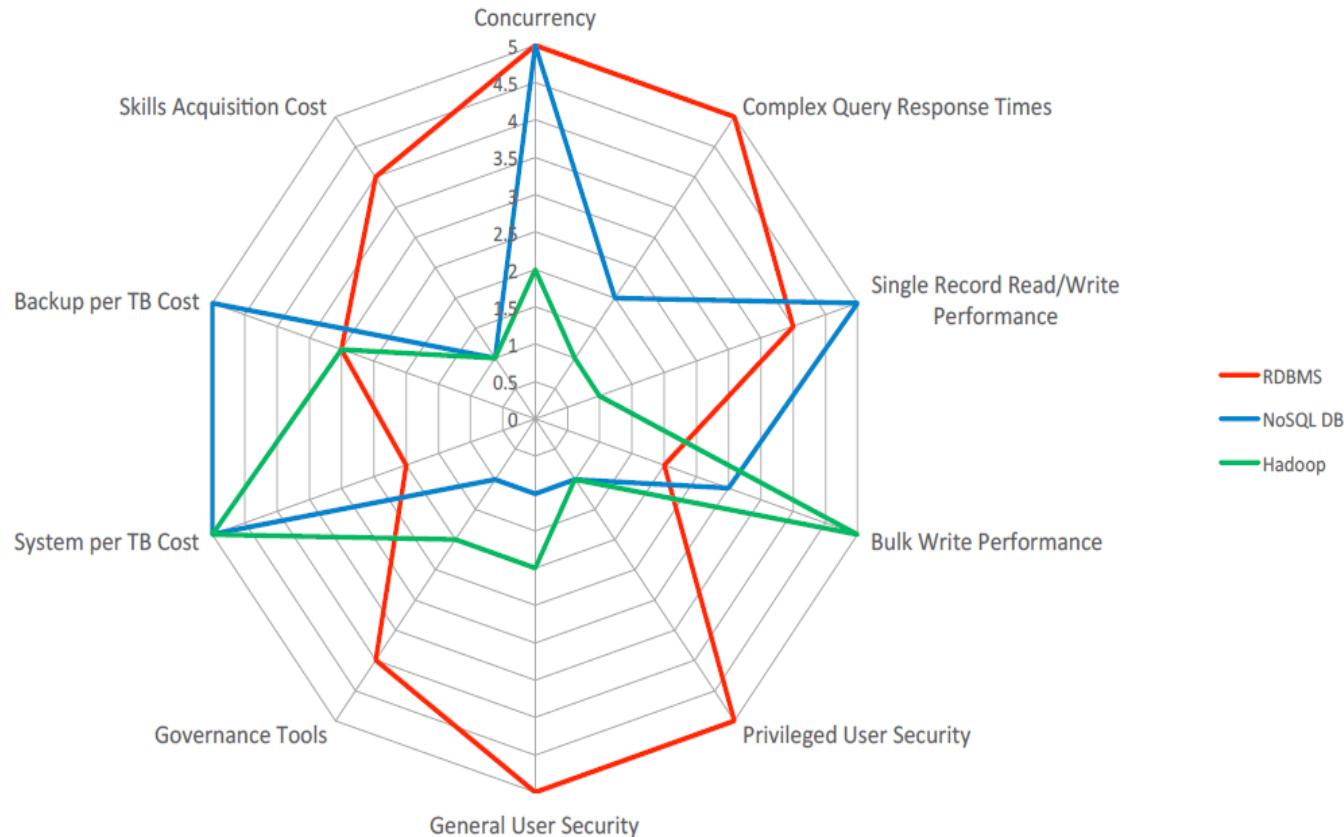
# Comparison SQL, NoSQL, Hadoop

Pass it on...  
Knowledge is power

Performance

Security

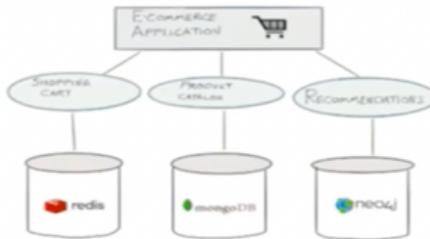
Cost



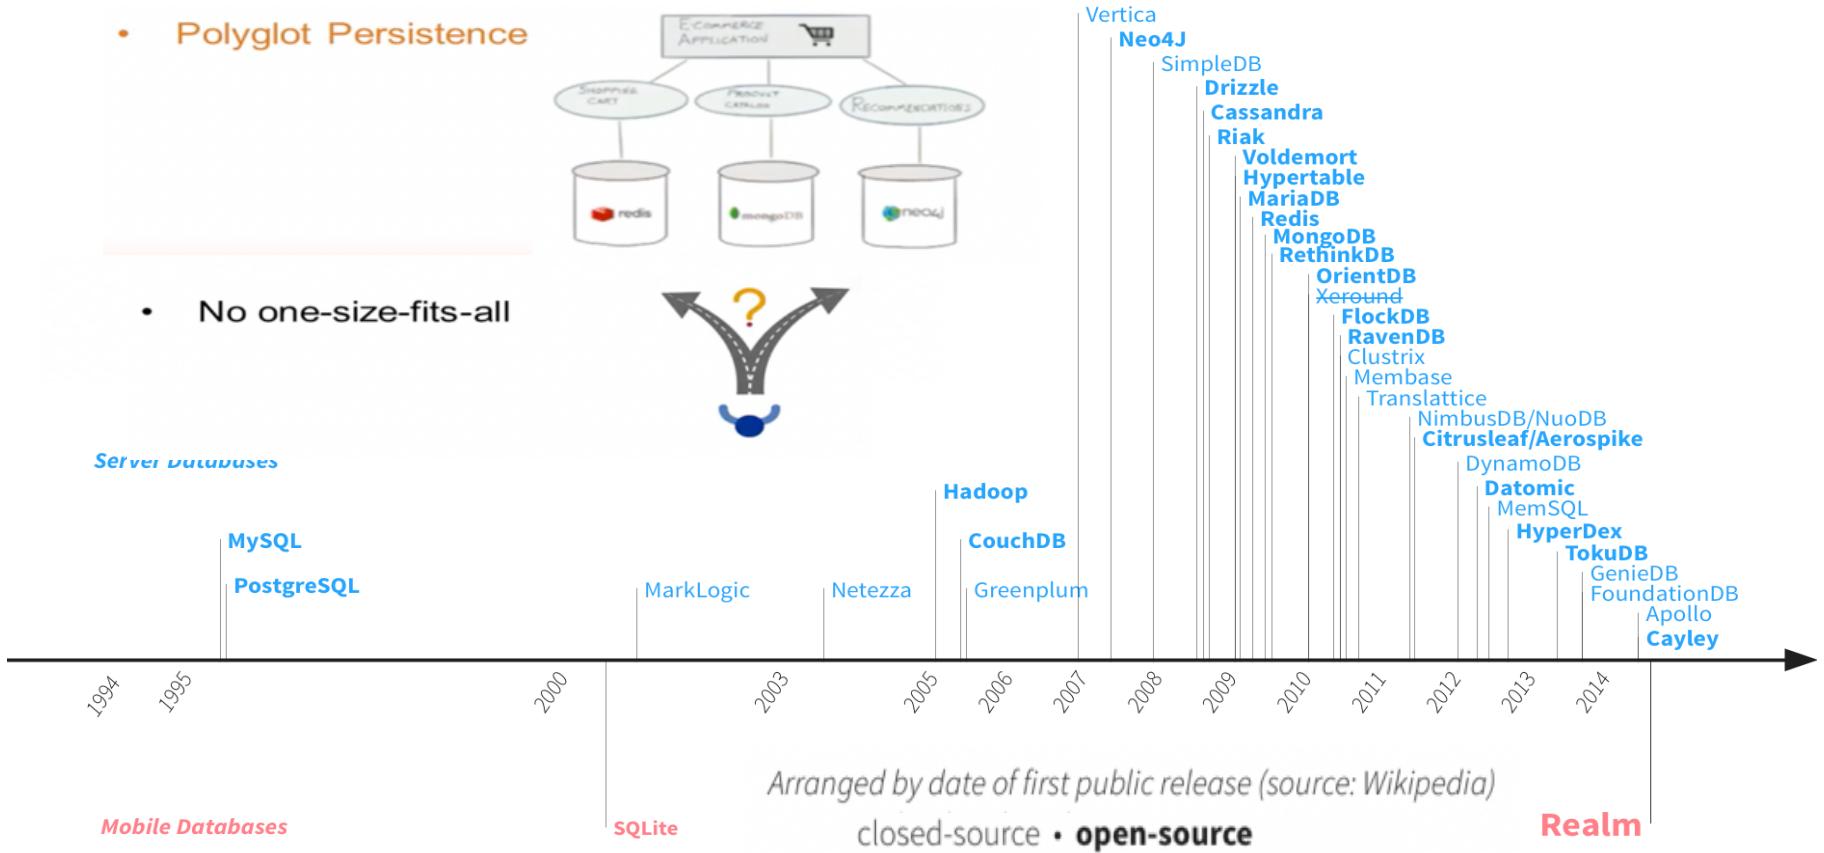
# No one-size-fits-all

Pass it on...  
Knowledge is power

- Polyglot Persistence



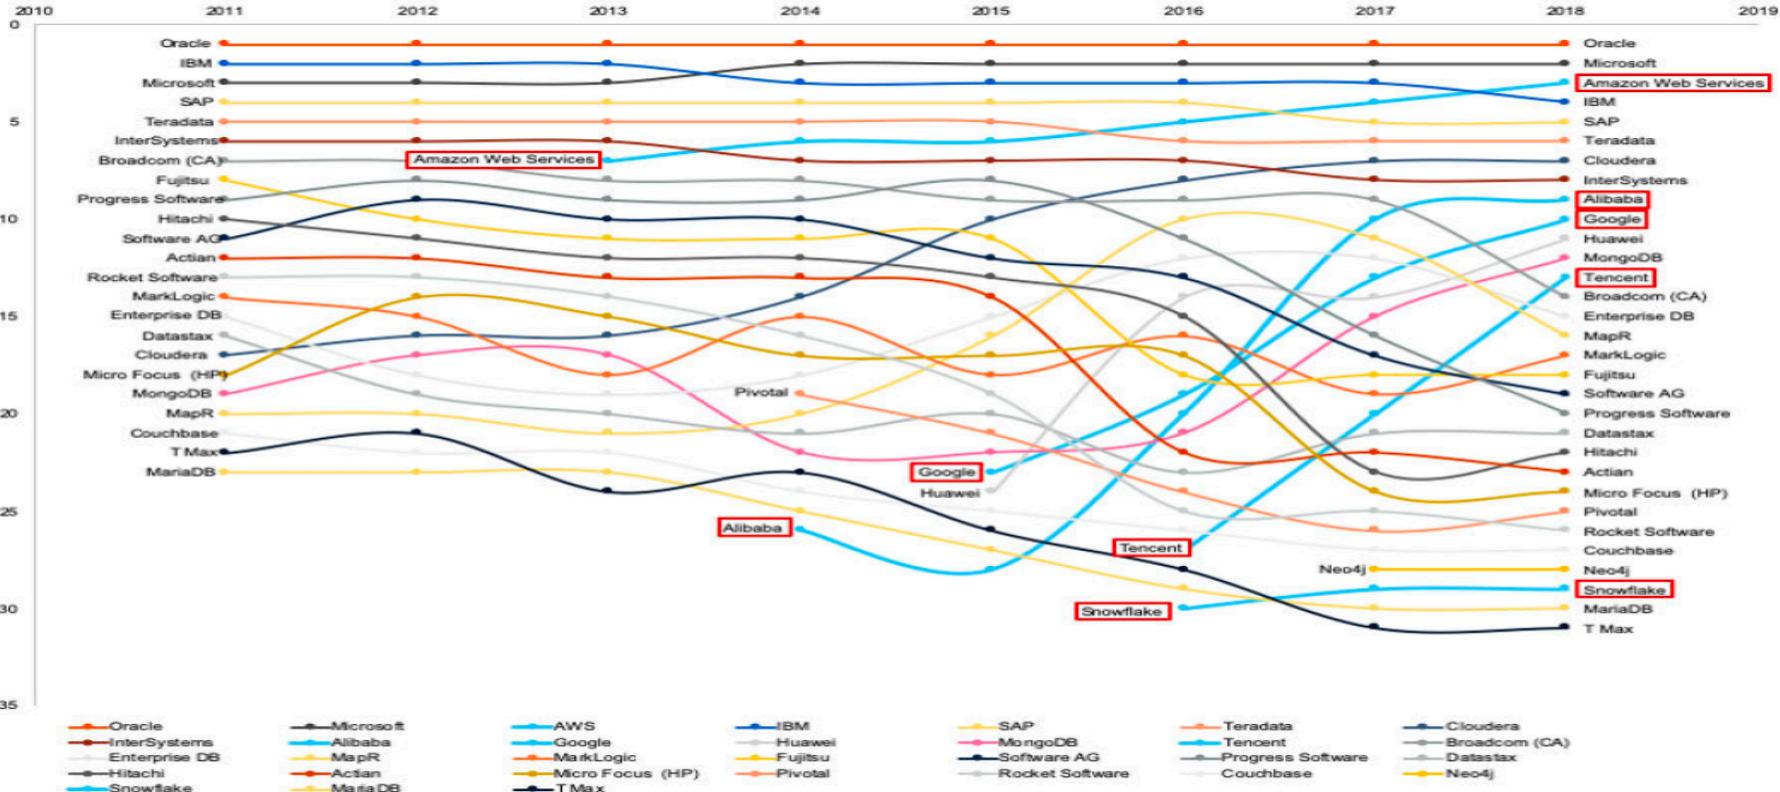
- No one-size-fits-all



# Gartner Market Share Ranking, 2011-2018

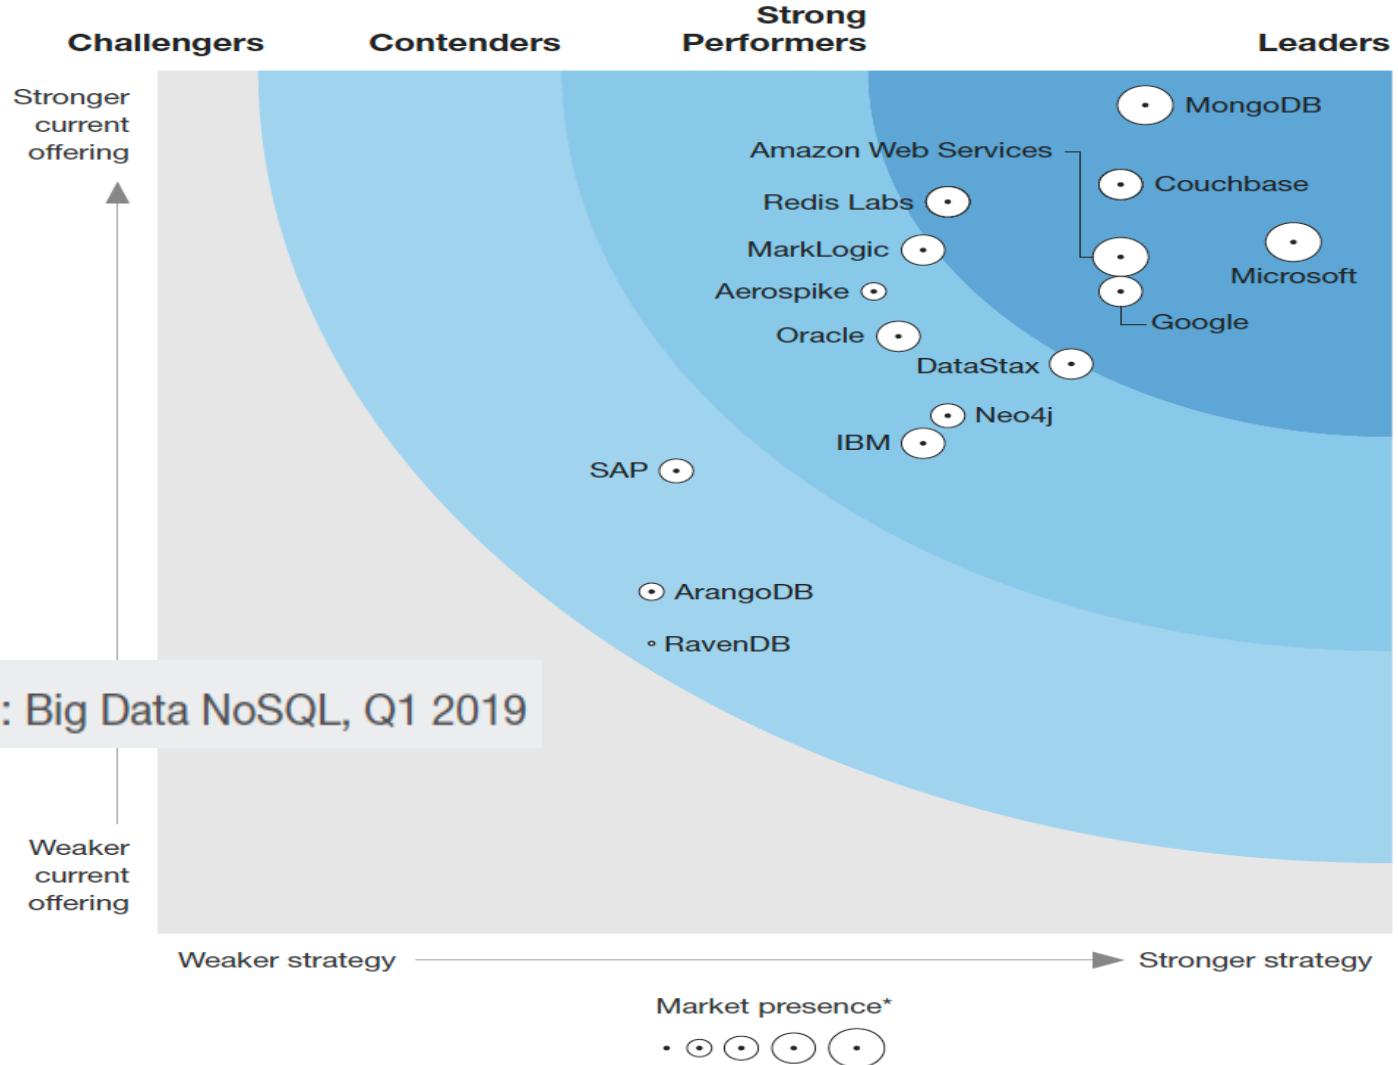
Rank

Pass it on...   
Knowledge is power



Source: Gartner (June 2019)

Note: The following historical vendor revenues were combined to reflect the state of the market in 2019:  
 Cloudera, reflecting the merger with Hortonworks; Micro Focus, reflecting the acquisition of HPE Vertica; Broadcom, reflecting the acquisition of CA.



# SUMMARY SQL or NoSQL?

## Projects where SQL is ideal:

- ✓ Logical related discrete data requirements which can be identified up-front.
- ✓ Data integrity is essential.
- ✓ Standards-based proven technology with good developer experience and support.

## Projects where NoSQL is ideal:

- ✓ Unrelated, indeterminate or evolving data requirements.
- ✓ Simpler or looser project objectives, able to start coding immediately.
- ✓ Speed and scalability is imperative.