



# Distributed Computing With Hadoop

## (Introduction)

Rashmi Kansakar

# What is Hadoop?



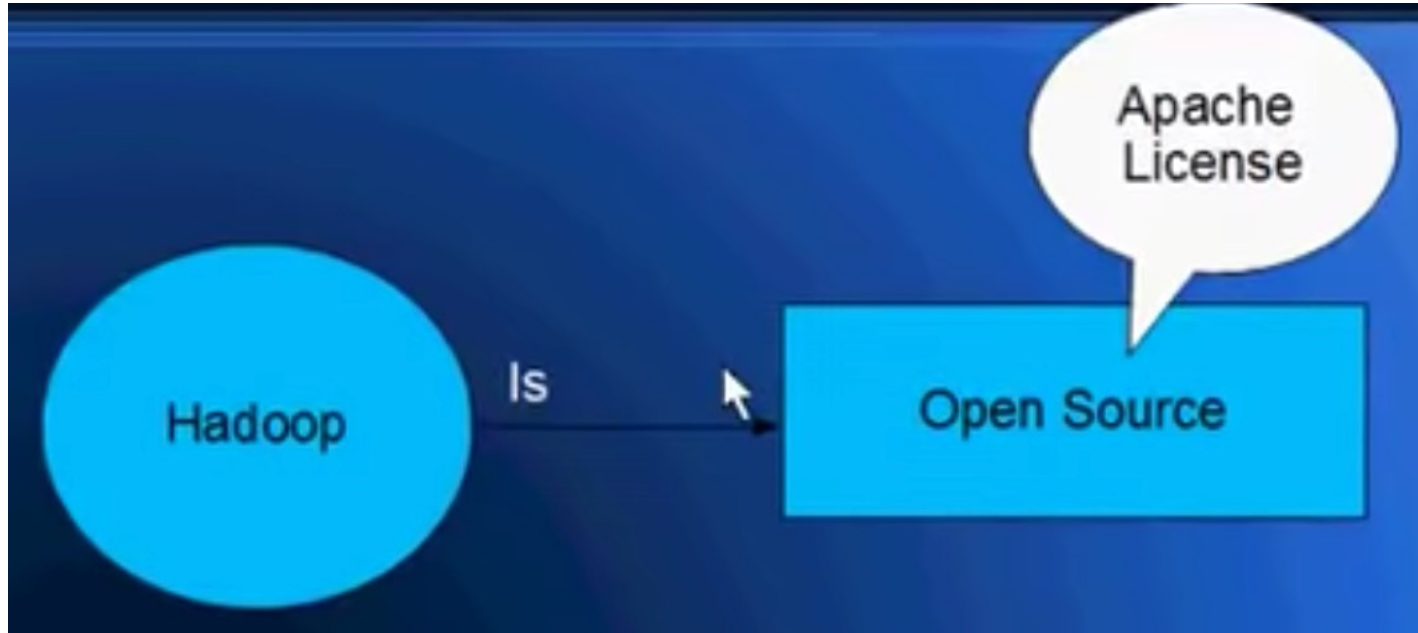
# Objective




# Open Source



# Apache



# Big Data Challenge Points



# Data Never Sleeps 7.0

Pass it on...   
Knowledge is power





# DATA NEVER SLEEPS 8.0

## How much data is generated *every minute*?

In 2020, the world changed fundamentally—and so did the data that makes the world go round. As COVID-19 swept the globe, nearly every aspect of life—from work to working out—moved online, and people depended more and more on apps and the Internet to socialize, educate and entertain ourselves. Before quarantine, just 15% of Americans worked from home. Now over half do. And that's not the only big shift. In our 8th edition of Data Never Sleeps, we bring you the latest stats on how much data is being created in every digital minute—a trend that shows no sign of stopping.

The world's Internet population is growing significantly year over year. As of April 2020, the internet reaches 59% of the world's population and now represents 4.57 billion people — a 6% increase from January 2019.




GLOBAL INTERNET POPULATION GROWTH 2014-2020  
(IN BILLIONS)

As the world changes, businesses need to change with the times—and that requires data. Every click, swipe, share or like tells you something about your customers and what they want, and Domo is here to help your business make sense of all of it. Domo gives you the power to make data-driven decisions at any moment, on any device, so you can make smart choices in a rapidly changing world.

Learn more at [domo.com](https://domo.com)




# Data Growth over the years



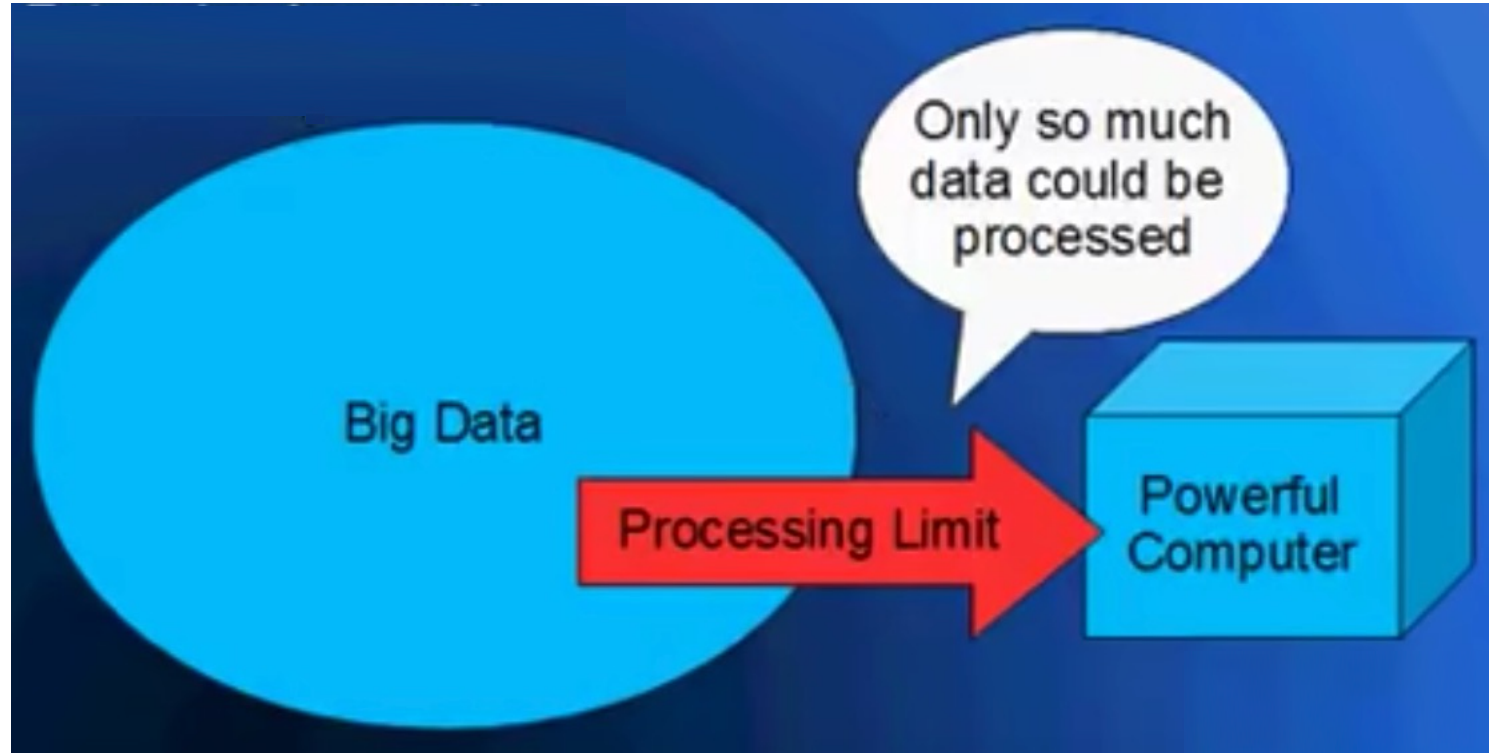
# Traditional Approach

## Monolithic Approach



# Traditional Approach

## Monolithic Approach




# Problems with Big Data



Storing huge and exponentially growing  
datasets



Processing data having complex structure  
(structured, un-structured, semi-  
structured)



Bringing huge amount of data to  
computation unit becomes a bottleneck

# 5 V's of Big Data?

## Volume

Data is being generated at an accelerating speed

## Variety

Different kinds of data is being generated from various sources

## Velocity


Data is being generated at an alarming rate

## Value

Mechanism to bring the correct meaning out of the data

## Veracity

Uncertainty and inconsistencies in the data



Age	Sex	Height	Wt
45	?	5.04	0.80
55	4.4	0.85	unknown
unknown	7.8	1.80	0.45
60	6.8	?	0.76

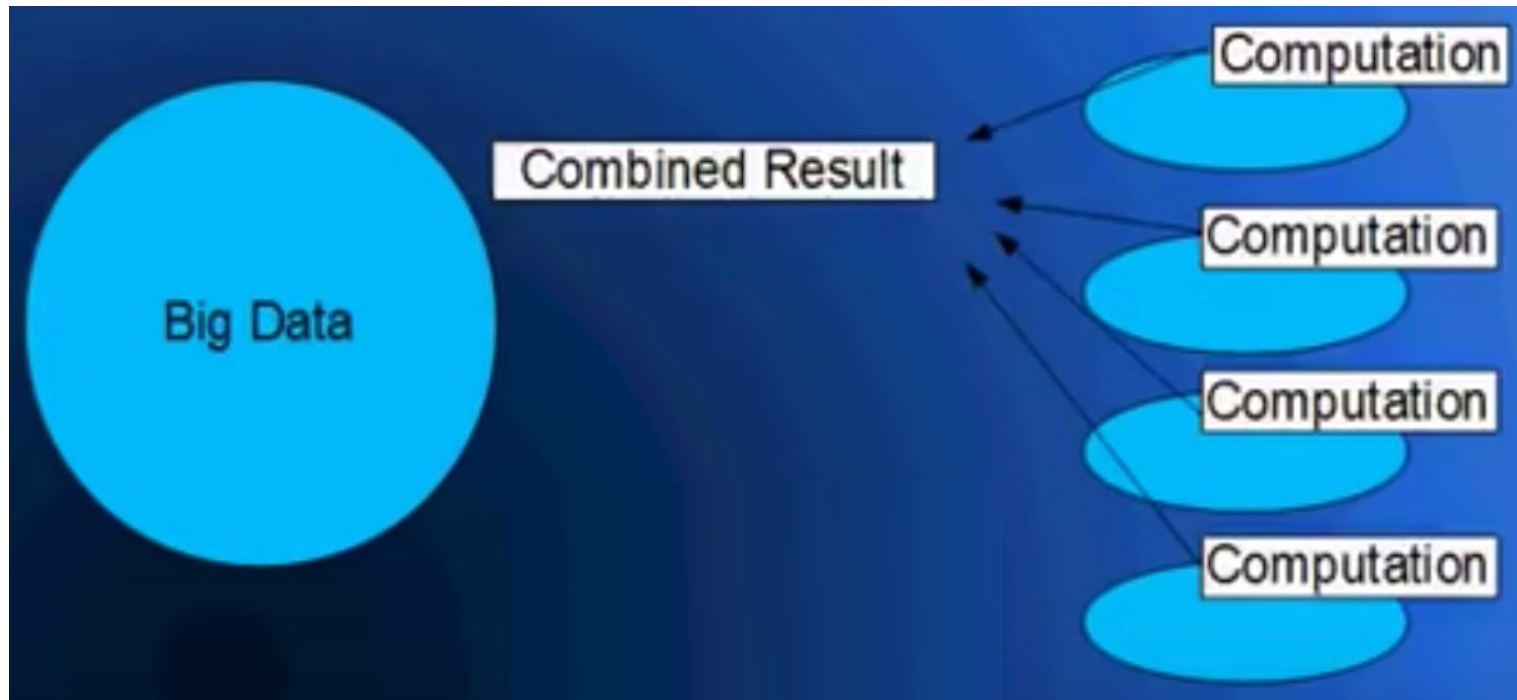
# Breaking the Data

## Hadoop's Approach




# Breaking the Data

## Hadoop's Approach







# Distributed Model



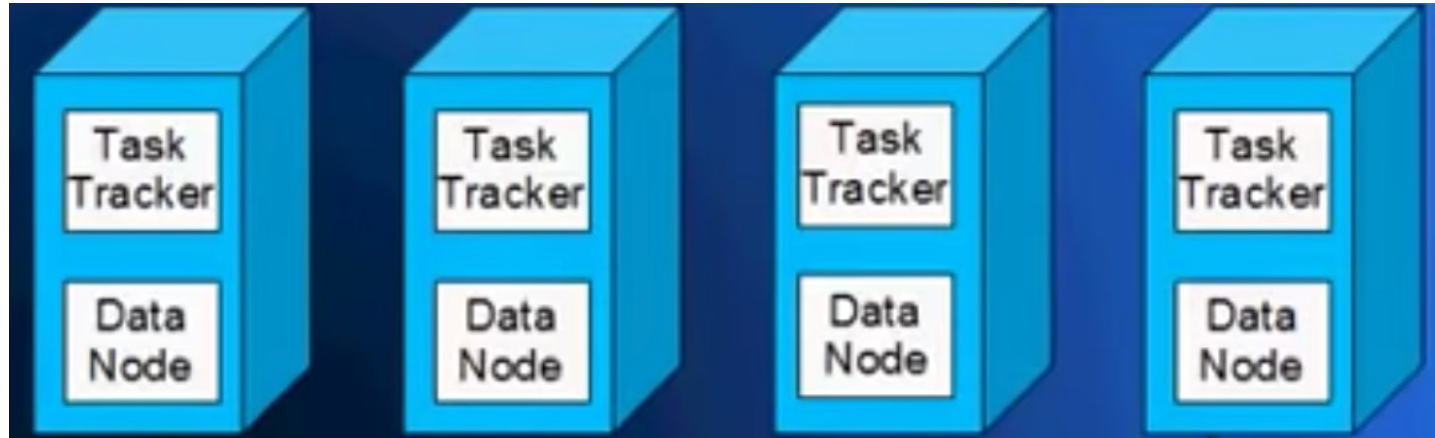
# Linux Based



# Task Trackers and Data Nodes




# Slaves

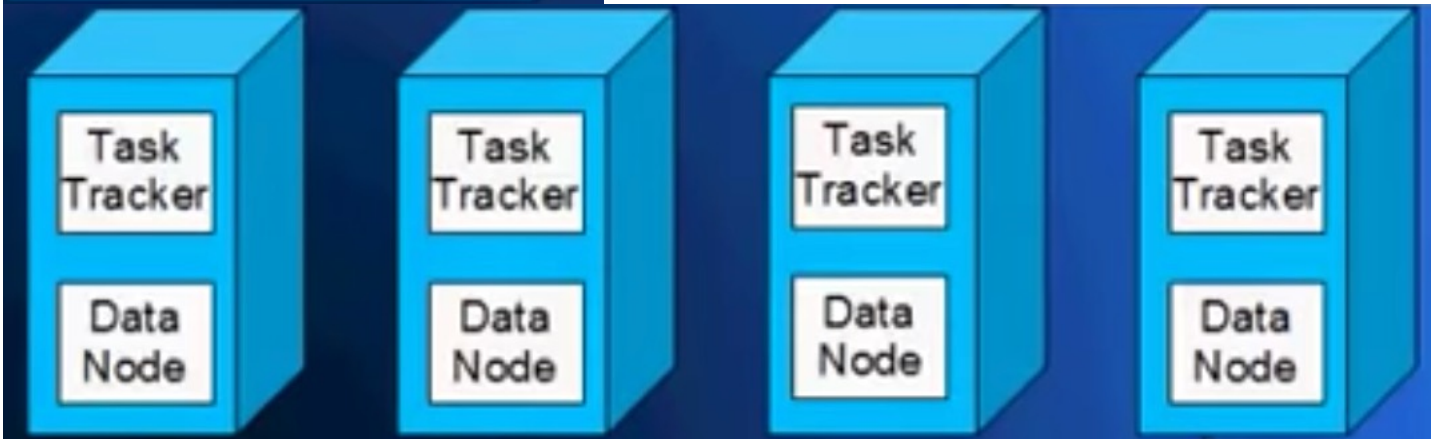


# Master Slave Architecture

Master



Slaves




# Components

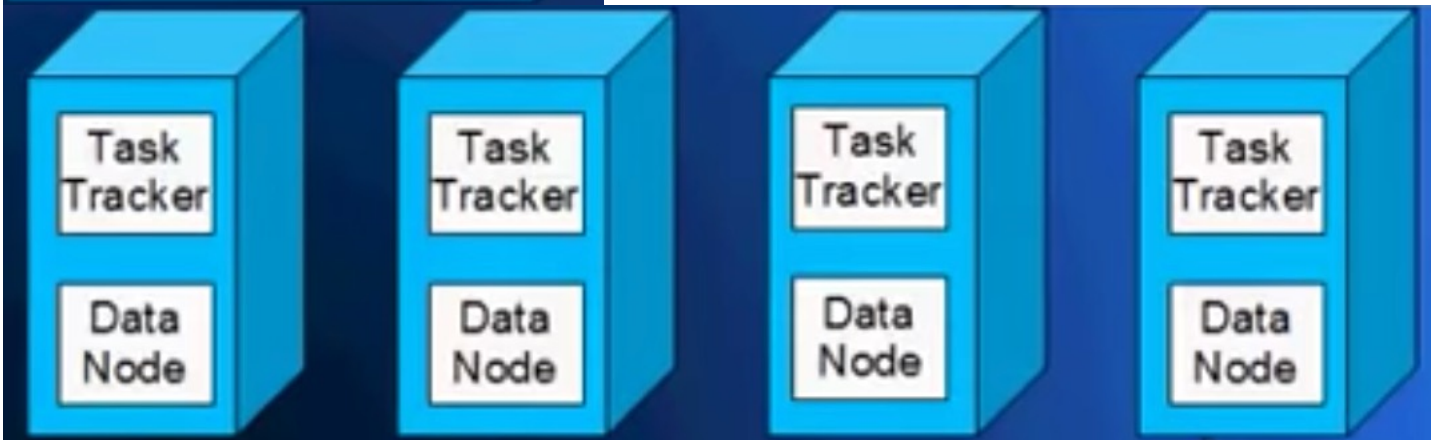


# HDFS

Master




Slaves



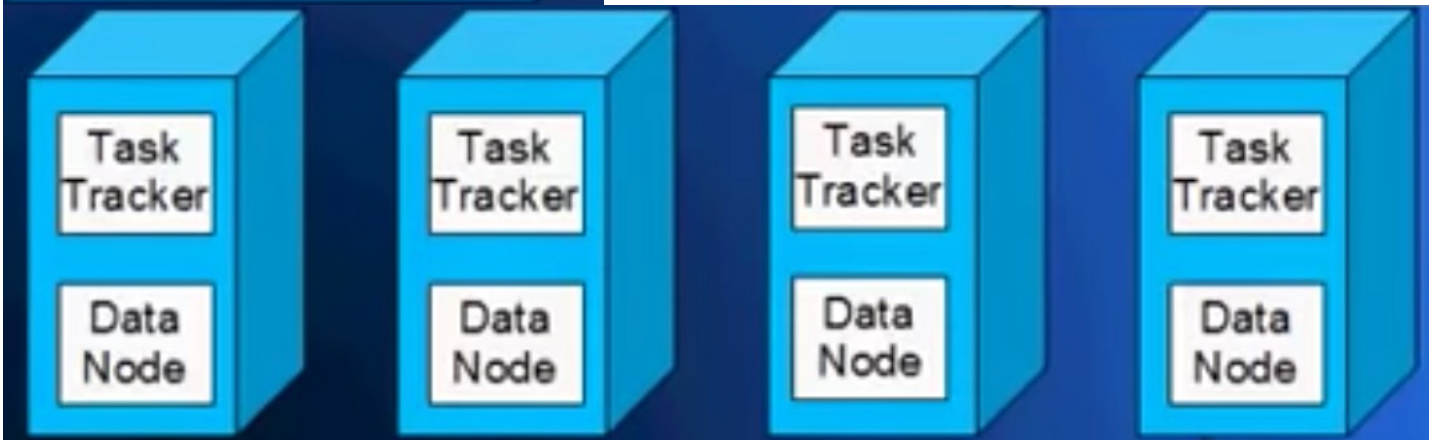


# MapReduce

Master




Slaves

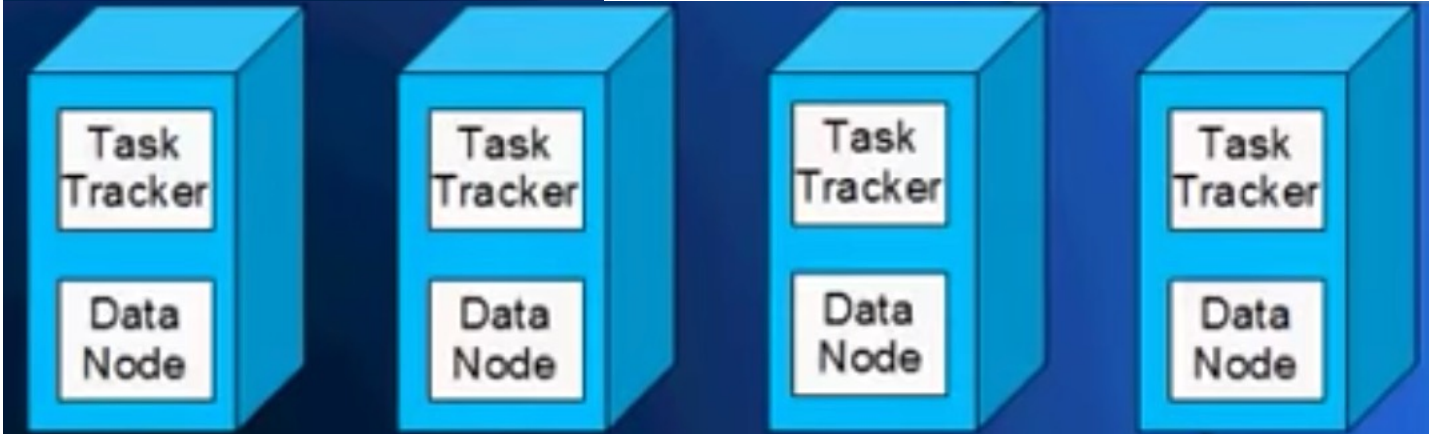


# Applications

Master




Slaves




# Batch Processing

Master

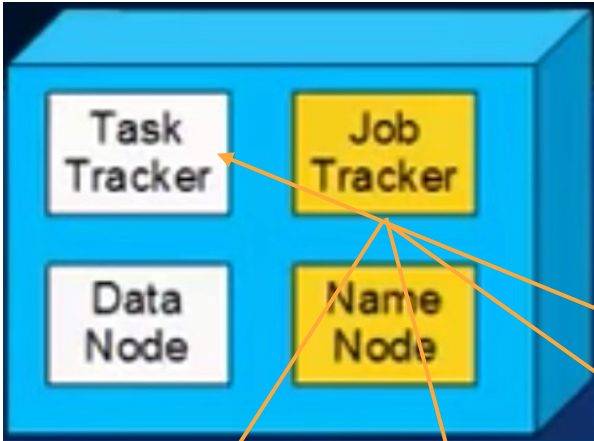


Slaves

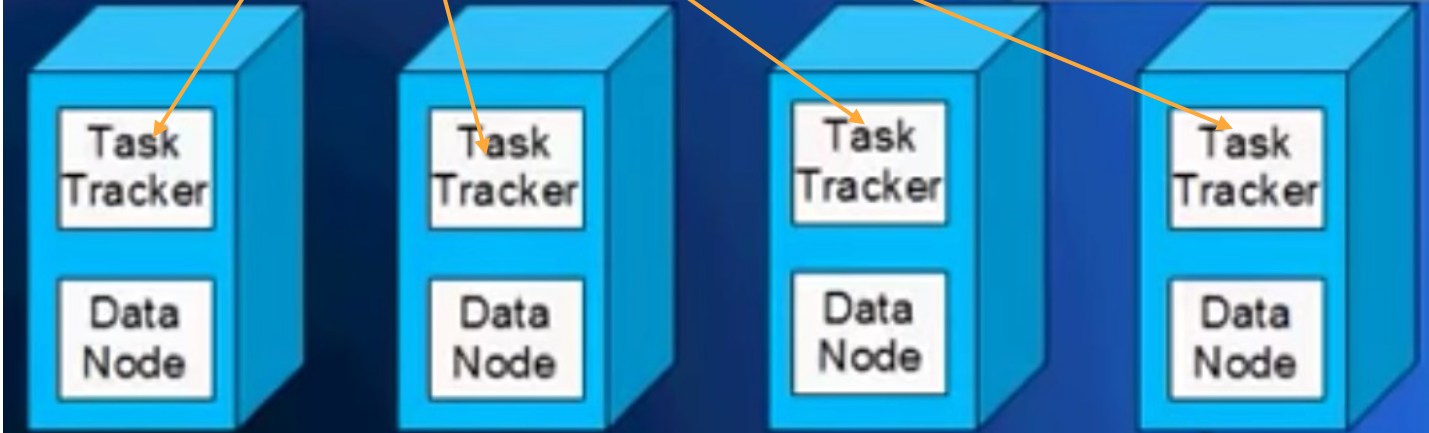


# Job Tracker

Master

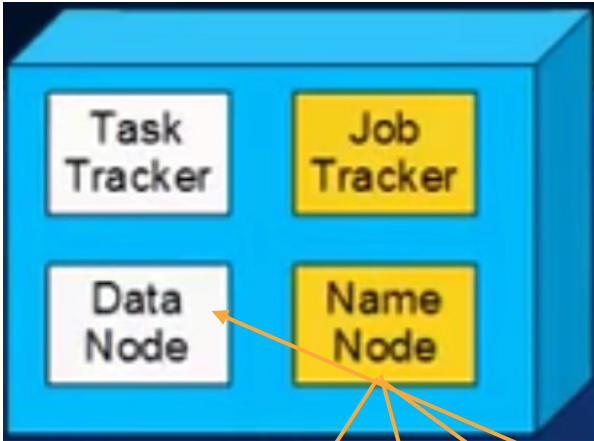


Slaves

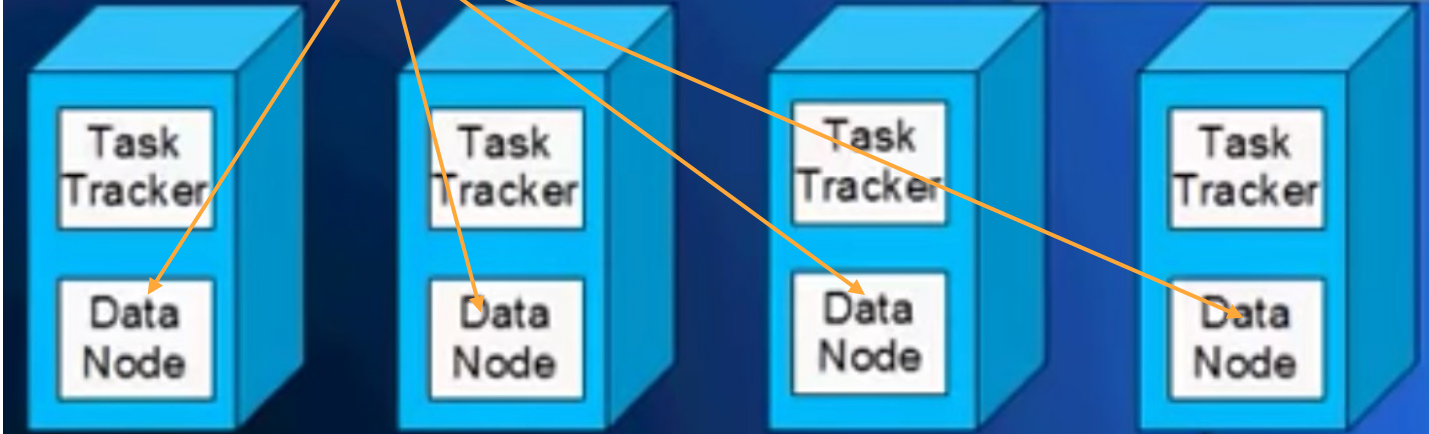


# Name Node

Master

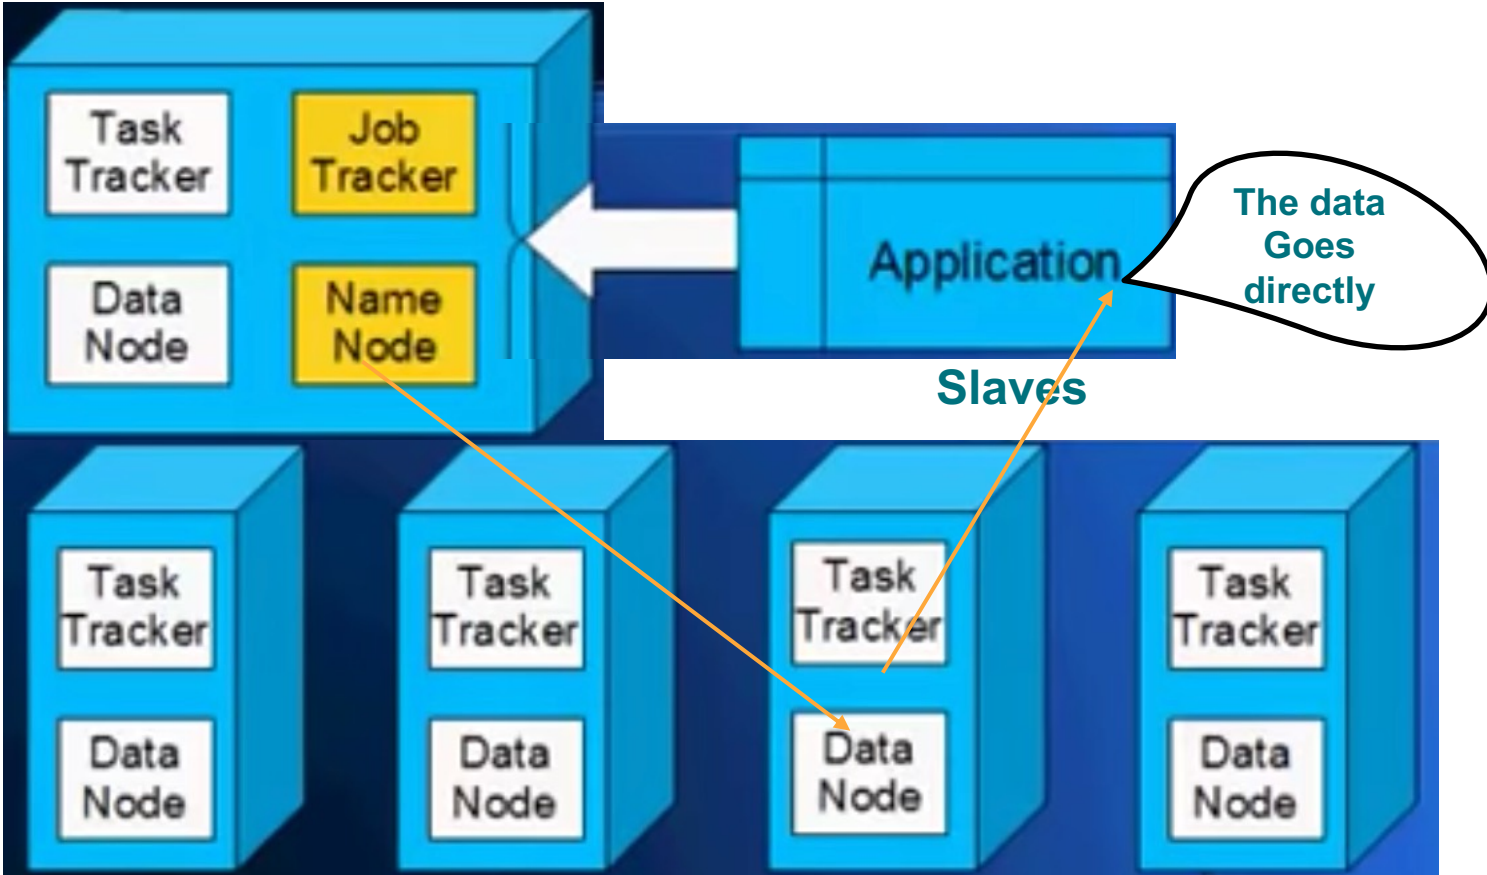


Slaves



# Data


Master





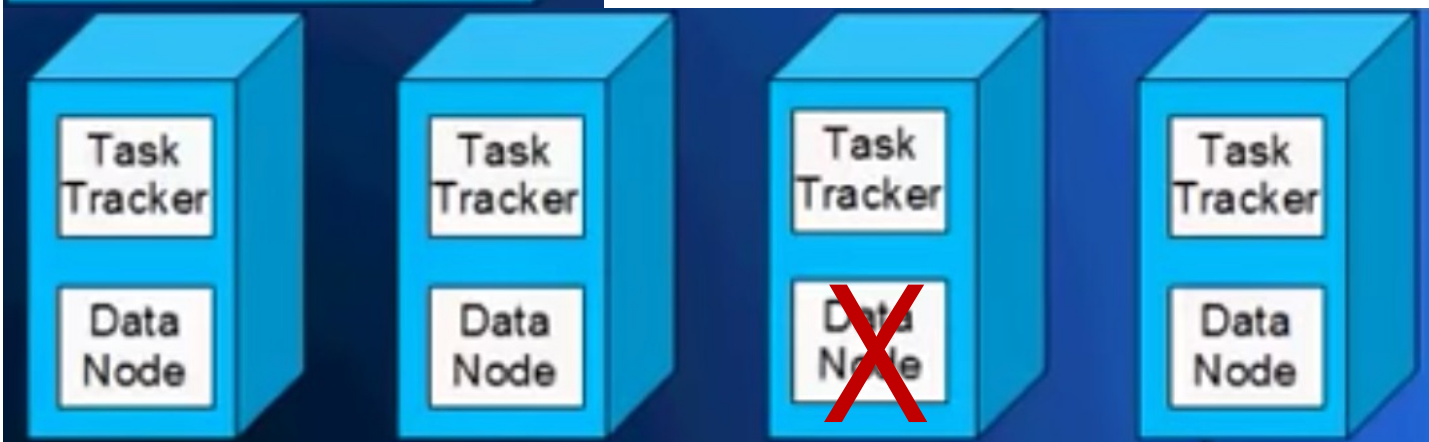
# Fault Tolerance for Data

Master



HDFS

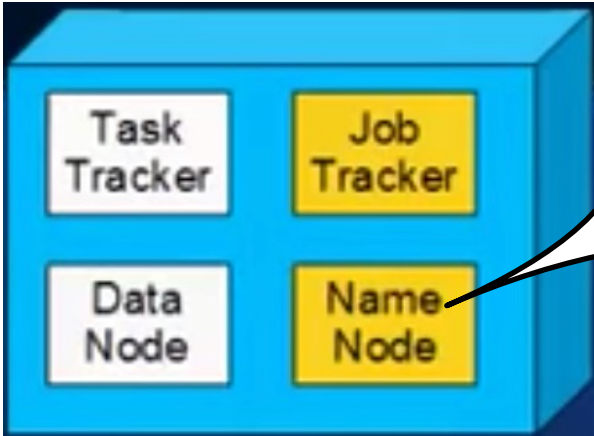
Slaves






# Master Backup

Master

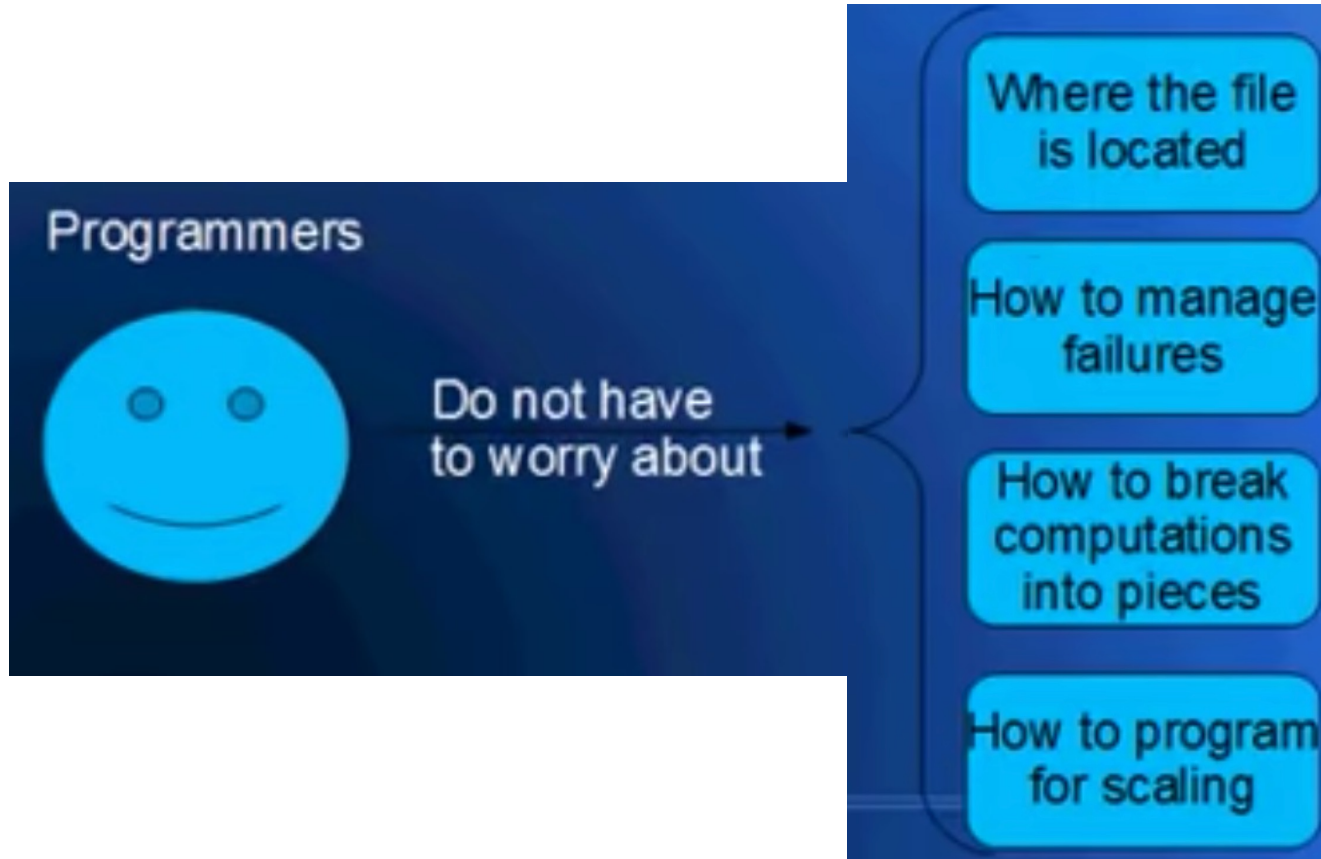


Tables are backed up

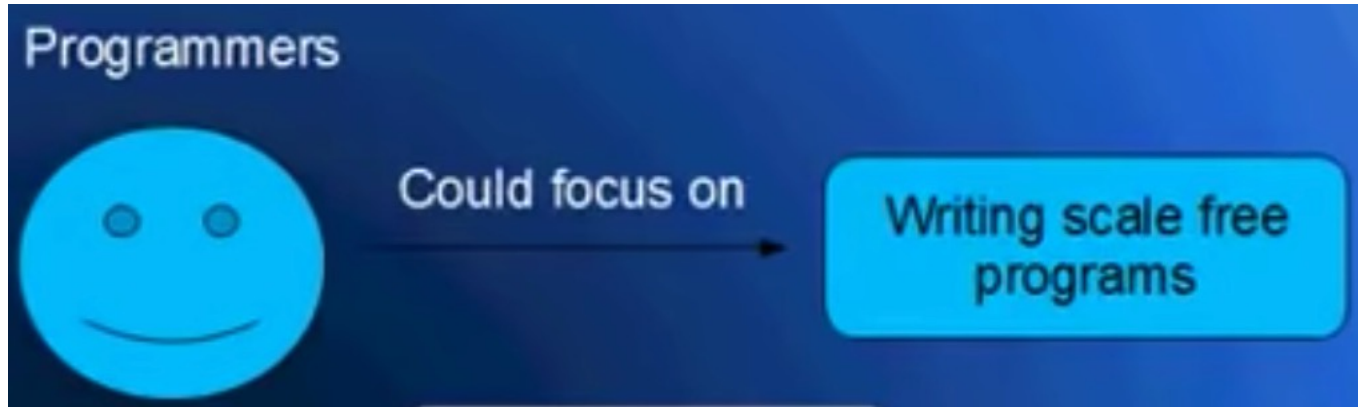
Slaves



# Easy Programming

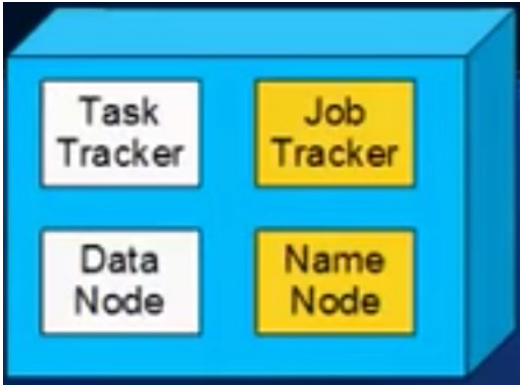


# Easy Programming



# Scalable


Master



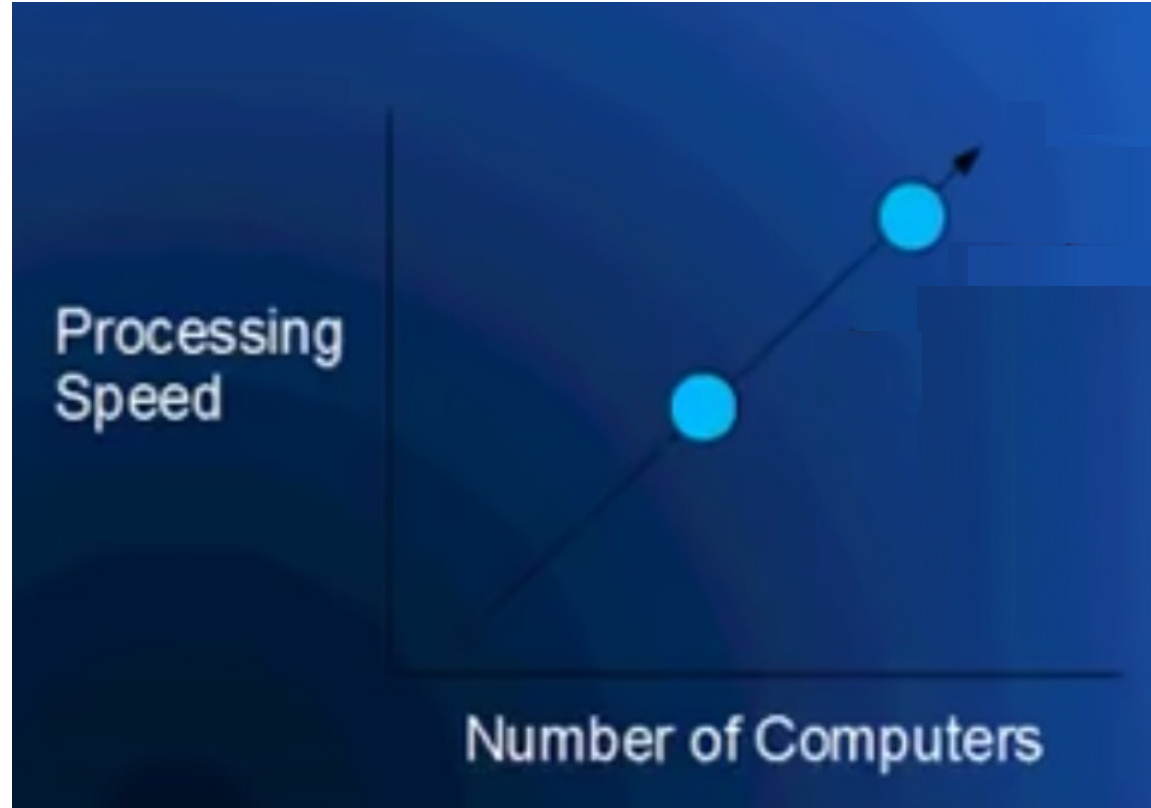
Slaves

Scalable

1000x



# Scalability Cost



# Solving Big Data problem w/ Hadoop


## Storing exponentially growing huge datasets

HDFS, storage unit of Hadoop is a Distributed File System




## Storing unstructured data

Allows to store any kind of data, be it structured, semi-structured or unstructured



## Processing data faster

Provides parallel processing of data present in HDFS  
Allows to process data locally i.e. each node works with a part of data which is stored on it



## **Needs management of components for distribution of work:**

- Compute resources (CPU/RAM)
- Long-Term Storage (Filesystem)
- User-supplied “algorithm” or “work”
- Scheduling/dependency



## **Provides framework/API for distributed computing**

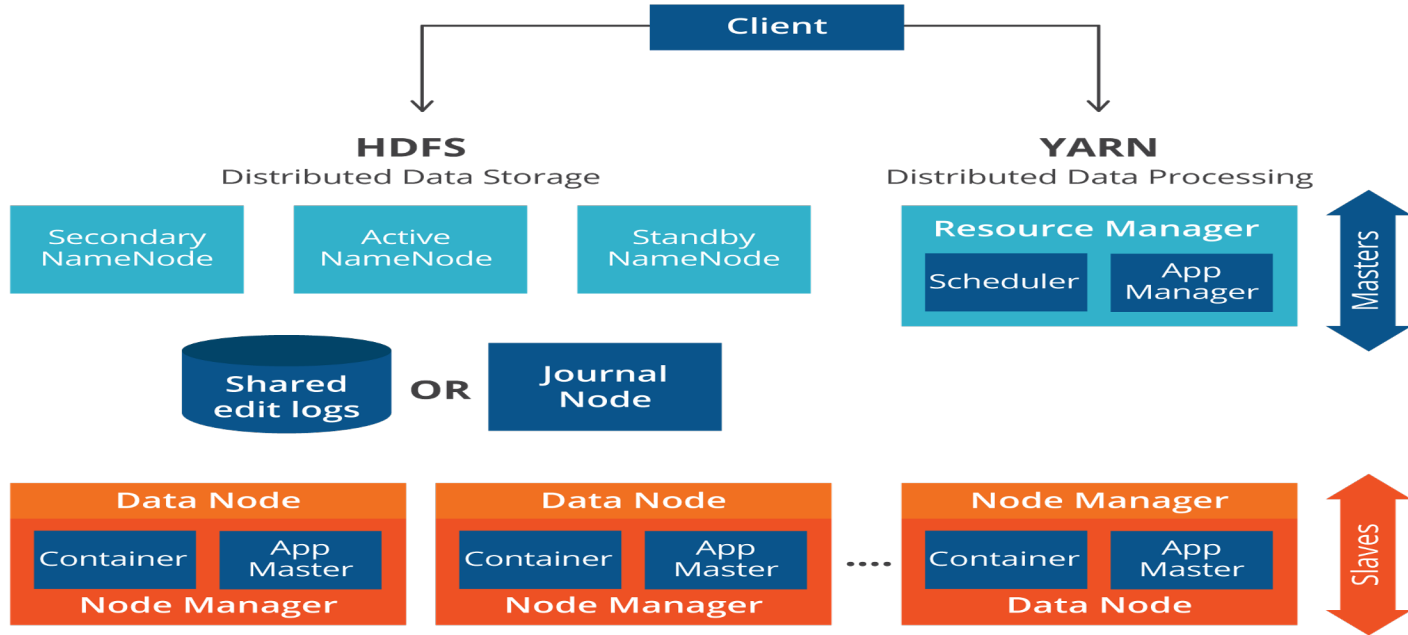
- Hide away distributed back-end
- Provides consolidated abstraction for:
  - Workload distribution
  - Storage and replication
  - Application development
  - Fault tolerance

# Hadoop Core Components

- Hadoop Common:
  - Common shared library
  - Runtime environment
- Hadoop Distributed File System (HDFS)
  - Storage abstraction
  - Fault tolerance
  - Data replication
- Hadoop YARN (Yet Another Resource Negotiator)
  - Compute management layer
  - Distributed OS

# Hadoop Architecture

## Apache Hadoop 2.0 and YARN



# Hadoop Properties

- At its core - Java application thus platform independent
- Data-centric
- Core tools & frameworks
- Cluster built out of “nodes” that provide resources - execution time, disk space, etc.
- Application development system, to build distributed solutions to user data problems

# Hadoop Common Components

- Basic primitives for programming environment:
  - Core data types (writable)
    - Text, Int (V), Boolean, Byte, Long (V), Double
  - Exceptions handling
  - Data structures
    - Arraywritable, TwoD Array Writable, Object Writable
  - IPC primitives
  - Logic, math, computation runtime
  - Data handling/encoding/transfer
    - Compressed files (.tar, .tar.gz, .tar.bz)

**Built from nodes, which serves processes deployed to nodes in your cluster:**

➤ **Namenode**

- Metadata is stored here
- Provides logical links to your data block

➤ **Datanode**

- Files are broken into smaller pieces and stored across these nodes

**Filesystem abstraction for Hadoop clusters Manage storage efficiently for:**

- Very large files
- Streaming data access
- Built on commodity hardware
- High availability
- Concurrent access

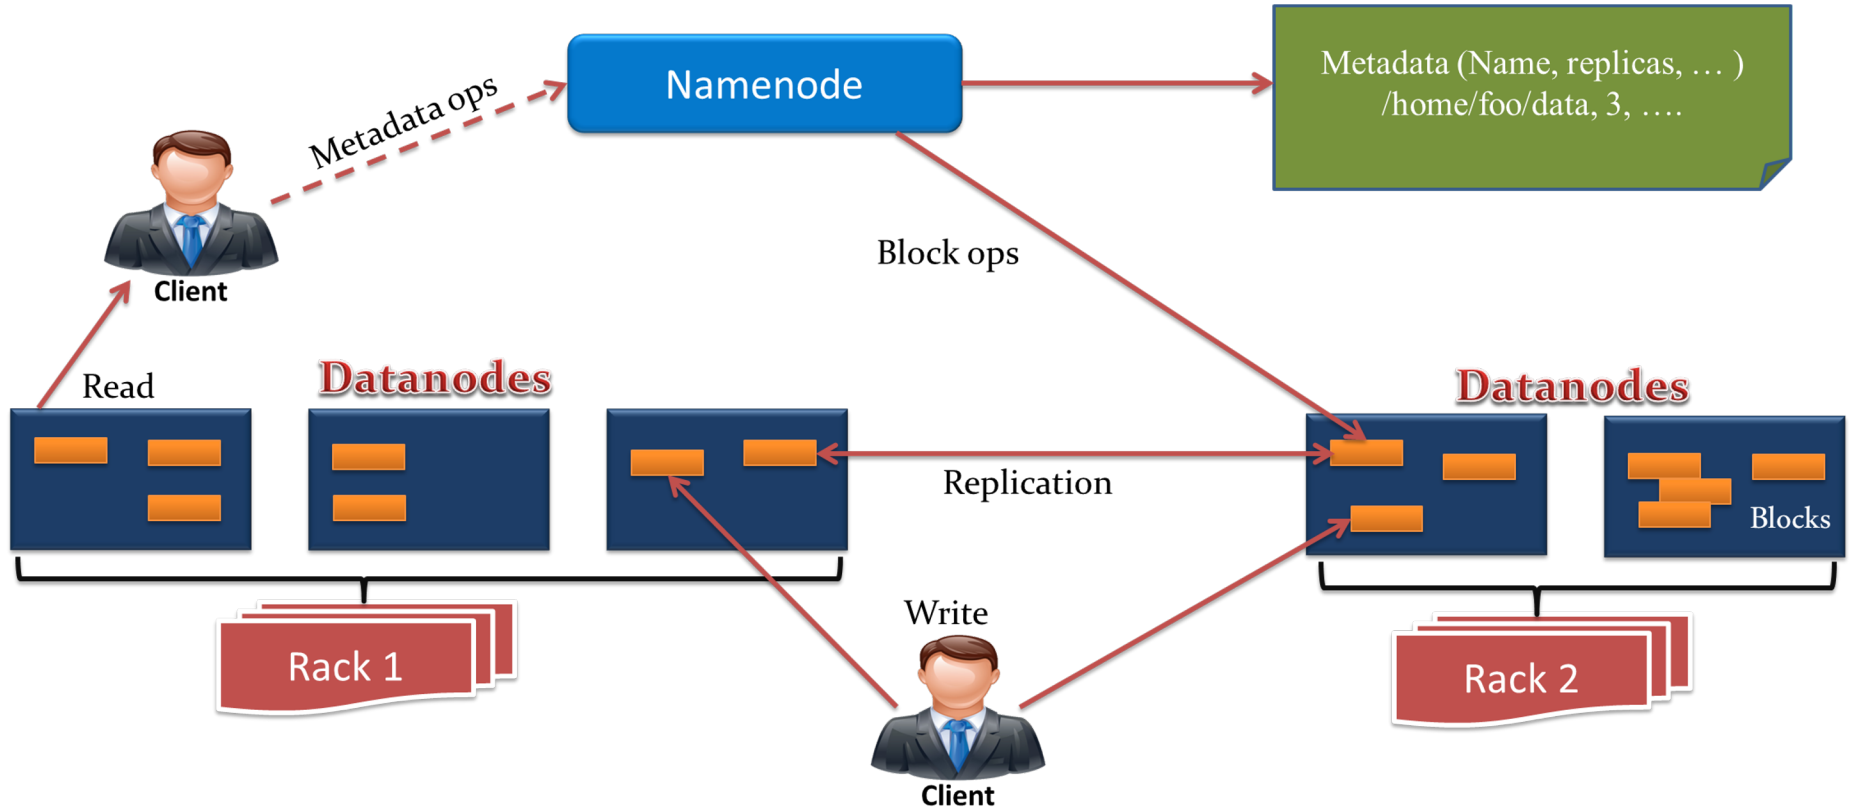


## **Doesn't work as well for**

- Many small files
- Low-latency data access

In these cases, you may always use another storage abstraction or engine

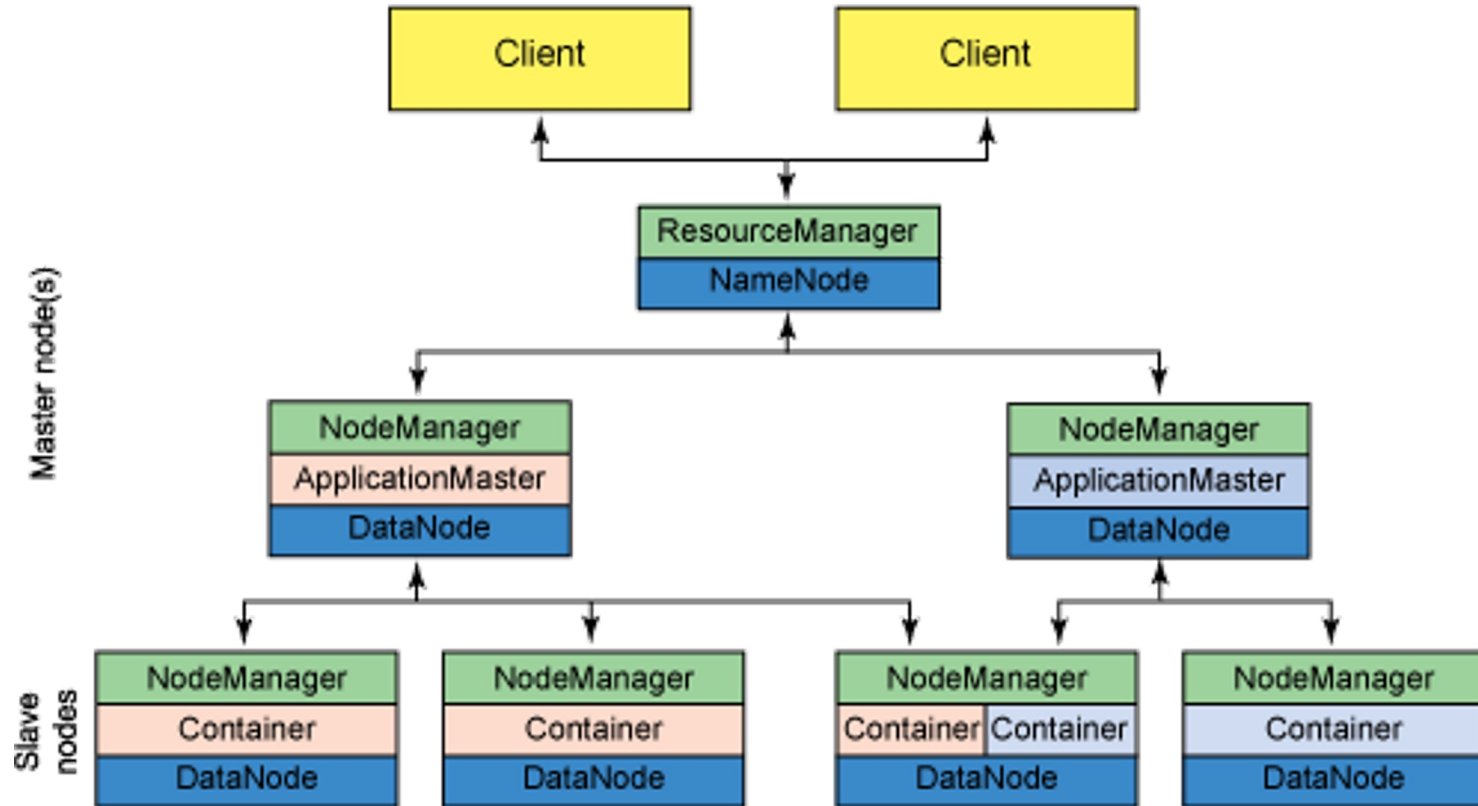
# HDFS Architecture



# YARN Architecture

- As Resource Manager
  - Manages the compute resources in the cluster
  - Receives work from clients
  
- As NodeManager
  - Manages workload within an execution node in the cluster
  - Instantiates containers to execute workloads
  - Distributes to other Node Managers on-demand

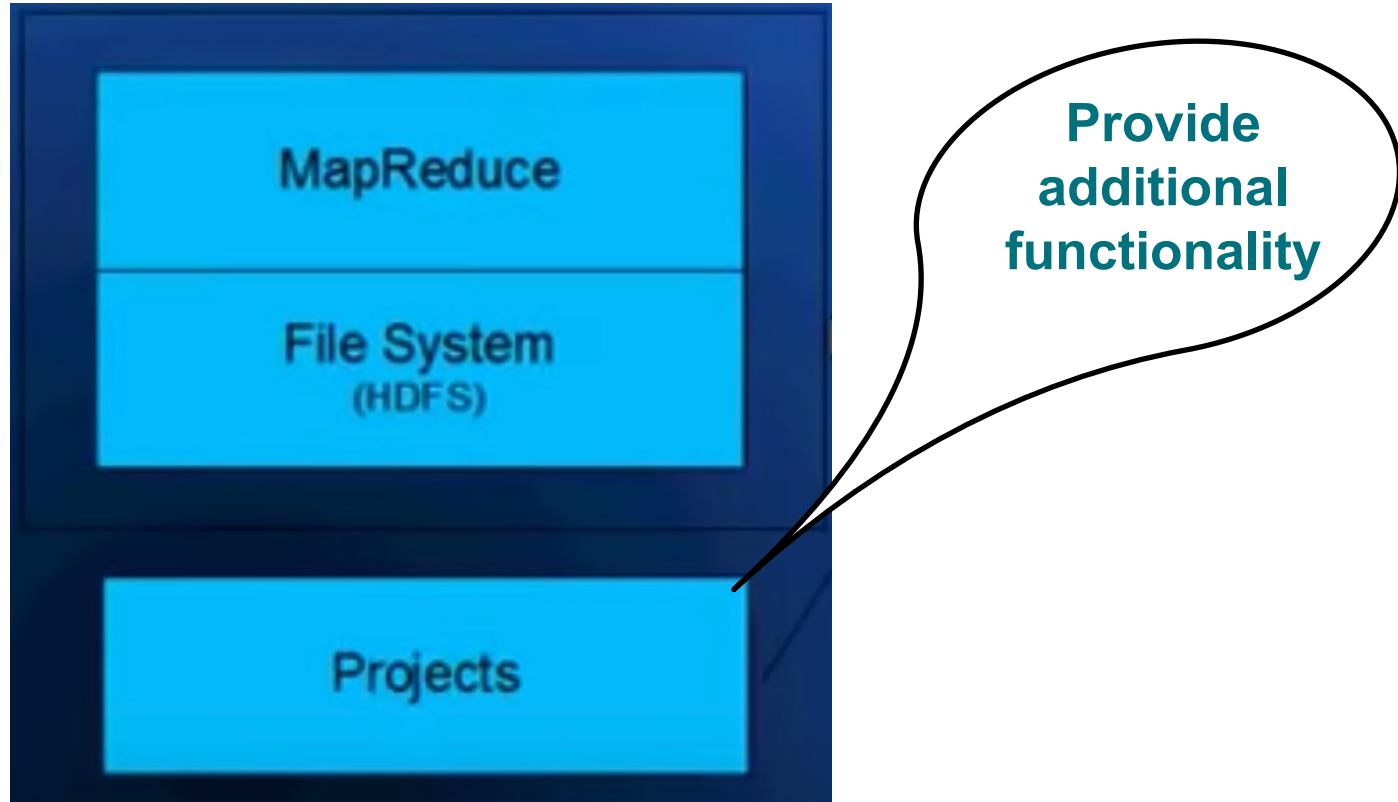
# YARN Architecture



## **The platform development layer for Hadoop - or “Cluster OS”**

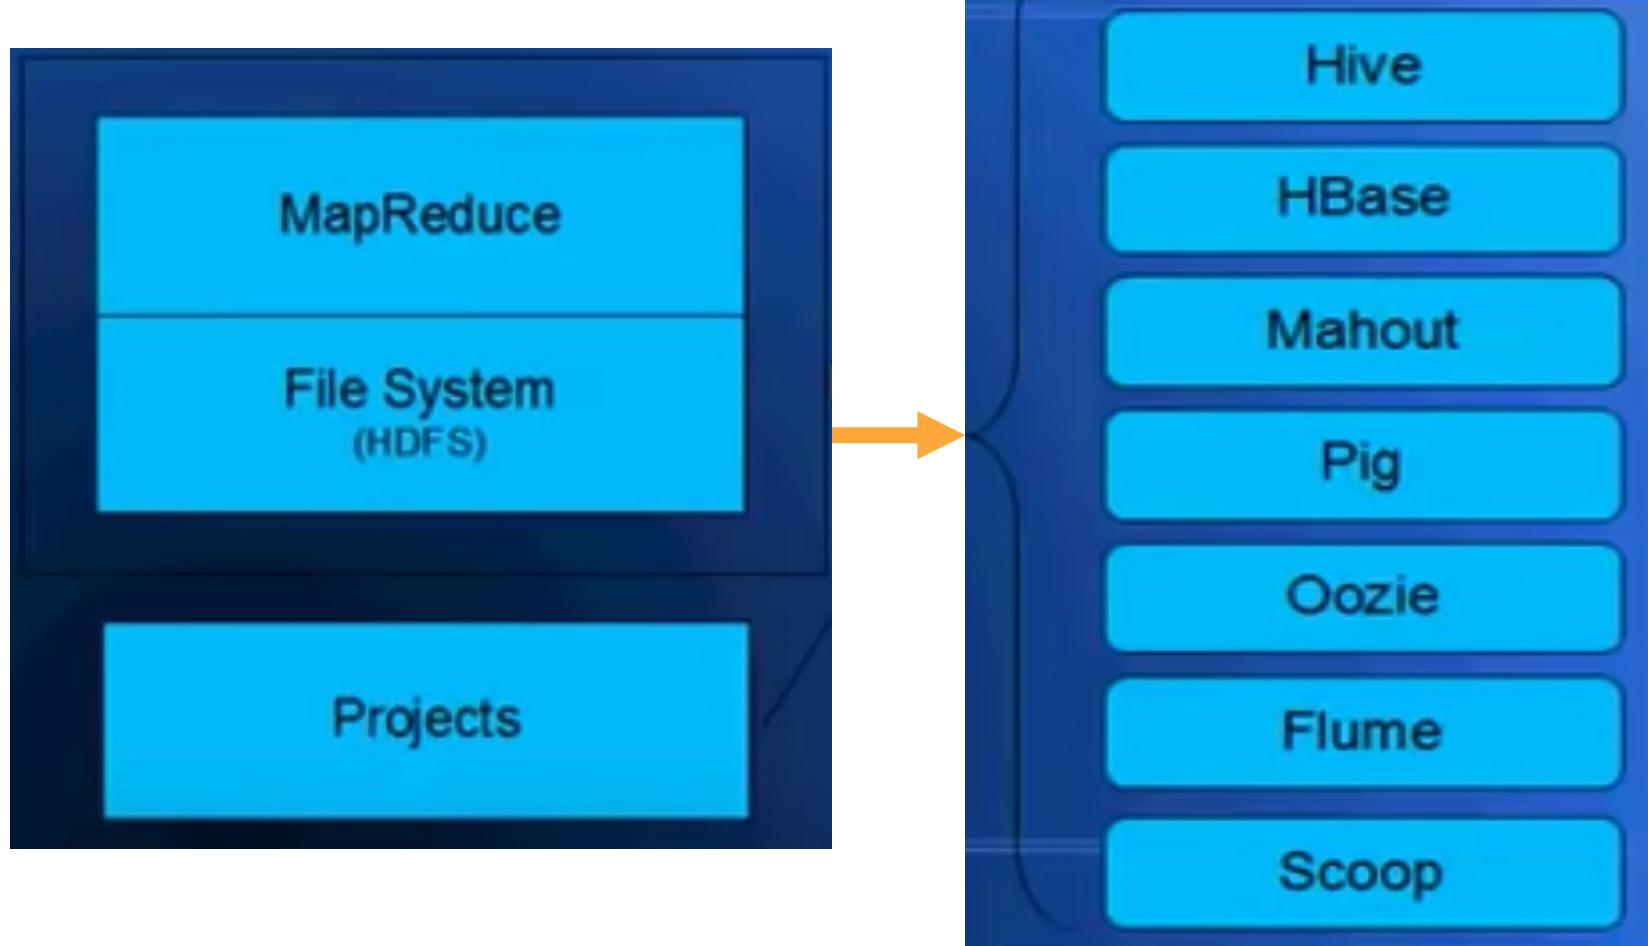
- In the context above,
  - It almost acts as a means of IaaS(M)
  - PaaS-like layers built atop this to provide “application interfaces” to Hadoop users.
- Examples include: MapReduce, Spark

# Projects



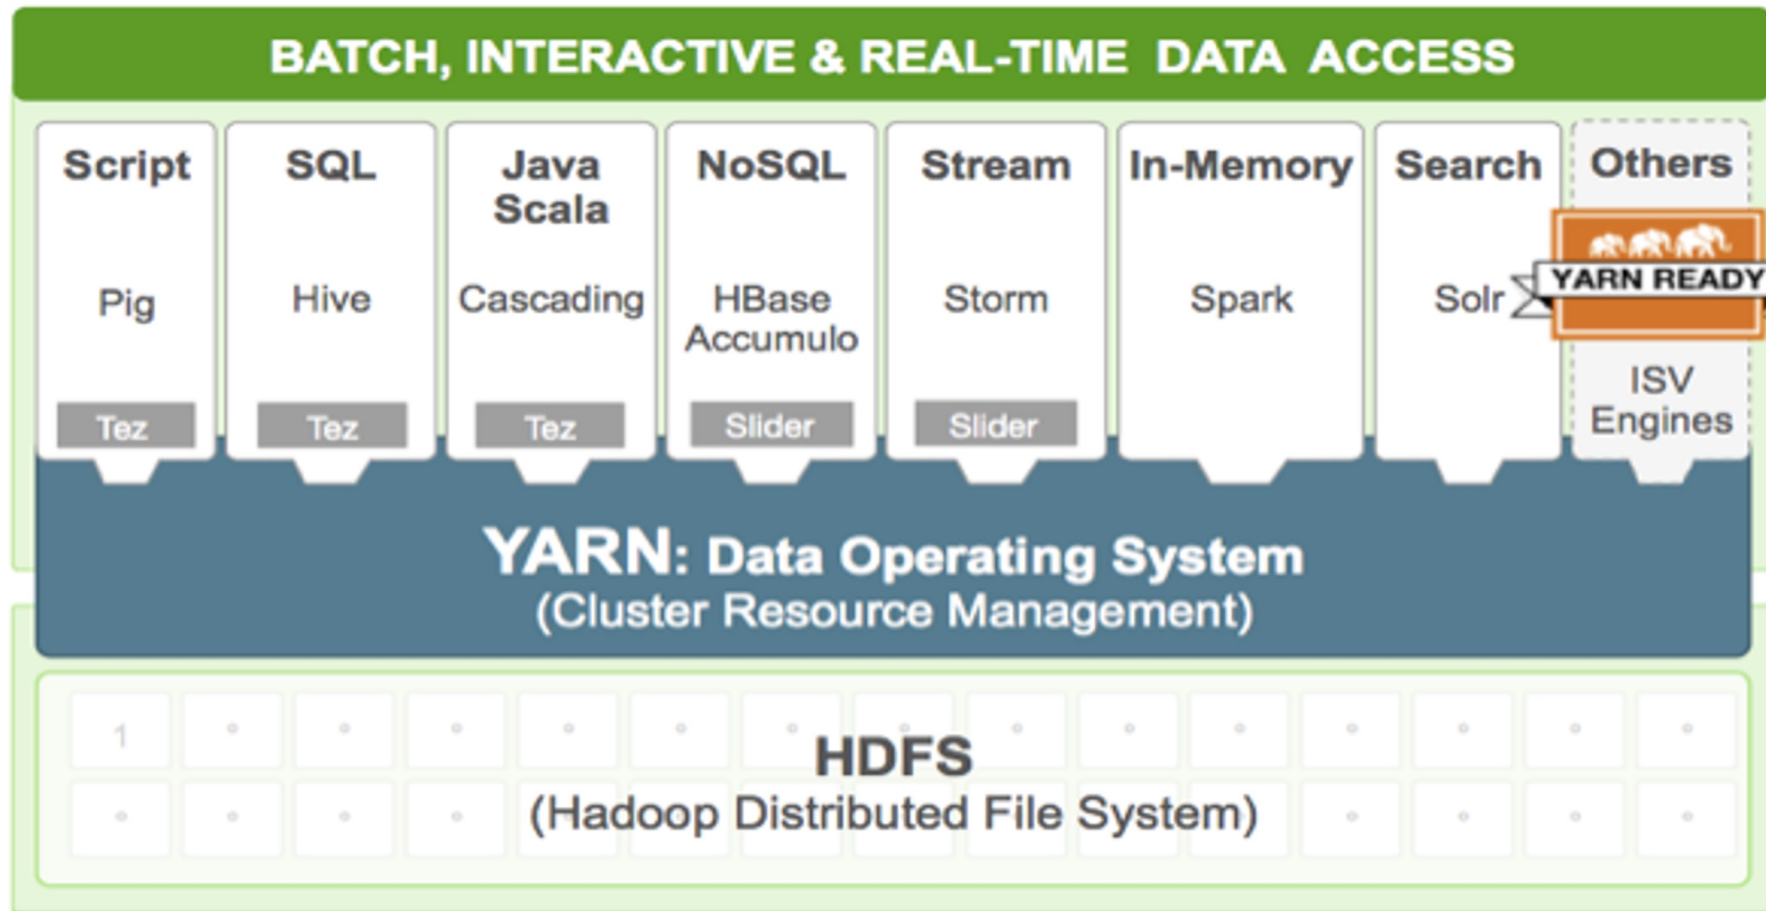
# Projects

Pass it on...   
Knowledge is power





# User, Hadoop, YARN, HDFS



# Application Flow in Hadoop

